

Distributional Semantics and Large Language Models

Original slides from João Sedoc (2019)
Updated 2023

Introduction to Human Language Technologies

Intuition of distributional word similarity

Nida example:

A bottle of ***tesgüino*** is on the table
Everybody likes ***tesgüino***
Tesgüino makes you drunk
We make ***tesgüino*** out of corn.

From context words humans can guess ***tesgüino*** means

- an alcoholic beverage like **beer**

Intuition for algorithm:

- Two words are similar if they have similar word contexts.

Distributional Hypothesis

If we consider **optometrist** and **eye-doctor** we find that, as our corpus of utterances grows, these two occur in almost the same environments. In contrast, there are many sentence environments in which **optometrist** occurs but **lawyer** does not...

It is a question of the relative frequency of such environments, and of what we will obtain if we ask an informant to substitute any word he wishes for **optometrist** (not asking what words have the same meaning).

These and similar tests all measure the probability of particular environments occurring with particular elements... If A and B have almost identical environments we say that they are synonyms.

—Zellig Harris (1954)

“You shall know a word by the company it keeps!”

—John Firth (1957)



Distributional models of meaning
= vector-space models of meaning
= vector semantics

Intuitions: Zellig Harris (1954):

- “oculist and eye-doctor ... occur in almost the same environments”
- “If A and B have almost identical environments we say that they are synonyms.”

Firth (1957):

- “You shall know a word by the company it keeps!”

Intuition

Model the meaning of a word by “embedding” in a vector space.

The meaning of a word is a vector of numbers

- Vector models are also called “**embeddings**”.

Contrast: word meaning is represented in many computational linguistic applications by a vocabulary index (“word number 545”)

$$\text{vec}(\text{"dog"}) = (0.2, -0.3, 1.5, \dots)$$

$$\text{vec}(\text{"bites"}) = (0.5, 1.0, -0.4, \dots)$$

$$\text{vec}(\text{"man"}) = (-0.1, 2.3, -1.5, \dots)$$

Term-document matrix

Term-document matrix

Each cell: count of term t in a document d : $\text{tf}_{t,d}$:

- Each document is a count vector in \mathbb{N}^v : a column below

	As You Like It	Twelfth Night	Julius Caesar	Henry V
battle	1	1	8	15
soldier	2	2	12	36
fool	37	58	1	5
clown	6	117	0	0

Term-document matrix

Two documents are similar if their vectors are similar

	As You Like It	Twelfth Night	Julius Caesar	Henry V
battle	1	1	8	15
soldier	2	2	12	36
fool	37	58	1	5
clown	6	117	0	0

The words in a term-document matrix

Each word is a **count vector** in \mathbb{N}^D : a row below

	As You Like It	Twelfth Night	Julius Caesar	Henry V
battle	1	1	8	15
soldier	2	2	12	36
fool	37	58	1	5
clown	6	117	0	0

The words in a term-document matrix

Two **words** are similar if their vectors are similar

	As You Like It	Twelfth Night	Julius Caesar	Henry V
battle	1	1	8	15
soldier	2	2	12	36
fool	37	58	1	5
clown	6	117	0	0

Brilliant insight: Use running text as implicitly supervised training data!

- A word s near *fine*
 - Acts as gold ‘correct answer’ to the question
 - “Is word w likely to show up near *fine*? ”
- No need for hand-labeled supervision
- The idea comes from **neural language modeling**
 - Bengio et al. (2003)
 - Collobert et al. (2011)

Word2vec

Popular embedding method

Very fast to train

Code available on the web


Idea: **predict** rather than **count**

Word2vec


- Instead of **counting** how often each word w occurs near “*fine*”
- Train a classifier on a binary **prediction task**:
 - Is w likely to show up near “*fine*”?
- We don’t actually care about this task
 - But we’ll take the learned classifier weights as the word embeddings

Word2vec

CBOW



SKIPGRAM




I am selling these fine leather jackets

CBOW – high level

The resulting projection matrix P is the embedding matrix


Goal: Predict the middle word given the words of the context



Skip-gram – high level

The resulting projection matrix P is the embedding matrix

Goal: Predict the context words given the middle word



Dense embeddings you can
download!



Word2vec (Mikolov et al.)

<https://code.google.com/archive/p/word2vec/>



Fasttext <http://www.fasttext.cc/>

Glove (Pennington, Socher, Manning)

<http://nlp.stanford.edu/projects/glove/>



Why vector models of meaning? Computing the similarity between words

“**fast**” is similar to “**rapid**”

“**tall**” is similar to “**height**”

Question answering:



*Q: “How **tall** is Mt. Everest?”*

*Candidate A: “The official **height** of Mount Everest is 29029 feet”*

Analogy: Embeddings capture relational meaning!

$\text{vector('king')} - \text{vector('man')} + \text{vector('woman')} \approx \text{vector('queen')}$

$\text{vector('Paris')} - \text{vector('France')} + \text{vector('Italy')} \approx \text{vector('Rome')}$



Evaluating similarity

Extrinsic (task-based, end-to-end) Evaluation:

- Question Answering
- Spell Checking
- Essay grading

Intrinsic Evaluation:

- Correlation between algorithm and human word similarity ratings
 - Wordsim353: 353 noun pairs rated 0-10. $sim(plane,car)=5.77$
- Taking TOEFL multiple-choice vocabulary tests
- Levied is closest in meaning to:
imposed, believed, requested, correlated

Evaluating embeddings

Compare to human scores on word similarity-type tasks:


- WordSim-353 (Finkelstein et al., 2002)
- SimLex-999 (Hill et al., 2015)
- Stanford Contextual Word Similarity (SCWS) dataset (Huang et al., 2012)
- TOEFL dataset: *Levied is closest in meaning to: imposed, believed, requested, correlated*

Intrinsic evaluation

cos sim	Psycho linguistic experiment		mean 10 judges	WordSim 353
	<u>↑</u>	<u>↓</u>		
	Love , sex		6.8	
:	tiger , cat		1.3	
:	tiger , tiger		10	
:	fertility , egg		6.7	
:	stock , egg		1.8	
	professor , cucumber		0.3	GOLD STANDARD

Intrinsic evaluation

Rate the following word pairs according to how similar they are by moving the slider



Measuring similarity

Given 2 target words v and w

We'll need a way to measure their similarity.

Most measure of vectors similarity are based on the:


Cosine between embeddings!

The similarity between two vectors v and w is:

$$\frac{v \cdot w}{\|v\| \|w\|}$$

- High when two vectors have large values in same dimensions.
- Low (in fact 0) for **orthogonal vectors** with zeros in complementary distribution

Visualizing vectors and angles



	large	data
apricot	2	0
digital	0	1
information	1	6

Bias in Word Embeddings

Extreme *she* occupations

- | | | |
|-----------------|-----------------------|------------------------|
| 1. homemaker | 2. nurse | 3. receptionist |
| 4. librarian | 5. socialite | 6. hairdresser |
| 7. nanny | 8. bookkeeper | 9. stylist |
| 10. housekeeper | 11. interior designer | 12. guidance counselor |

Extreme *he* occupations

- | | | |
|----------------|------------------|----------------|
| 1. maestro | 2. skipper | 3. protege |
| 4. philosopher | 5. captain | 6. architect |
| 7. financier | 8. warrior | 9. broadcaster |
| 10. magician | 11. figher pilot | 12. boss |

More to come on bias later ...


Embeddings are the workhorse of NLP

- Used as pre-initialization for language models, neural MT, classification, NER systems...
- Downloaded and easily trainable
- Available in ~100s of languages
- And ... **contextualized word embeddings** are built on top of them

Contextualized Word Embeddings




BERT - Deep Bidirectional Transformers




From last time...

- Multiple (N) layers
- For encoder-decoder attention, Q: previous decoder layer, K and V: output of encoder
- For encoder self-attention, Q/K/V all come from previous encoder layer
- For decoder self-attention, allow each position to attend to all positions up to that position
- Positional encoding for word order



Transformer



Training BERT

- BERT has two training objectives:
 1. Predict missing (masked) words
 2. Predict if a sentence is the next sentence

BERT- predict missing words


Use the output of the masked word's position to predict the masked word

Possible classes:
All English words

0.1%	Aardvark
...	...
10%	Improvisation
...	...
0%	Zyyzyva

FFNN + Softmax

Randomly mask
15% of tokens



Input

[CLS] Let's stick to improvisation in this skit


BERT's clever language modeling task masks 15% of words in the input and asks the model to predict the missing word.

BERT- predict is next sentence?





Predict likelihood
that sentence B
belongs after
sentence A

Tokenized
Input


Input




The second task BERT is pre-trained on is a two-sentence classification task. The tokenization is oversimplified in this graphic as BERT actually uses WordPieces as tokens rather than words --- so some words are broken down into smaller chunks.




BERT on tasks




(a) Sentence Pair Classification Tasks:
MNLI, QQP, QNLI, STS-B, MRPC,
RTE, SWAG



(b) Single Sentence Classification Tasks:
SST-2, CoLA



(c) Question Answering Tasks:
SQuAD v1.1



(d) Single Sentence Tagging Tasks:
CoNLL-2003 NER

Take-Aways

- Distributional semantics – learn a word’s “meaning” by its context
- Simplest representation is frequency in documents
- Word embeddings (Word2Vec, GloVe, FastText) predict rather than use co-occurrence
- Similarity measured often using cosine
- Intrinsic evaluation uses correlation with human judgements of word similarity
- Contextualized embeddings are the new stars of NLP

Language Models is All You Need


- Contextualized word embeddings (such as BERT) were shown to be useful to build actual applications
Summarization, Question Answering, Information Extraction, ...
- Language models were trained to obtain such word embeddings
- **But why not just use the language models?**
- Seminal paper from OpenAI (2018): “Language Models are Unsupervised Multitask Learners”
Convert any NLP into a text continuation problem

Stages of Training

- Train a language model on raw text
 - Up to trillion of tokens
 - Arms race on building bigger and bigger models
- Fine-tune on instruction data
- Reinforcement learning from human feedback

Decoder-Only Models

- Alternative architecture: Just decoder of Transformer model
⇒ no input, only self-attention
- Trained with next-word prediction



Instruction Training


- Examples of requests and responses constructed by human annotators
- May be collected from actual user requests and edited by experts
- May be generated from existing data sets

Question Answering What is the highest mountain in the world? The highest mountain in the world is Mount Everest.	Summarization Summarize the following paragraph into one sentence. The Federal Reserve paused its campaign of interest rate increases for the first time in more than a year. But officials suggested that rates would rise more in 2023, as inflation remains "well above" the central bank's target. Summary: No interest rate rise for now but maybe later in the year.	Translation Translate from English to German. English: My name is Ozymandias, King of Kings; Look on my Works, ye Mighty, and despair! German: Mein Name ist Ozymandias, König der Könige; Schau auf meine Werke, du Mächtiger, und verzweifle!
--	--	---

Reinforcement Learning from Human Feedback



- Machine generates multiple responses to a prompt
- Human annotators rank them
- Train a reward model
 - predict human annotation
 - can be applied to any text generated by model
 - normalized as reward function
- Fine-tune model with reward model
 - model generates response to prompt
 - reward function assesses response
 - if low reward, model needs to change



Prompt Engineering

- How a task is presented to the language model matters
- Black art: often unclear what is in the training data (especially the instruction data)
- “Think step by step”
 - Language model has very limited working memory
 - Complex reasoning may require several inference steps
 - By allowing model to verbalize intermediate steps, working memory is created
- “Imagine you are...” creates style or sets a context