

EN. 601.467/667

Introduction to Human Language Technology

End-to-End Speech Recognition

Shinji Watanabe



Today's agenda

- Introduction to end-to-end speech recognition
- Connectionist Temporal Classification (CTC)
- Attention



Frederick Jelinek (1932 –2010)

Statistical speech recognition and machine
translation

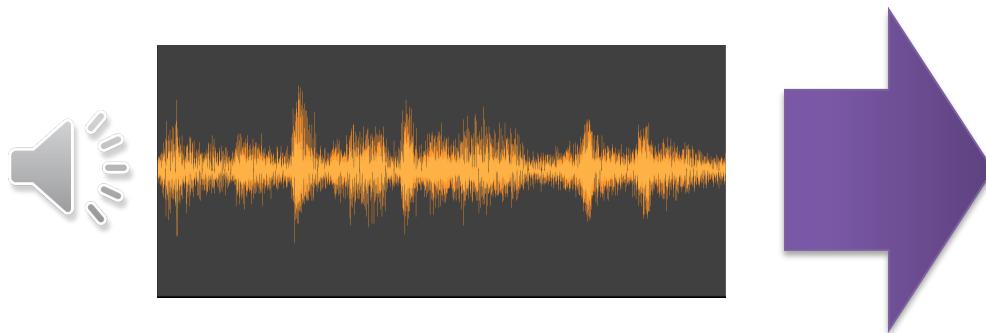
1972 - 1993: IBM

1993 - 2010: JHU

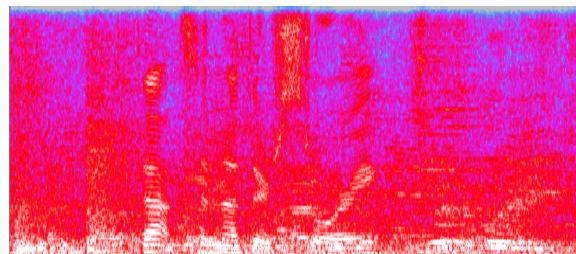
Jelinek methodology (1970s-)

Jelinek methodology (1970s-)

- Automatic Speech Recognition: Mapping *physical signal sequence* to *linguistic symbol sequence*



$$X = \{x_l \in \mathbb{Z} | l = 1, \dots, L\}$$
$$L = 43263$$



$$X = \{\mathbf{x}_t \in \mathbb{C}^D | t = 1, \dots, T\}$$
$$T = 268$$

“That’s another story”

$$W = \{w_n \in \mathcal{V} | n = 1, \dots, N\}$$
$$N = 3$$

Jelinek methodology (1970s-)

$$\arg \max_W p(W|O)$$

O : Speech sequence
 W : Text sequence

Jelinek methodology (1970s-)

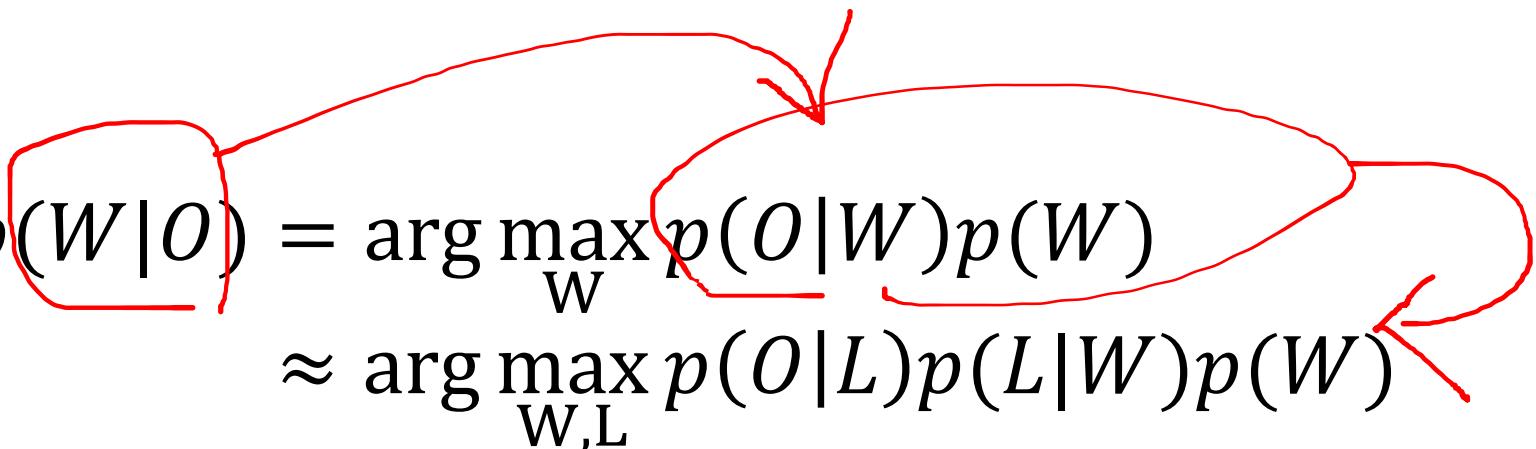
$$\arg \max_W p(W|O) = \arg \max_W p(O|W)p(W) \approx \arg \max_{W,L} p(O|L)p(L|W)p(W)$$

L: Phoneme sequence

- **Speech recognition**

- $p(O|L)$: Acoustic model (Hidden Markov model)
- $p(L|W)$: Lexicon
- $p(W)$: Language model (n-gram)

Jelinek methodology (1970s-)

$$\begin{aligned} \arg \max_W p(W|O) &= \arg \max_W p(O|W)p(W) \\ &\approx \arg \max_{W,L} p(O|L)p(L|W)p(W) \end{aligned}$$


- **Speech recognition**

- $p(O|L)$: Acoustic model (Hidden Markov model)
- $p(L|W)$: Lexicon
- $p(W)$: Language model (n-gram)

- Factorization
- **Conditional independence (Markov) assumptions, CIA**

Jelinek methodology (1970s-)

$$\arg \max_W p(W|O) = \arg \max_W p(O|W)p(W)$$

- **Machine translation**
 - $p(O|W)$: Translation model
 - $p(W)$: Language model

Jelinek methodology (1970s-)

$$\begin{aligned}\arg \max_W p(W|O) &= \arg \max_W p(O|W)p(W) \\ &\approx \arg \max_{W,L} p(O|L)p(L|W)p(W)\end{aligned}$$

- **Speech recognition**
 - $p(O|L)$: Acoustic model (Hidden Markov model)
 - $p(L|W)$: Lexicon
 - $p(W)$: Language model (n-gram)
- Continued 40 years

Jelinek methodology (1970s-)

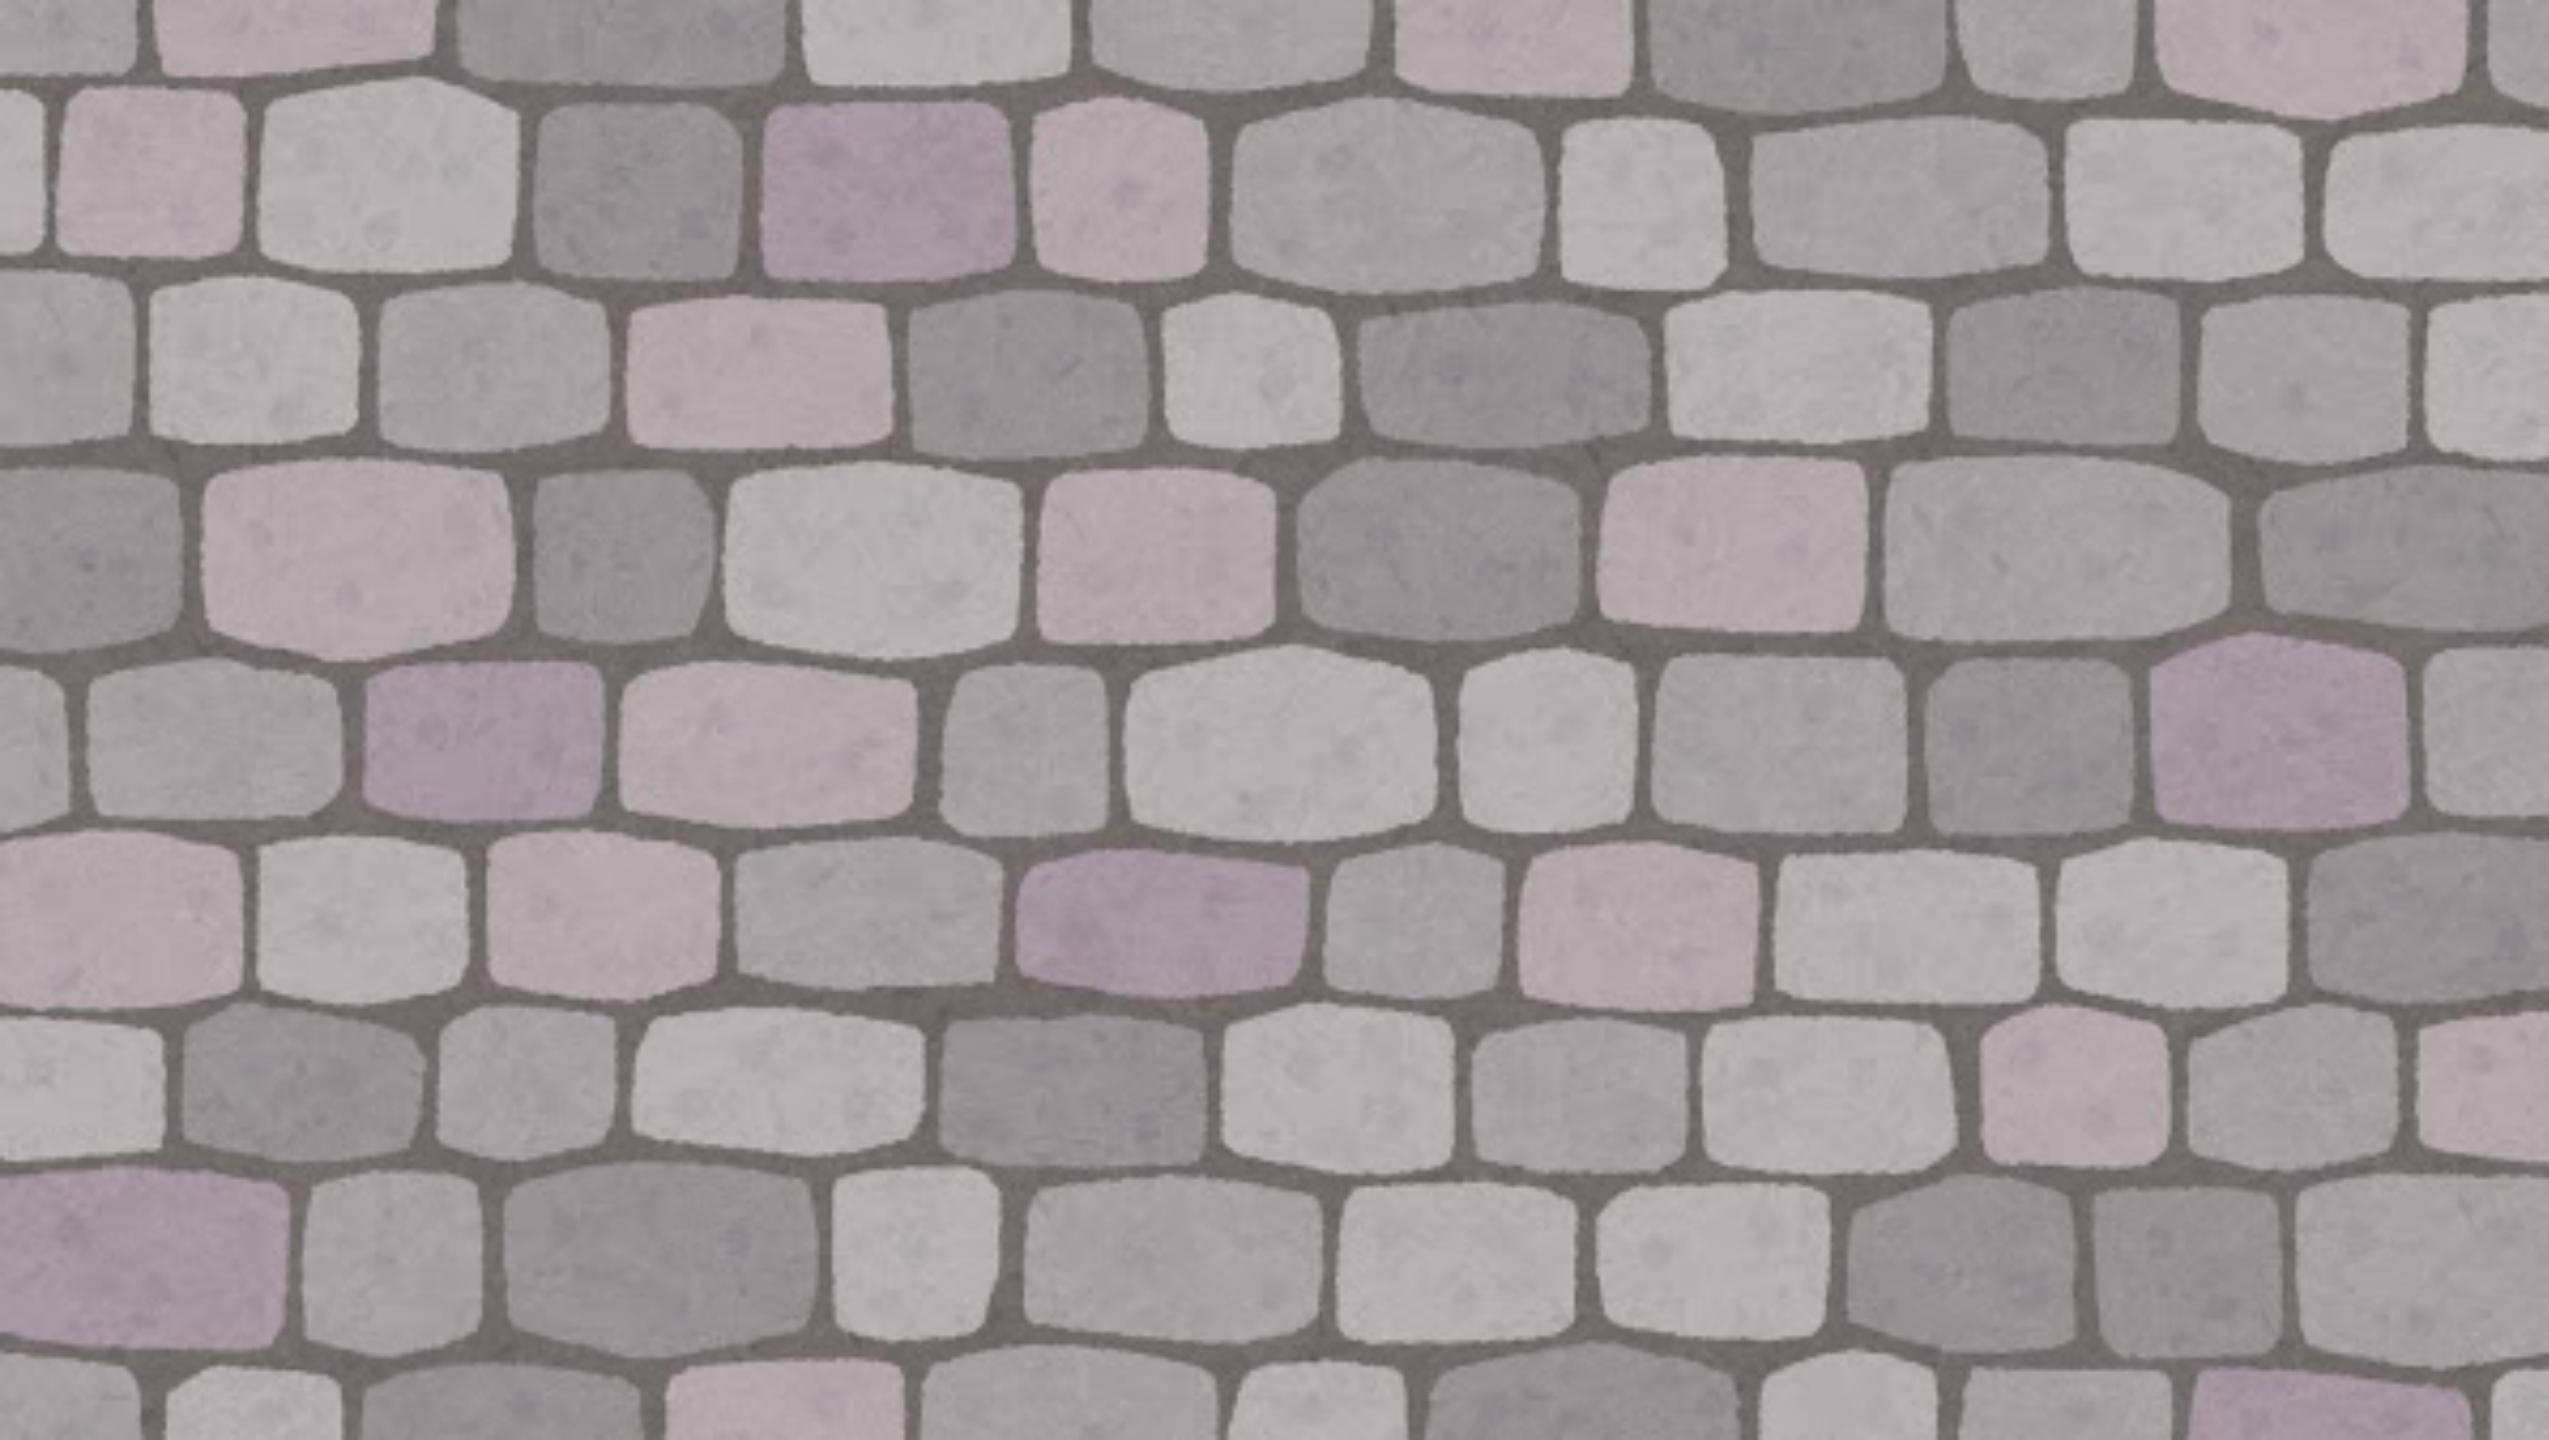
$$\begin{aligned}\arg \max_W p(W|O) &= \arg \max_W p(O|W)p(W) \\ &\approx \arg \max_{W,L} p(O|L)p(L|W)p(W)\end{aligned}$$

- **Speech recognition**
 - $p(O|L)$: Acoustic model
 - $p(L|W)$: Lexicon
 - $p(W)$: Language model
- Continued 40 years

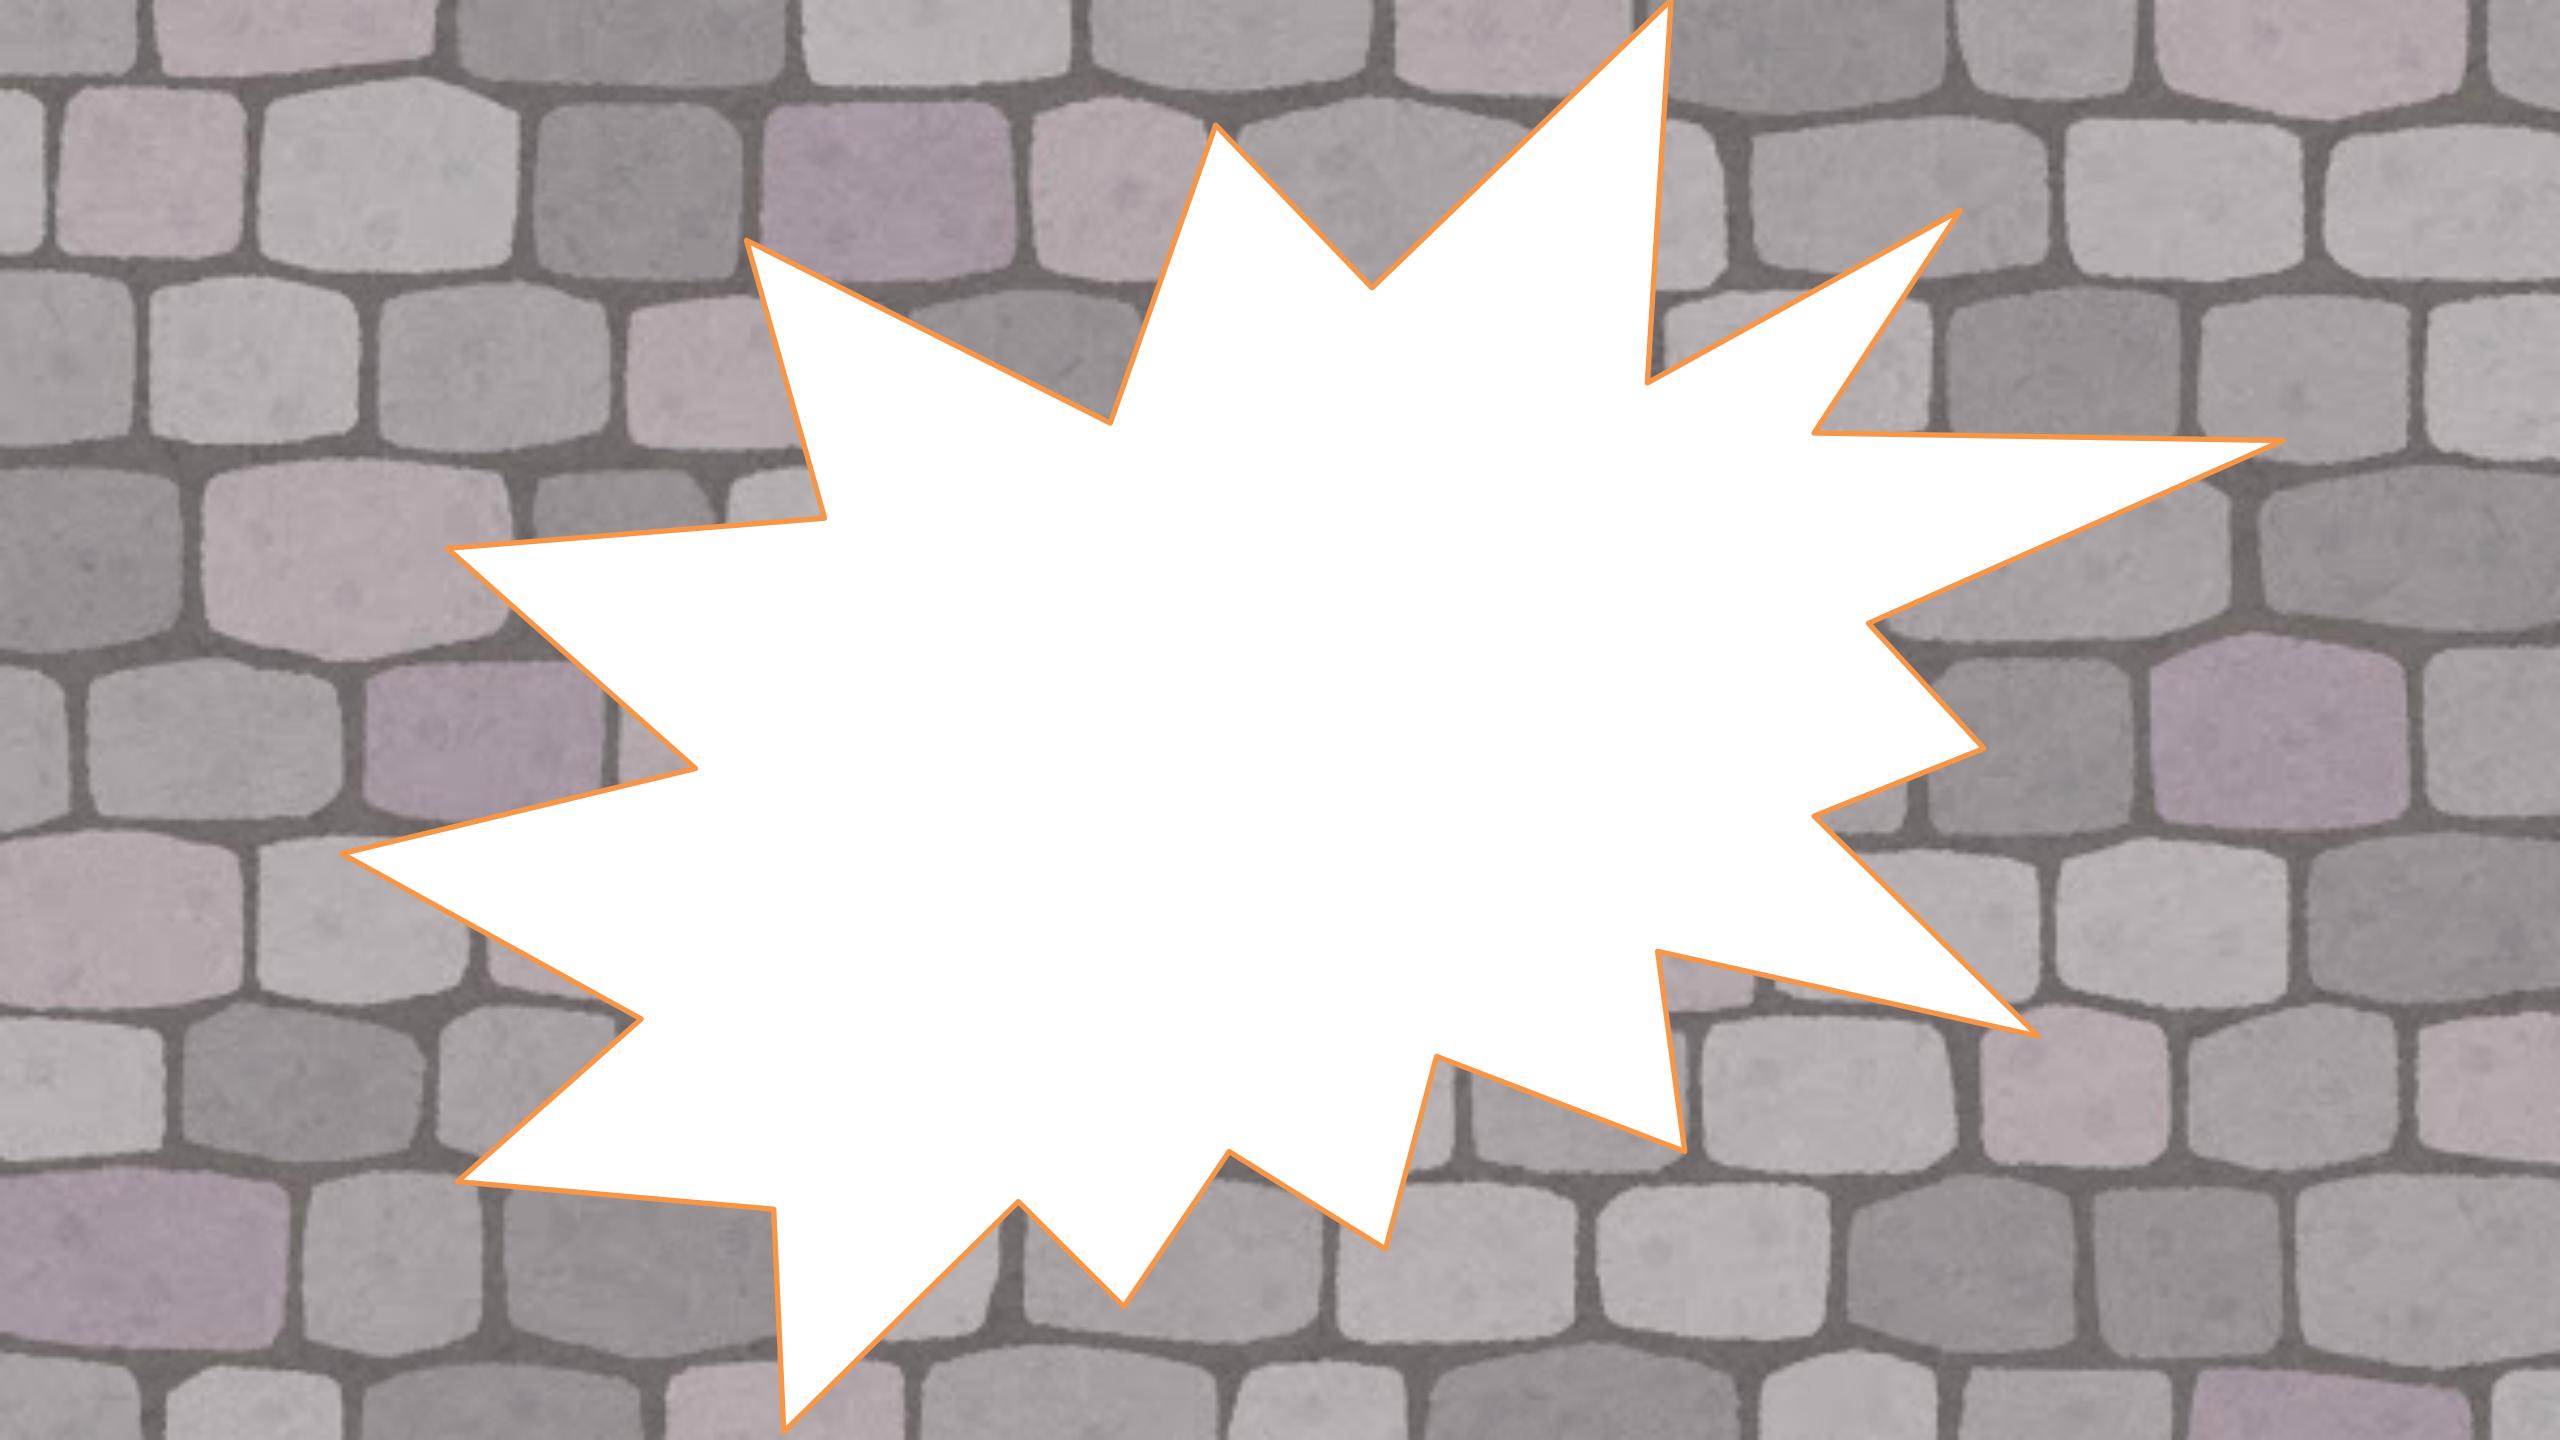


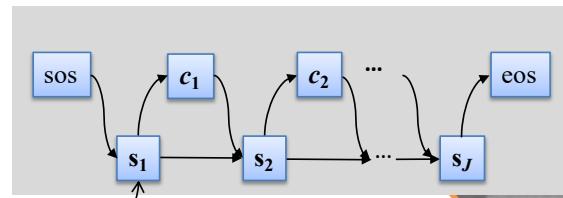
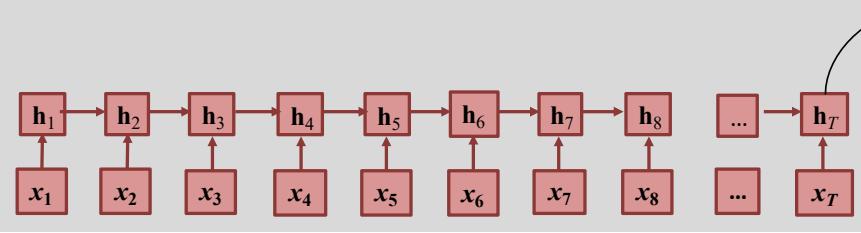
Big barrier:
noisy channel model
HMM
n-gram
etc.

However,

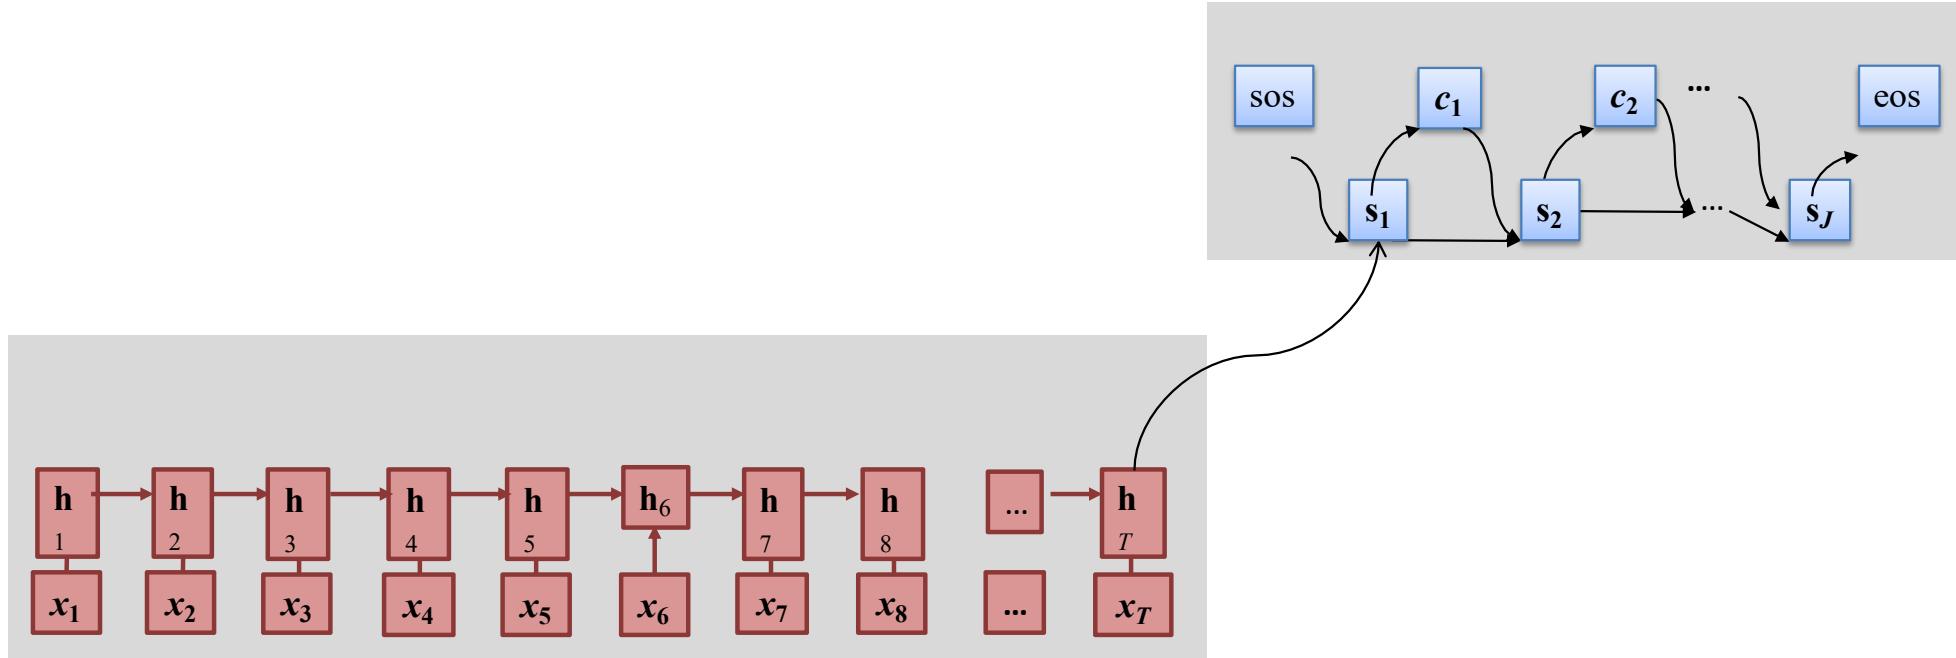






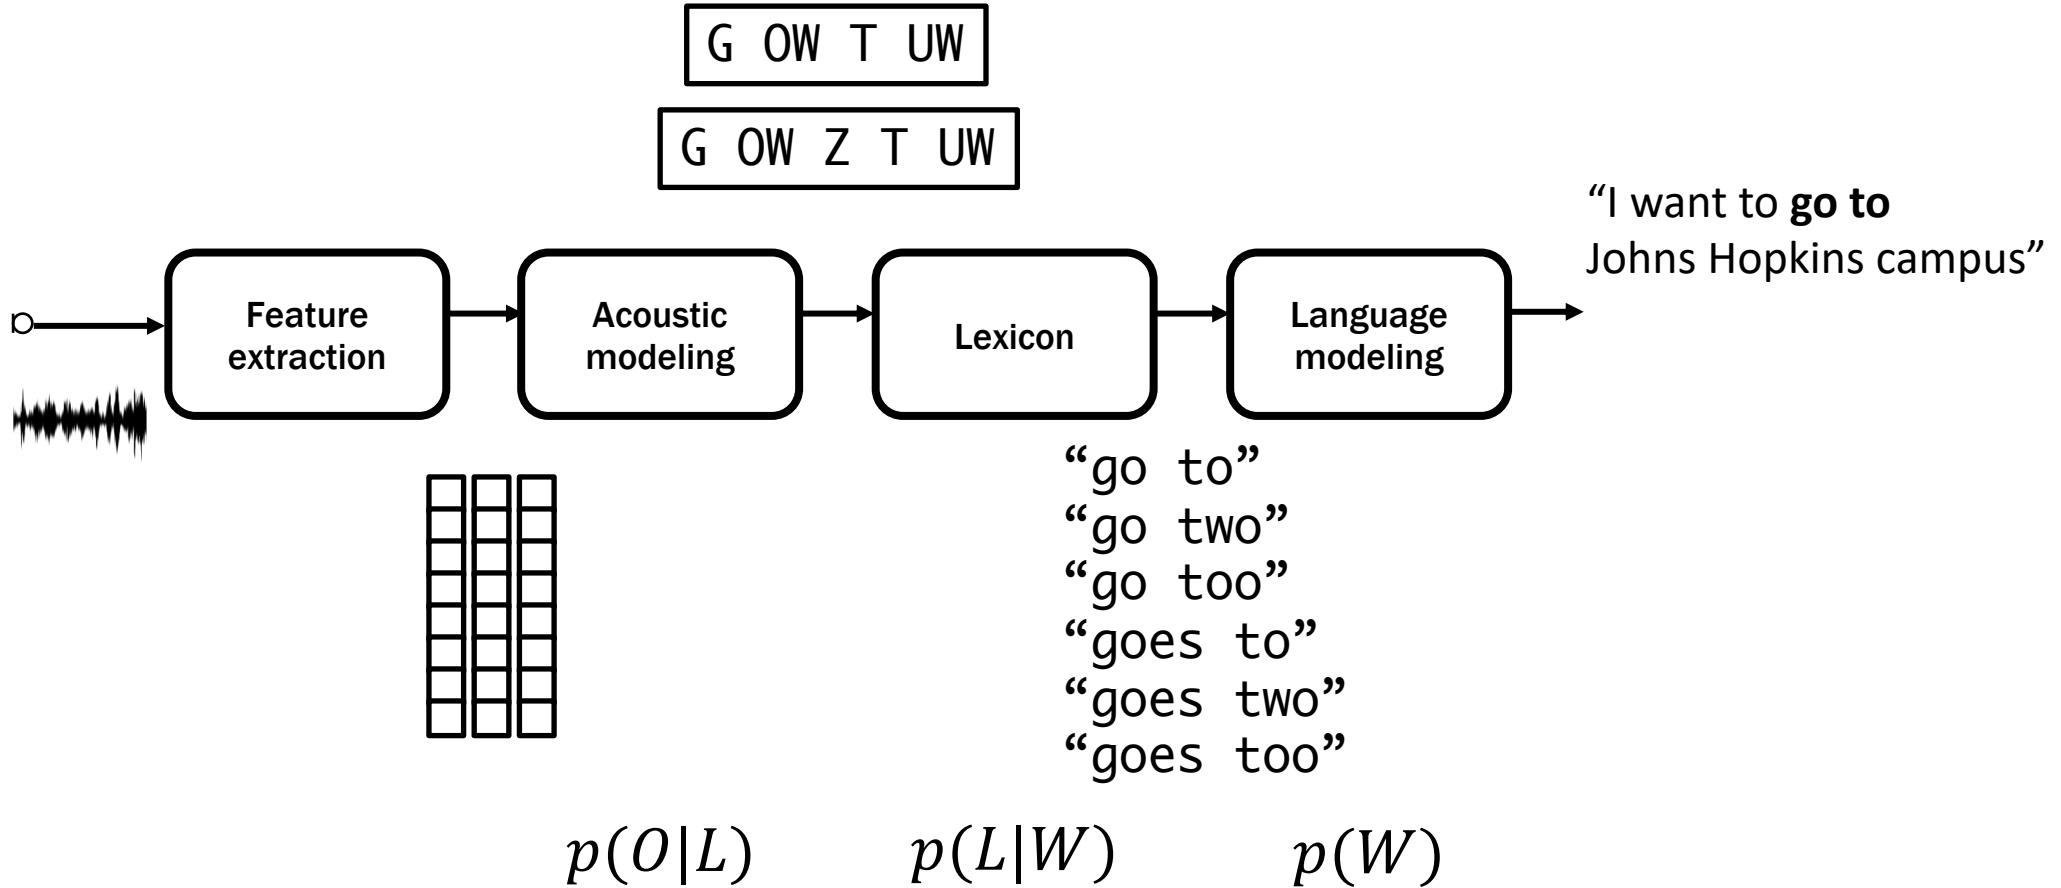


“End-to-End” Processing Using Sequence to Sequence

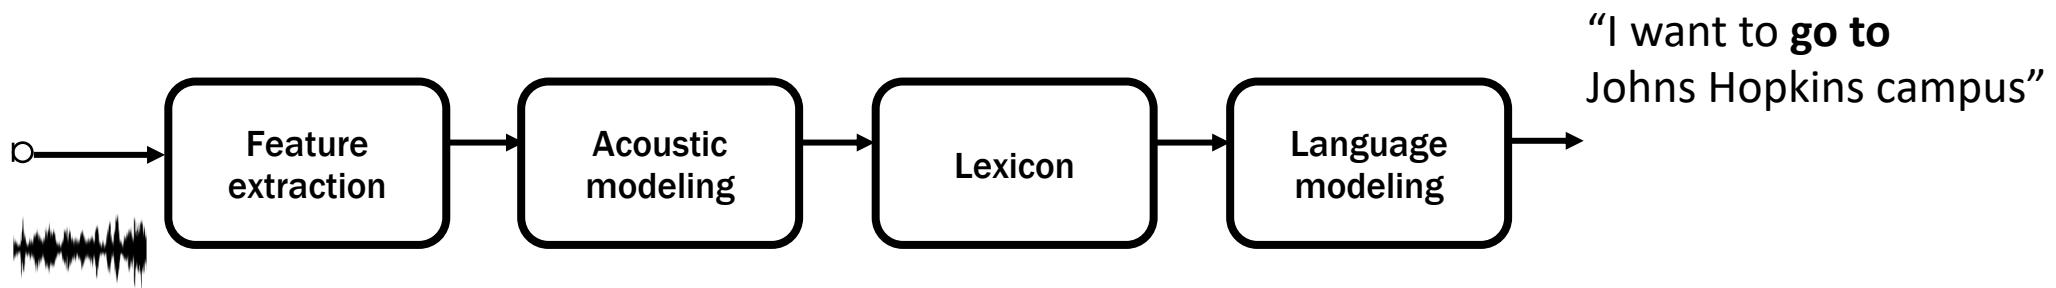


- Directly model $p(W|O)$ with a **single neural network**
- Great success in neural machine translation

Speech recognition pipeline

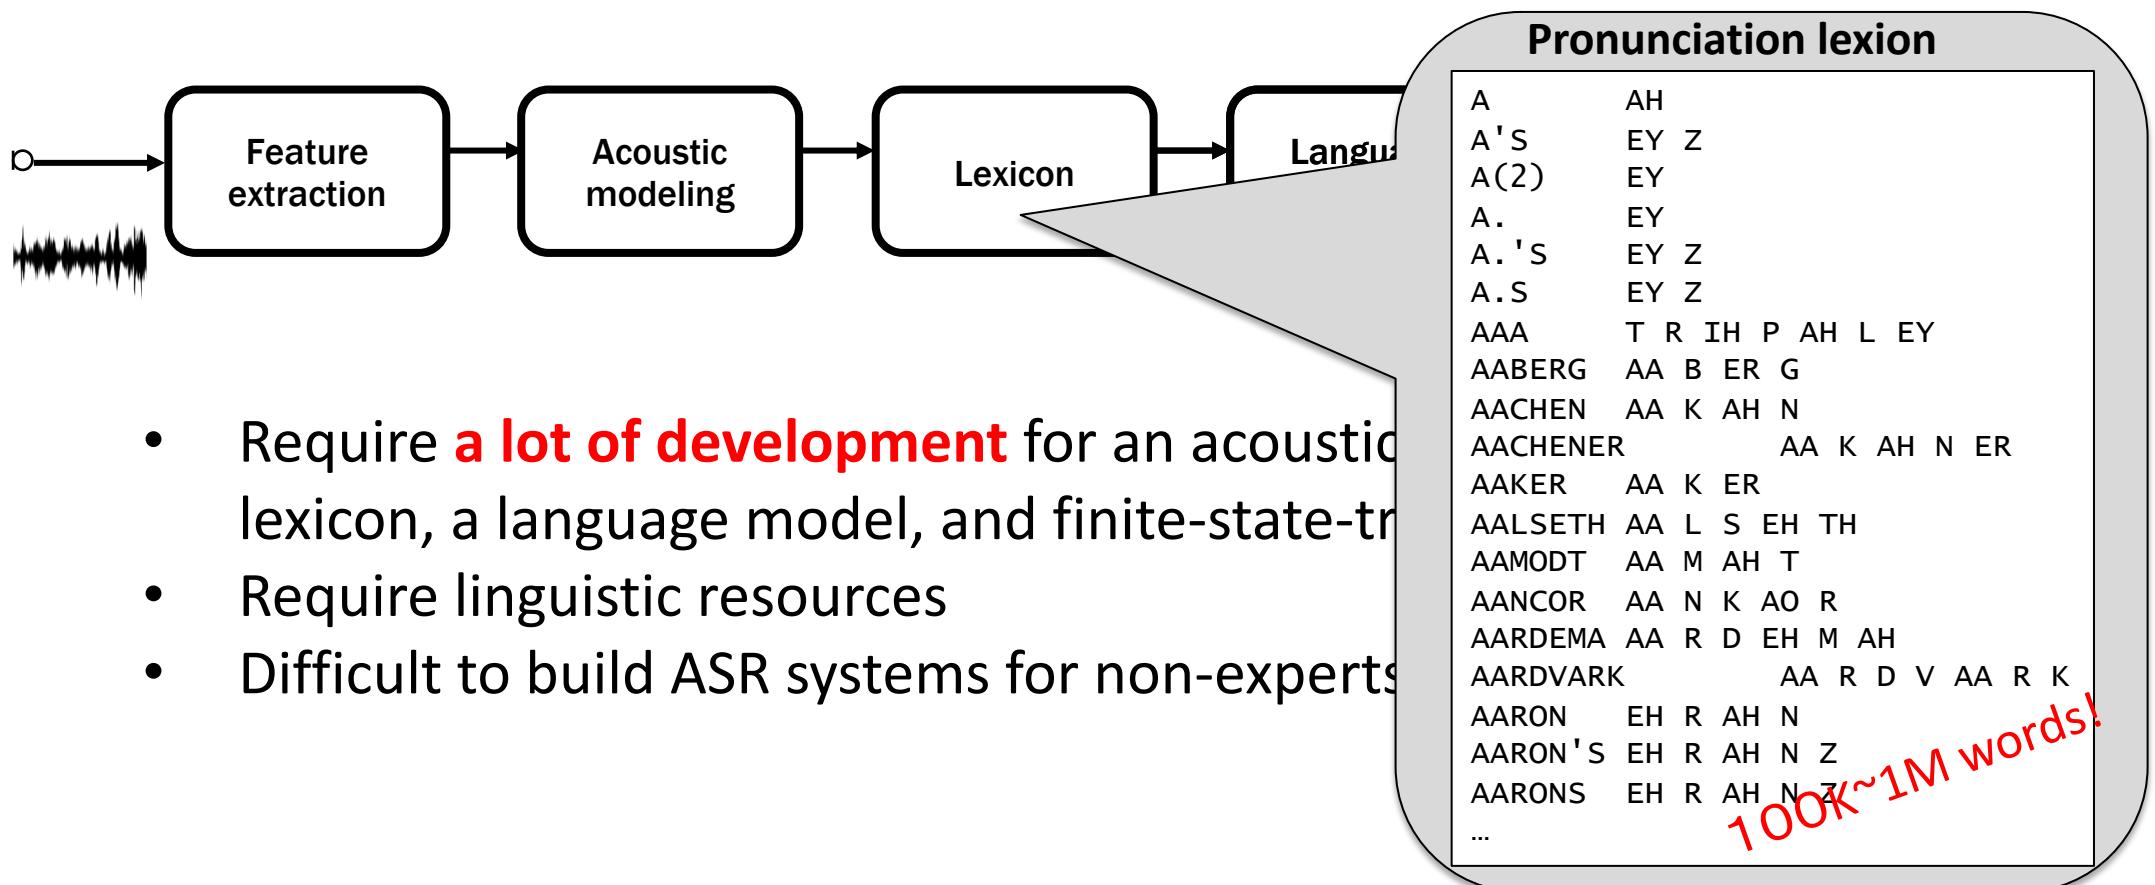


Speech recognition pipeline

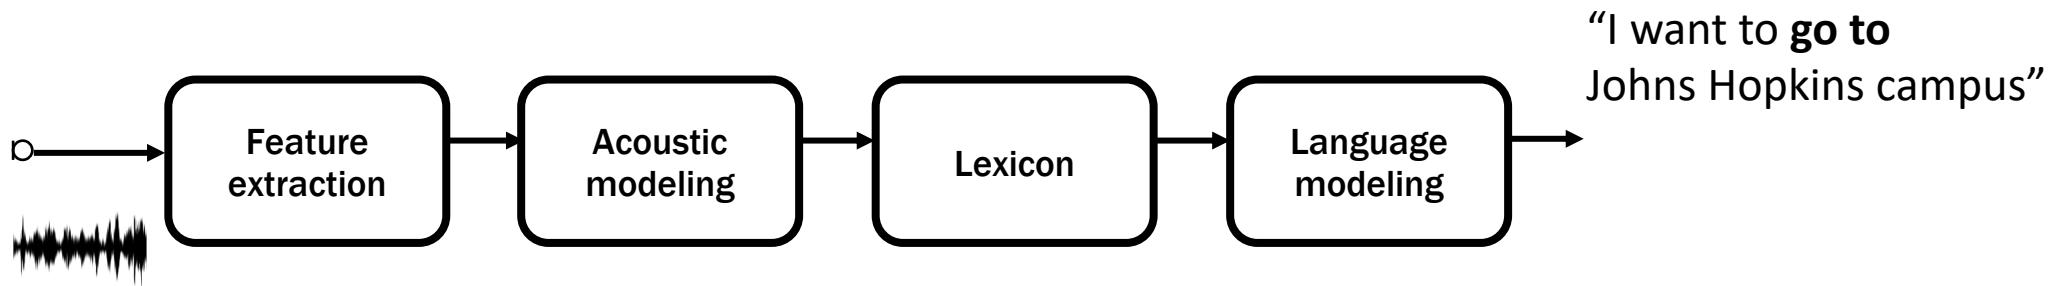


- Require **a lot of development** for an acoustic model, a pronunciation lexicon, a language model, and finite-state-transducer decoding
- Require linguistic resources
- Difficult to build ASR systems for non-experts

Speech recognition pipeline

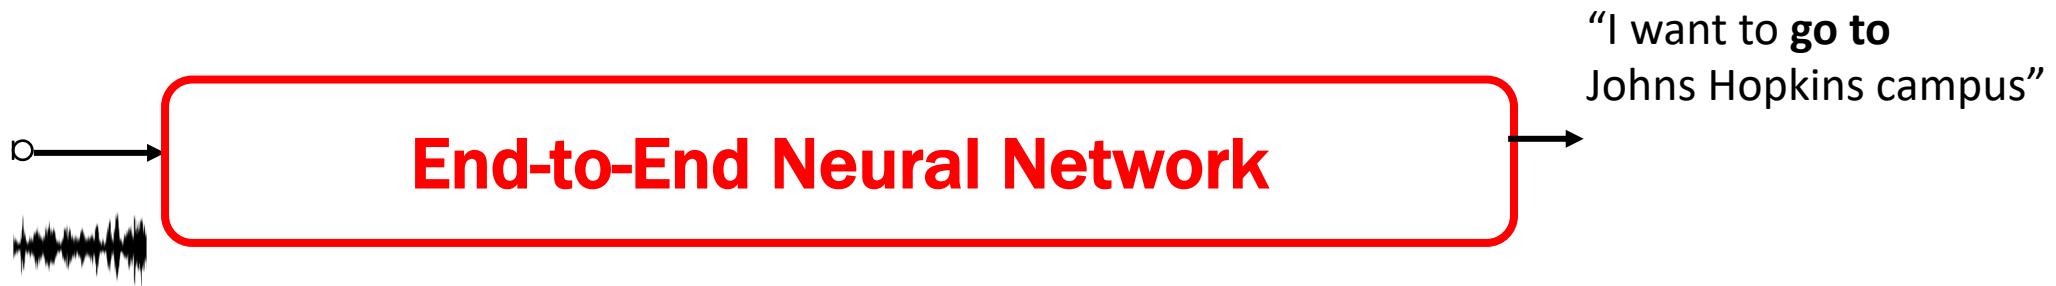


Speech recognition pipeline



- Require **a lot of development** for an acoustic model, a pronunciation lexicon, a language model, and finite-state-transducer decoding
- Require linguistic resources
- Difficult to build ASR systems for **non-experts**

From pipeline to integrated architecture

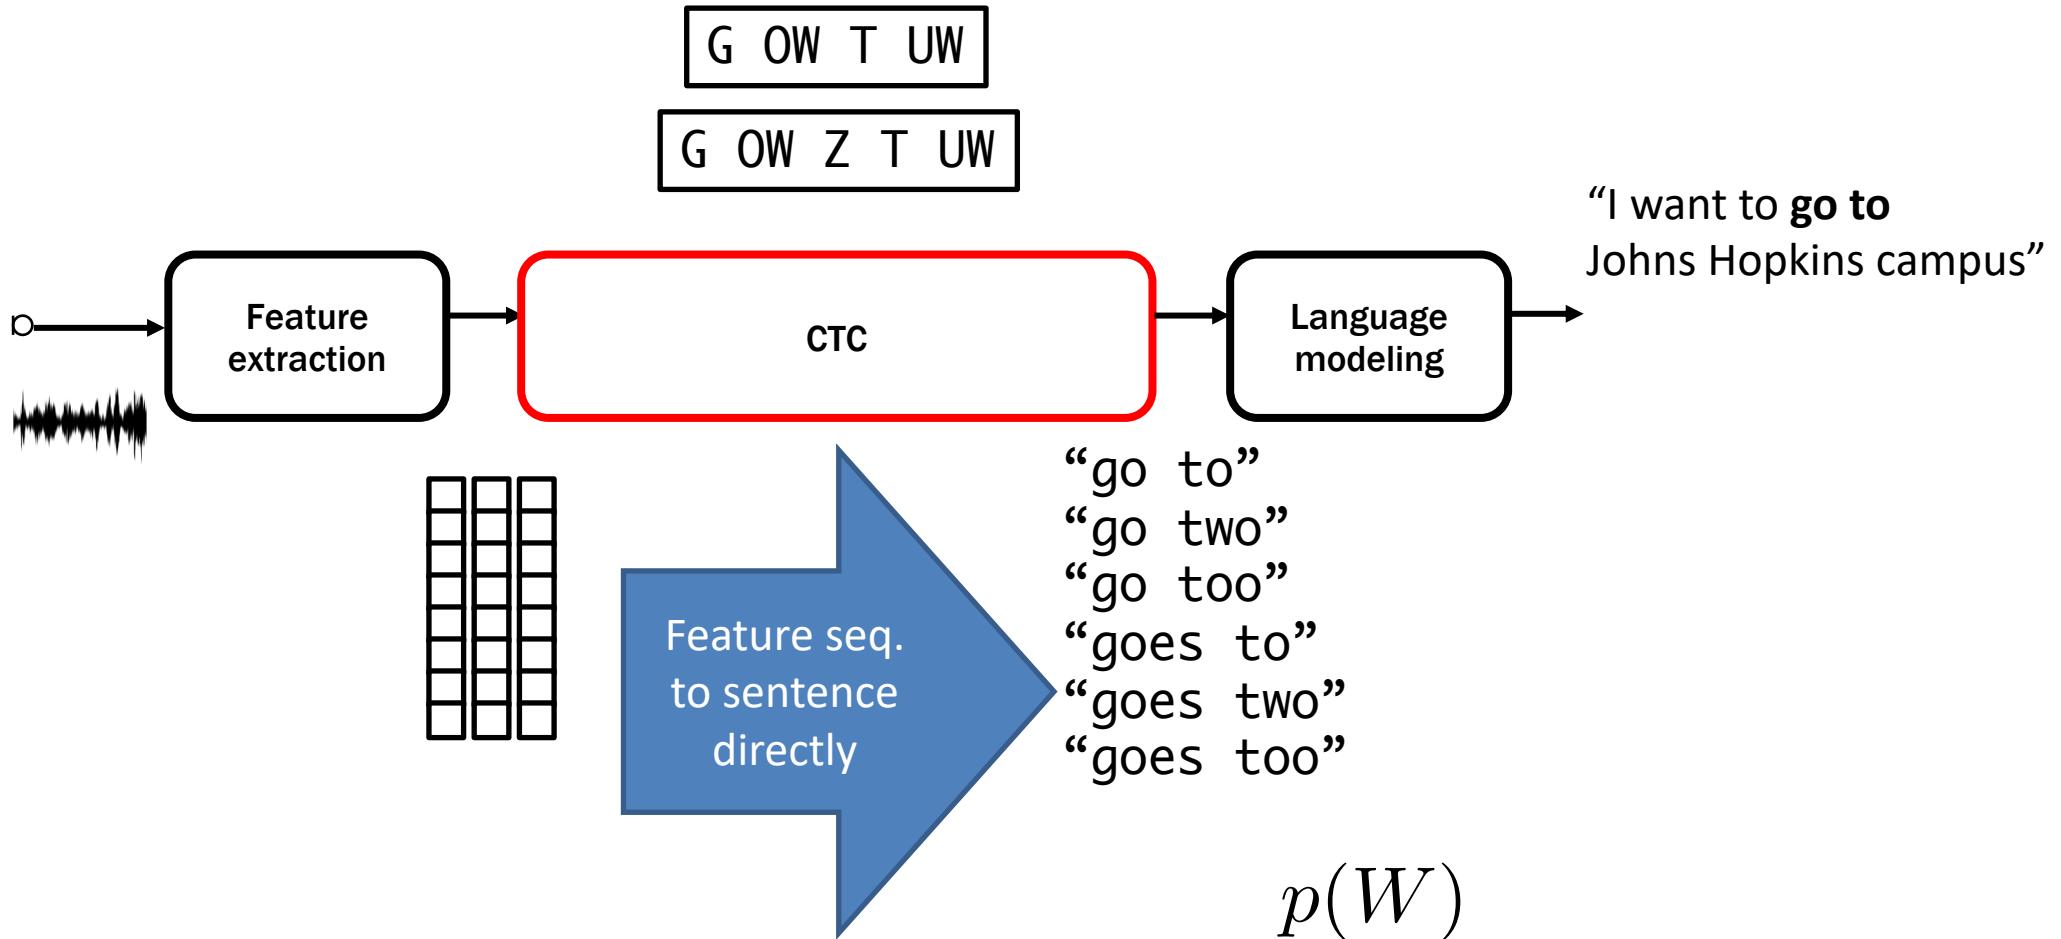


- Train a deep network that directly maps speech signal to the target letter/word sequence
- Greatly simplify the complicated model-building/decoding process
- Easy to build ASR systems for new tasks **without expert knowledge**
- Potential to outperform conventional ASR by **optimizing the entire network** with a single objective function

Today's agenda

- Introduction to end-to-end speech recognition
- Connectionist Temporal Classification (CTC)
- Attention

Speech recognition pipeline



Character seq. vs. word seq.

- Example: “I see”
 - $W = (w_i \in \{\text{“i”}, \text{“see”}, \dots\} | i = 1, 2)$
 - $C = (c_j \in \mathbb{U} | j = 1, \dots, 5)$, where $\mathbb{U} = \{\text{“a”}, \text{“b”}, \text{“c”}, \text{“d”}, \text{“e”}, \dots\}$ (Latin alphabet)
- Low/zero count problem
 - Word “bitcoin” is not appeared in old WSJ sentences, but character seq. can cover it
- Semantic context, lexicon constraint
 - Word unit can handle them, but not in the character unit
- No word unit in some languages
 - Some languages do not have word boundaries (no explicit word units)
- Remark: Subword/token (e.g., “_i”, “_s”, “ee”), data-driven ways to tokenize a character sequence to consider the benefits of both character and word units

Connectionist temporal classification

- Formulation
 - Let character seq. be $C = (c_t \in \mathbb{U} | j = 1, \dots, J)$ and feature seq. be $O = (\mathbf{o}_t \in \mathbb{R}^D | t = 1, \dots, T)$
 - Focus on the posterior distribution $p(C|O)$, and we can start from the Bayes decision theory:
$$\hat{C} = \operatorname{argmax}_C p(C|O)$$
 - This is the same start as the HMM except that we use the character seq. instead of the word seq.

Connectionist temporal classification

- Formulation
 - Focus on the posterior distribution $p(C|O)$, and rewrite it as

$$\begin{aligned} p(C|O) &= \sum_Z p(C|Z) p(Z|O) \\ &\approx \underbrace{\sum_Z p(C|Z)}_{\text{CTC LM}} \underbrace{p(Z|O)}_{\text{CTC AM}} \end{aligned}$$

- No Bayes theorem, but use conditional independence assumption
- Introduce latent variable seq. $Z = (z_t \in \{\mathbb{U}, \langle \mathbf{b} \rangle\} | t = 1, \dots, T)$ that has **the same length** as input feature seq.
 - We can use a conventional RNN to model this $p(Z|O)$
- Similarly to HMM, we'll consider the summation of all possible Z

Introduction of blank symbol

- First, we insert to the character seq.

"see"

$\rightarrow C = ("s", "e", "e")$, where $|C| = J$

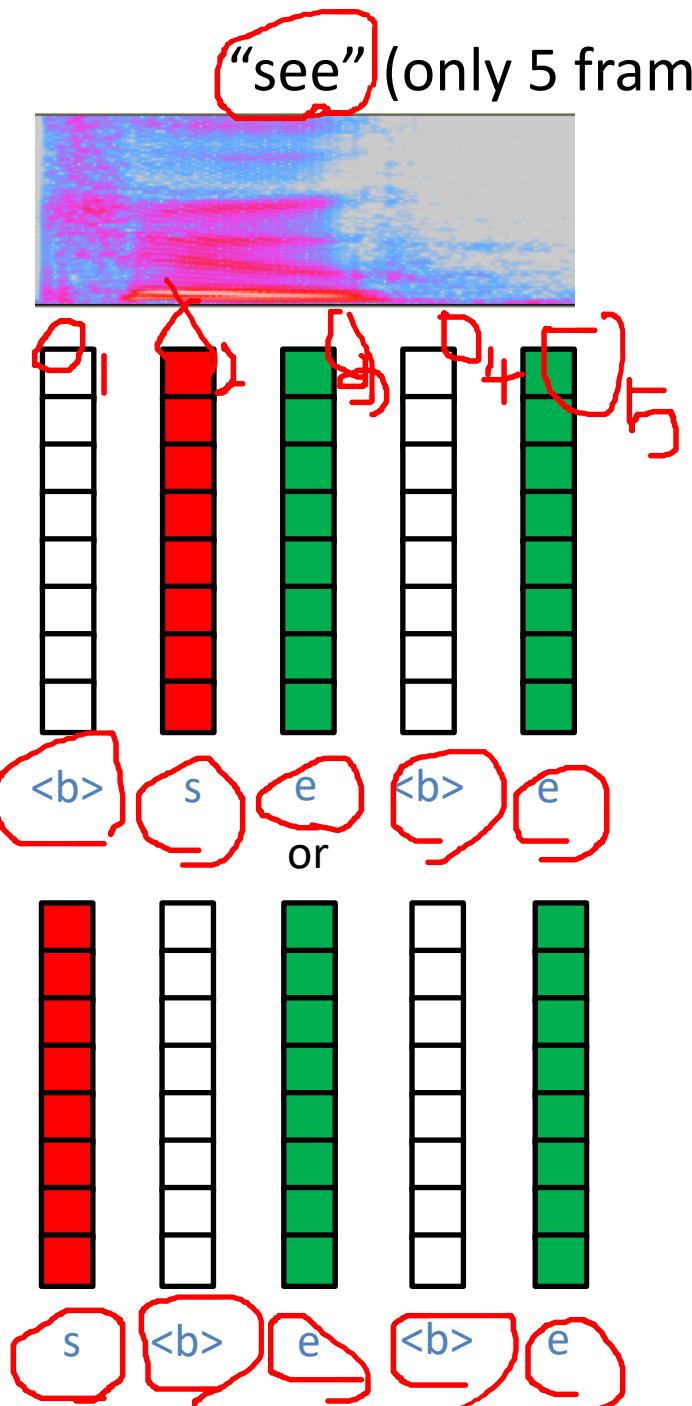
$\rightarrow C' = (, "s", , "e", , "e",)$, where $|C'| = 2J + 1$

- Then, expand C' to the frame length T to form Z
 - All characters can be repeated
 - can be skipped except when it is inserted between repeated character
 - "s", "", "e": we can skip
 - "e", "", "e": we **cannot** skip
 - See A. Graves et al. "Connectionist temporal classification: labelling unsegmented sequence data with recurrent neural networks." in ICML

“see” (only 5 frames)

Example of Z

- $C = ("s", "e", "e")$
- $C' = ("", "s", "", "e", "", "e", "")$
- $T = 5$
- $Z = ("", "s", "e", "", "e"), ("s", "", "e", "", "e"), ("s", "s", "e", "", "e"), \dots$
- This is an alignment problem
- $C \rightarrow Z$: one to many mapping



CTC Formulation

- CTC acoustic model

$$p(Z|O) = \prod_{t=1}^T p(z_t | z_1, \dots, \cancel{z_{t-1}}, O)$$

$$\approx \prod_{t=1}^T \underline{p(z_t | O)}.$$

- Using conditional independence assumption to factorize the posterior $p(Z|O)$ but this is not bad assumption compared with HMM
- This can be realized by Bidirectional LSTM or self attention

$$p(z_t = j | O) = [\text{softmax}(\mathbf{W}\mathbf{h}_t + \mathbf{b})]_j,$$

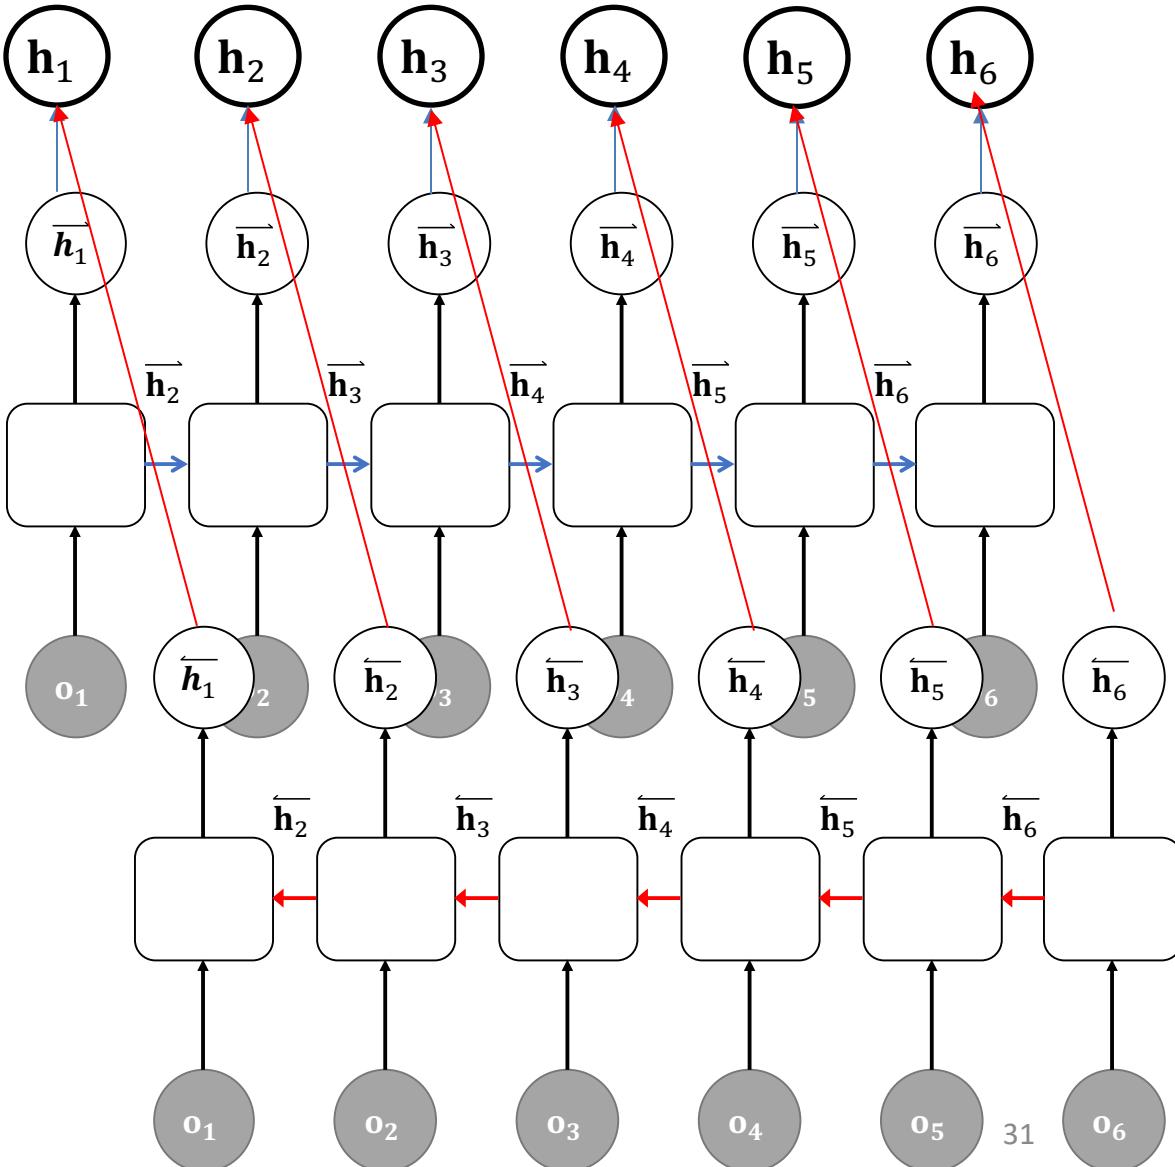
$$\mathbf{h}_t = \text{BLSTM}(O) \text{ for } t = 1, \dots, T.$$

Bidirectional RNN

- $\mathbf{h}_t = \begin{bmatrix} \overrightarrow{\mathbf{h}}_t \\ \overleftarrow{\mathbf{h}}_t \end{bmatrix} = f(O = (\mathbf{o}_1, \dots, \mathbf{o}_T))$

Then,

- $p(z_t | O)$
 $\approx p(z_t | \mathbf{h}_t)$



CTC Formulation

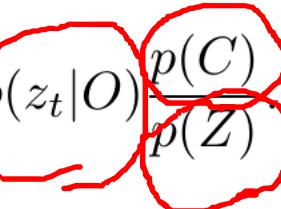
- CTC Language model (generative model view)

$$\begin{aligned} p(C|Z) &= \frac{p(Z|C)p(C)}{p(Z)} \\ &= \prod_{t=1}^T p(z_t|z_1, \dots, z_{t-1}, C) \frac{p(C)}{p(Z)} \\ &\approx \underbrace{\prod_{t=1}^T p(z_t|z_{t-1}, C)}_{\text{red underline}} \frac{p(C)}{p(Z)}, \end{aligned}$$

- Using conditional independence assumption (1st order Markov) to factorize the posterior, same as the HMM
- $p(C)$: Letter language model (we can also combine the word language model)
- $p(Z)$: Prior probability for the state sequence

Summary of CTC formulation

- $p(C|O)$ is rewritten as follows

$$p(C|O) \approx \sum_Z \prod_{t=1}^T p(z_t|z_{t-1}, C) p(z_t|O) \frac{p(C)}{p(Z)}$$


- In general, prior probabilities $p(C)$ and $p(Z)$ are separately obtained (not fully end-to-end)
- We can further eliminate the prior probabilities by assuming the uniform distributions as follows ($\mathcal{Z}(C)$ denotes all possible CTC paths given C):

$$p(C|O) \approx \underbrace{\sum_{Z \in \mathcal{Z}(C)} \prod_{t=1}^T p(z_t|O)}_{\triangleq p_{ctc}(C|O)}$$

- Basically, we can use a **forward-backward algorithm** to estimate the parameter

HMM/DNN vs. CTC

- Conditional independence assumptions
- Language models
- Use of pronunciation lexicon information
- Implementation

Let's discuss the difference

CTC vs. HMM

HMM

$$p(W|O)$$

- With Bayes rule and CIA (separate acoustic, lexicon, and language models)

$$\sum_{Z,L} p(O, Z|L) p(L|W) p(W)$$

- 1st order Markov and frame-level decomposition

$$p(O|Z)p(Z|L) \rightarrow \prod_t p(o_t|z_t)p(z_t|z_{t-1}, L)$$

- Replace the likelihood function $p(o_t|z_t)$ with a DNN based on the pseudo likelihood trick

CTC

$$p(C|O)$$

- No Bayes rule, but CIA (separate acoustic and language model)

$$\sum_Z p(C|Z)p(Z|O)$$

- 1st order Markov and frame-level decomposition
 $p(Z|O)p(C|Z) \rightarrow \prod_t p(z_t|O) p(z_t|z_{t-1}, C)p(C)$
- Replace the frame-level posterior distribution $p(z_t|O)$ with a DNN

Basic assumptions are very similar (CIA, 1st order Markov assumptions)

Implementation of CTC

- During training
 - Major toolkit supports CTC
 - Tensorflow, Pytorch, Chainer, etc.
 - Nvidia cuDNN also supports CTC
- During recognition
 - You have to implement the following search:
 $\text{argmax}_C p(C|Z)p(Z|O)$
 - This can be efficiently performed by using a finite state transducer

The screenshot shows the PyTorch documentation for the `CTCLoss` class. The page includes a sidebar with links to various PyTorch modules and a search bar. The main content area has the following sections:

- CLASS `torch.nn.CTCLoss(blank=0, reduction='mean')`** [SOURCE] ⚡
- CTCLoss**
- Parameters:**
 - `blank (int, optional)` – blank label. Default 0.
 - `reduction (string, optional)` – Specifies the reduction to apply to the output: 'none' | 'mean' | 'sum'. 'none': no reduction will be applied, 'mean': the output losses will be divided by the target lengths and then the mean over the batch is taken. Default: 'mean'
- Inputs:**
 - `log_probs: Tensor of size (T, N, C) where C=number of characters in alphabet including blank,`
T = input length, and N = batch size. The logarithmized probabilities of the outputs (e.g. obtained with `torch.nn.functional.log_softmax()`).
 - `targets: Tensor of size (N, S) or (sum(target_lengths)).`
Targets (cannot be blank). In the second form, the targets are assumed to be concatenated.
 - `input_lengths: Tuple or tensor of size (N).`
Lengths of the inputs (must each be $\leq T$)
 - `target_lengths: Tuple or tensor of size (N).`
Lengths of the targets
- Example:**

```
>>> ctc_loss = nn.CTCLoss()
>>> log_probs = torch.randn(50, 16, 20).log_softmax(2).detach().requires_grad_()
>>> targets = torch.randint(1, 20, (16, 30), dtype=torch.long)
>>> input_lengths = torch.full((16,), 50, dtype=torch.long)
>>> target_lengths = torch.randint(10, 30, (16,), dtype=torch.long)
```
- Shortcuts** on the right side, listing various PyTorch modules and functions.

Baidu CTC

[Amodei+(2015)]

- Optimization of computational cost of CTC dynamic programming
- Multiple GPUs
- Architecture optimization (BLSTM -> GRU, use of CNN)
- Use 12,000 hours of data for training
- Data augmentation (noise)

Test set	Read Speech		
	DS1	DS2	Human
WSJ eval'92	4.94	3.60	5.03
WSJ eval'93	6.94	4.98	8.08
LibriSpeech test-clean	7.89	5.33	5.83
LibriSpeech test-other	21.74	13.25	12.69

Google CTC

[Soltan+ (2016)]

- Word-level CTC, conventional BLSTM
- No language model
- 125,000 hours of training data (!) from YouTube

Model	HMM	Layers	Outputs	Params	Vocab	OOV(%)	Spoken WER(%) w/ LM	w/o LM
CTC CD phone		7x1000	6400	43m	500000	0.24	12.3	—
CTC spoken words		7x1000	82473	116m	82473	0.63	11.6	12.0

- Word-level CTC obtains comparable performance (even without LM)

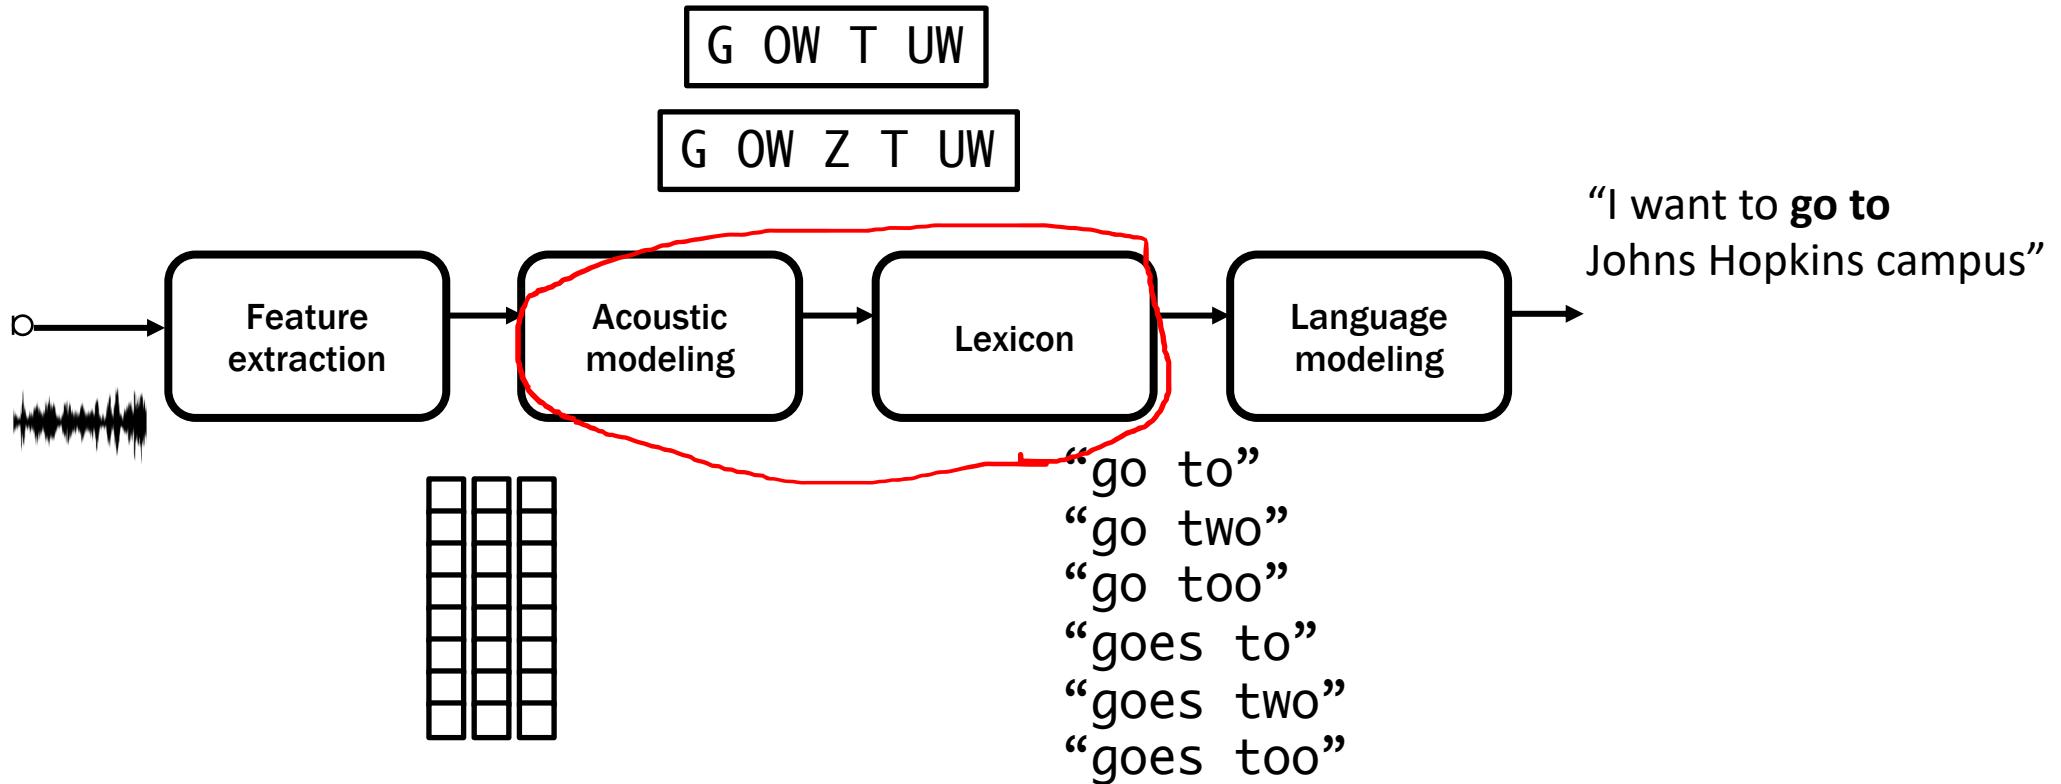
Summary

- CTC
 - One promising direction of end-to-end
 - No language model (but it can be combined with an LM)
 - Still based on conditional independence assumptions and Markov model
 - CTC is really end-to-end?
- Can we use it to any of sequence to sequence task?
 - The alignment should be monotonic (HMM like task, it cannot be applied to machine translation)
 - The input length must be longer than the output length (it cannot be applied to speech synthesis)
- Attention
 - Another end-to-end

Today's agenda

- Introduction to end-to-end speech recognition
- Connectionist Temporal Classification (CTC)
- Attention

Speech recognition pipeline

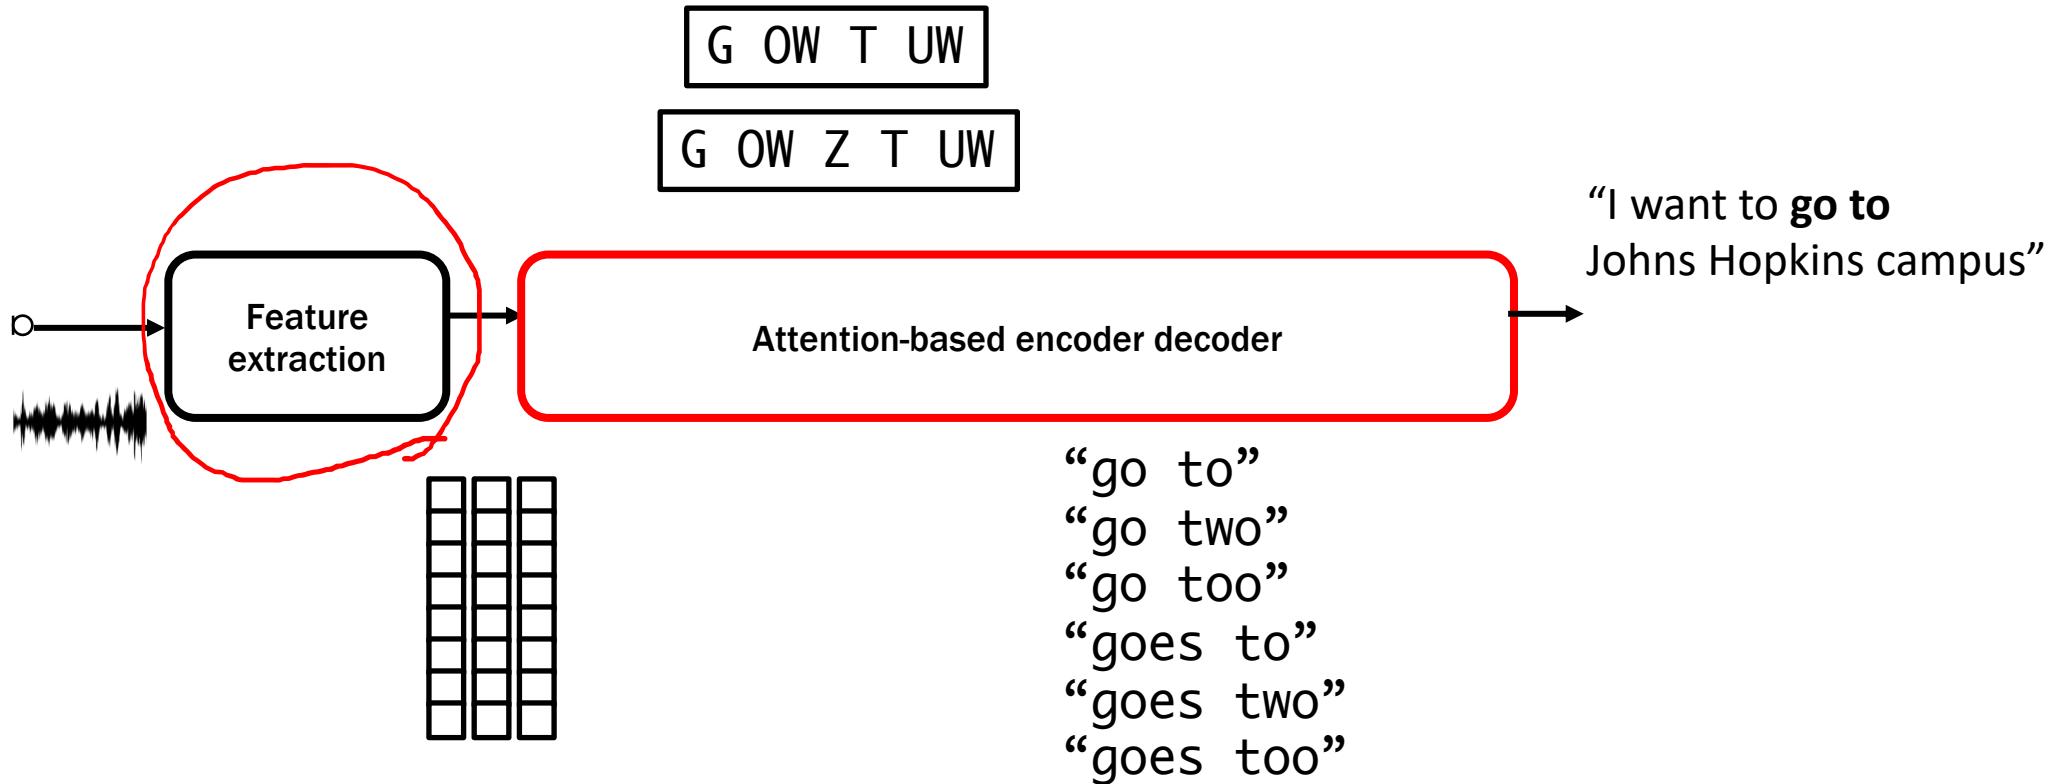


$$p(O|L)$$

$$p(L|W)$$

$$p(W)$$

Speech recognition pipeline



$$p(O|L)$$

$$p(L|W)$$

$$p(W)$$

Attention based encoder-decoder

- Let $C = (c_j \in \mathbb{U} | j = 1, \dots, J)$, be a character sequence
 - \mathbb{U} : set of characters
- Let $O = (\mathbf{o}_t \in \mathbb{R}^D | t = 1, \dots, T)$, be a sequence of D dimensional feature vectors
$$\hat{C} = \operatorname{argmax}_C p(C|O)$$
- Problem: T and J are different, and we cannot use normal neural networks
- Sequence to sequence is a solution to deal with it

Problem of original encoder-decoder architecture

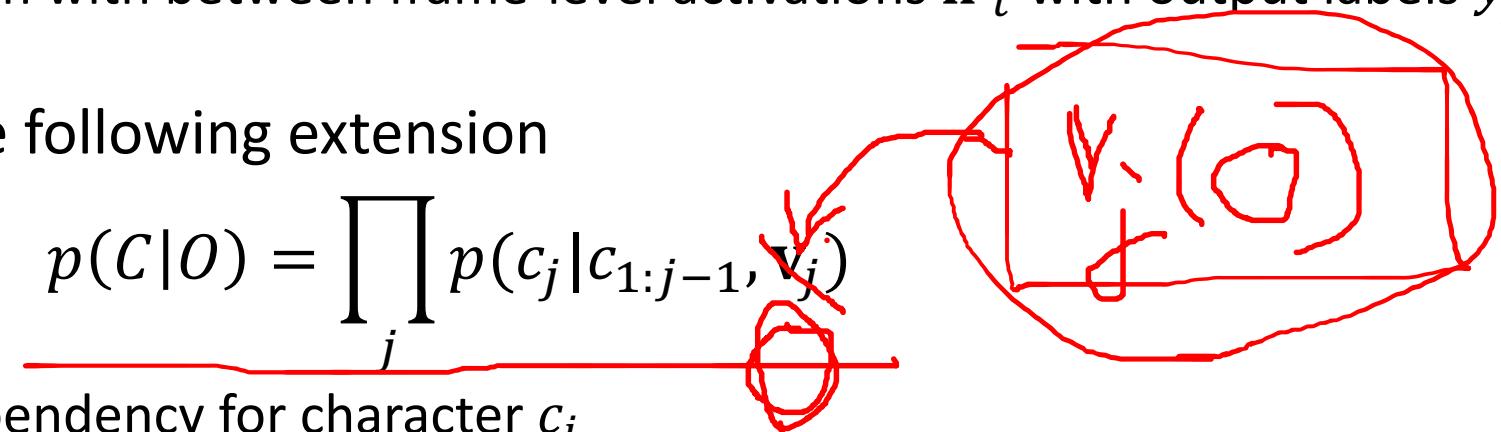
$$p(O|C) = \prod_j p(c_j|c_{1:j-1}, \mathbf{h}'_T)$$

- We cannot explicitly connect the relationship between input and output (an alignment property)
 - No explicit connection with between frame-level activations \mathbf{h}'_t with output labels y_j

Instead, we consider the following extension

$$p(O|C) = \prod_j p(c_j|c_{1:j-1}, \mathbf{v}_j)$$

- \mathbf{v}_j has an explicit dependency for character c_j



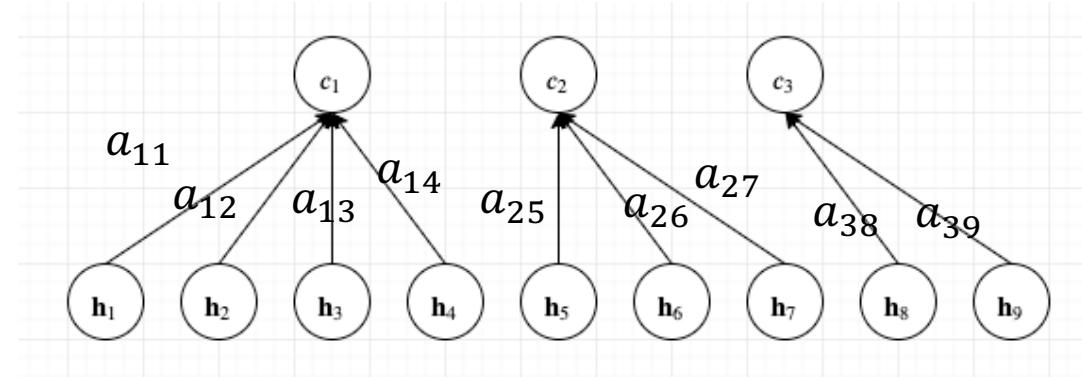
Attention mechanism

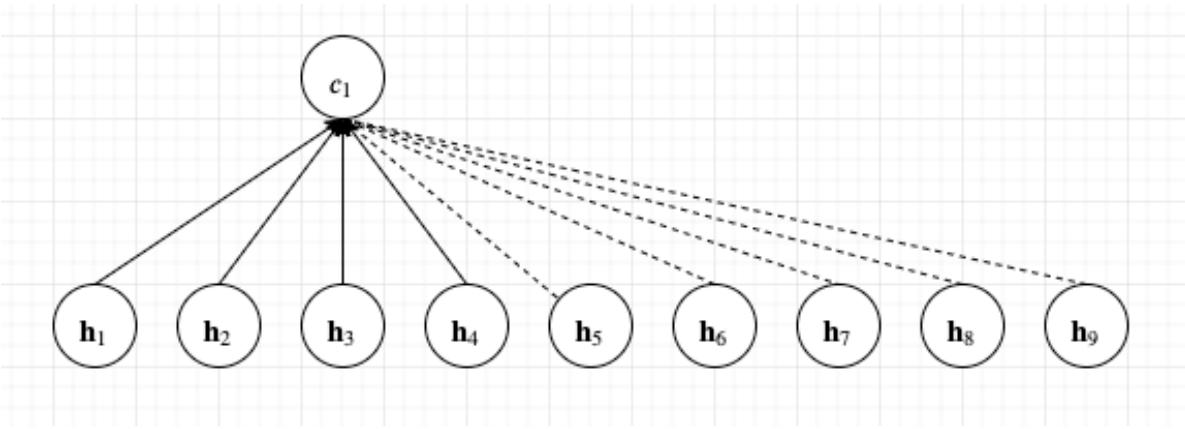
$$p(C|O) = \prod_j p(c_j | c_{1:j-1}, \mathbf{v}_j)$$

- Obtain the context vector

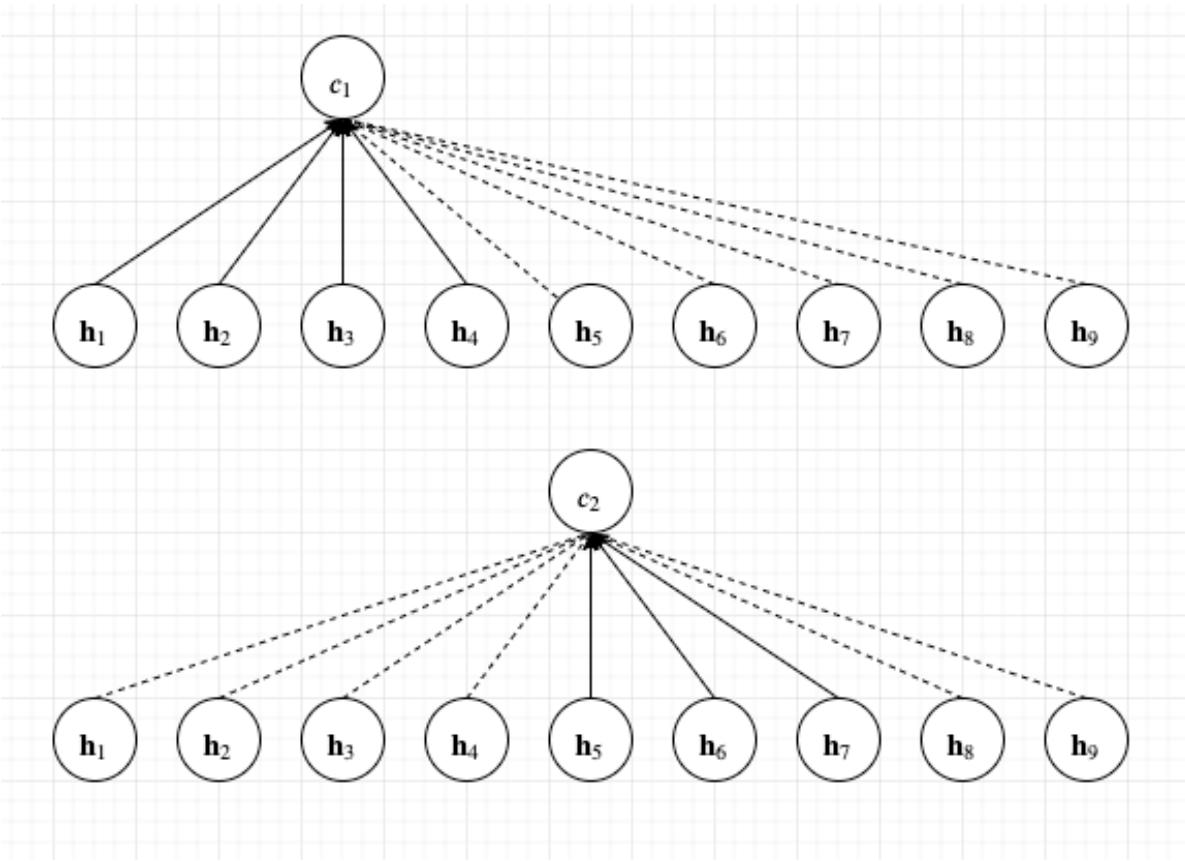
$$\mathbf{v}_j = \sum_{t=1}^T a_{jt} \mathbf{h}'_t$$

- Compute the assignment probability for each output j from a neural network
- $\mathbf{a}_j = \{a_{jt} | t = 1, \dots, T\} \in \mathbb{R}^T, 0 < a_{jt} < 1, \sum_{t=1}^T a_{jt} = 1$
- a_{jt} is obtained by using a neural network

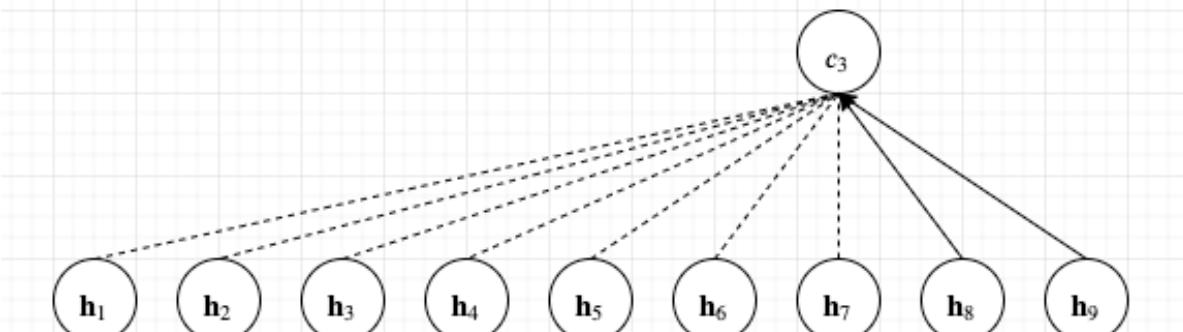
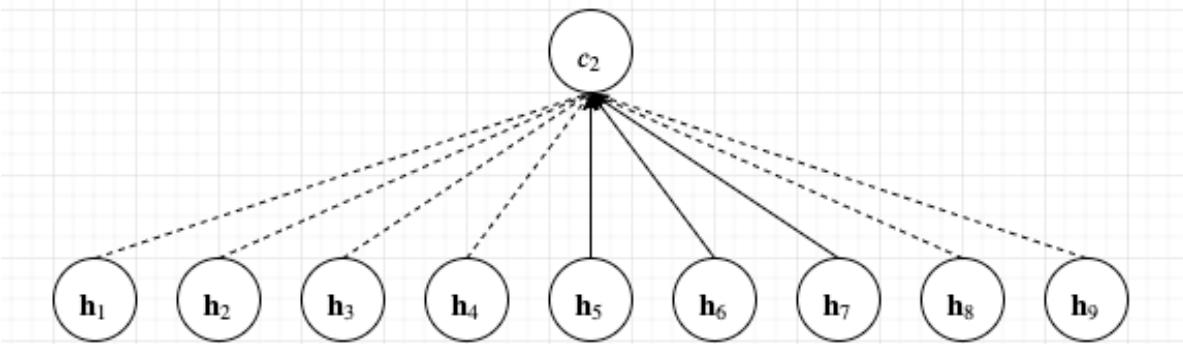
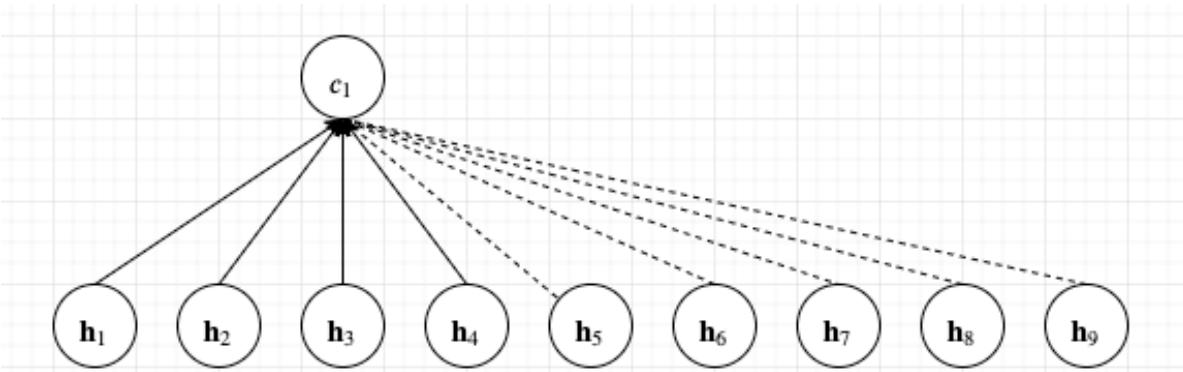




Normal arrow:
high probability
Dashed arrow:
low probability

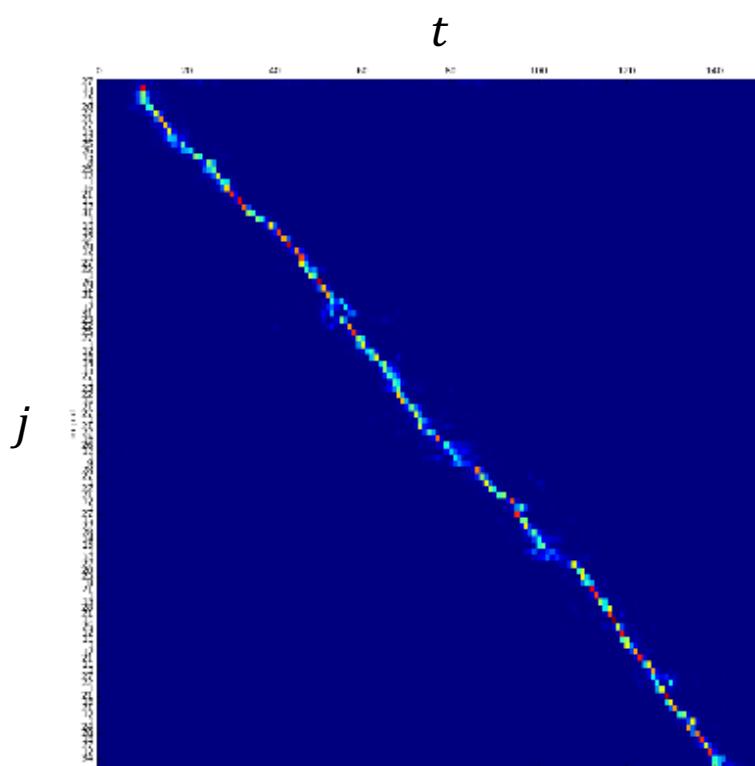


Normal arrow:
high probability
Dashed arrow:
low probability



Normal arrow:
high probability
Dashed arrow:
low probability

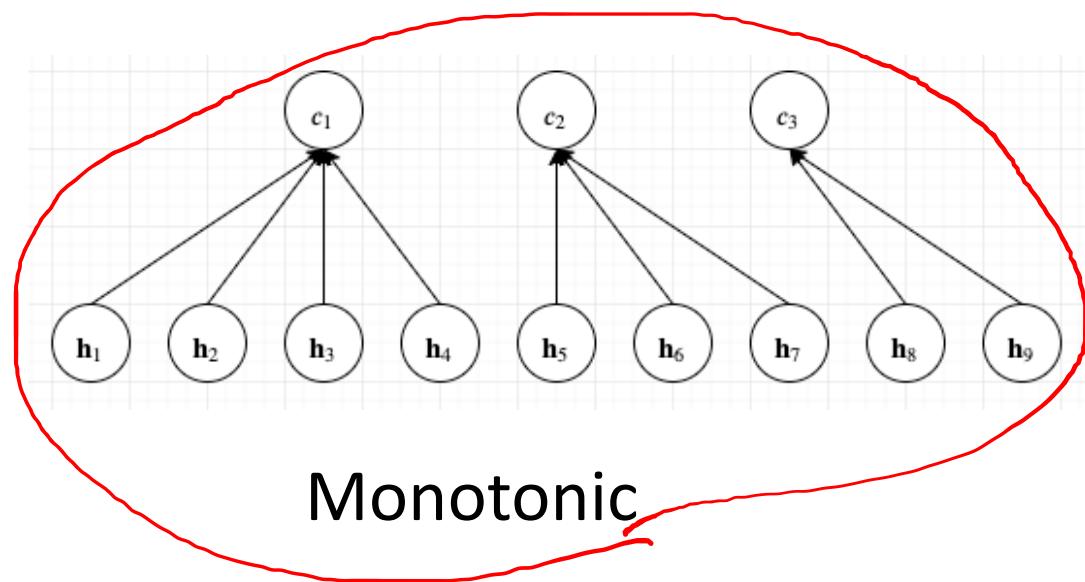
The attention mechanism performs a soft alignment



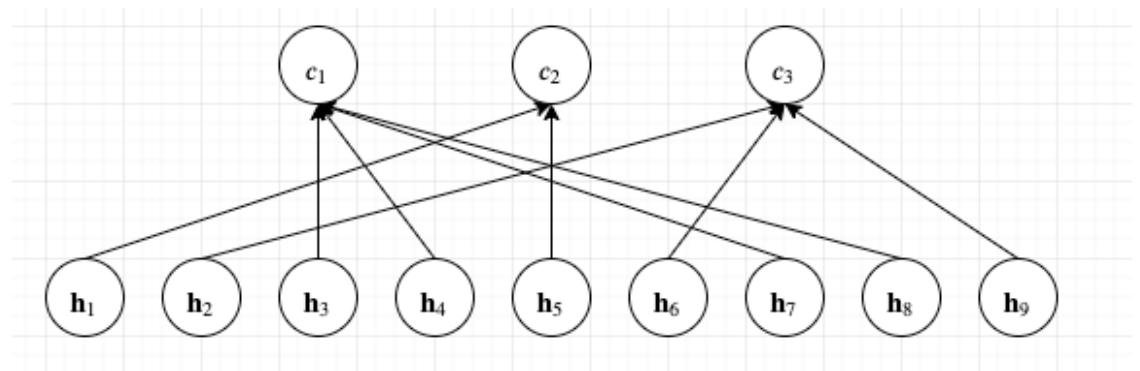
- $\mathbf{v}_j = \sum_{t=1}^T a_{jt} \mathbf{h}'_t$
- Attention weight a_{jt} determines whether encoder \mathbf{h}'_t is assigned to a character c_j or not
 - $a_{jt} \approx 0$: no assignment
 - $a_{jt} \neq 0$: assigned

The attention mechanism performs a soft alignment

- There is no constraint for the alignment
- The order can be changed (good for machine translation, but it does not happen in speech recognition)

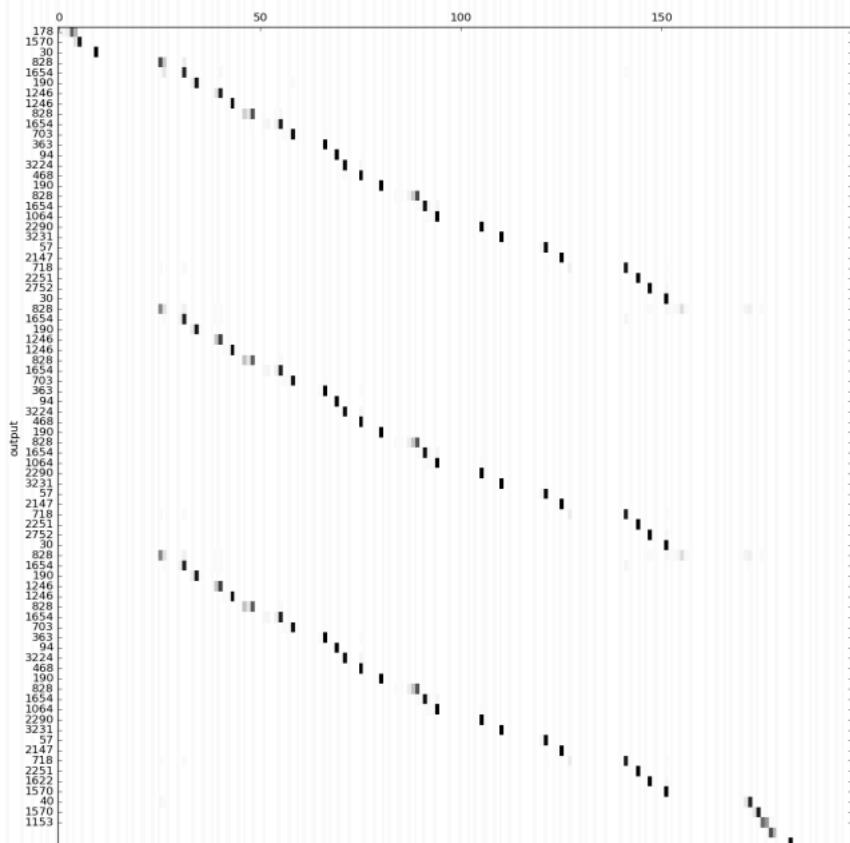


Monotonic



Non monotonic

Examples of wrong alignments



id: (20040717_152947_A010409_B010408-A-057045-057837)

Reference

但是如果你想如果回到了过去你如果带着这个现在的记忆是不是很痛苦啊

MTL

Scores: (#Correctness #Substitution #Deletion #Insertion) 28 2 3 45

但是如果你想如果回到了过去你如果带着这个现在的节
如果你想如果回到了过去你如果带着这个现在的节如果
你想如果回到了过去你如果带着这个现在的机是不是很

•

Difference in training and recognition

- **During training**
 - $\hat{\theta} = \operatorname{argmax}_{\theta} \prod_j p(c_j | c_{1:j-1}, \mathbf{v}_j)$
 - We use the transcriptions for $c_{1:j-1}$
- **During recognition**
 - $\hat{C} = \operatorname{argmax}_c \prod_j p(c_j | c_{1:j-1}, \mathbf{v}_j)$
 - However, we don't know the correct transcription $c_{1:j-1}$ during recognition
 - We use estimated history $\hat{c}_{1:j-1}$ instead of the correct transcription
→ mismatch of training and recognition
 - Recognition result is stopped when we observe the "eos" symbol
 - argmax_C is impossible → Approximately only consider possible high-score hypotheses (beam search)

Summary of attention encoder-decoder

- No conditional independence assumption
 - No need for pronunciation lexicon
 - Attention & Encoder: acoustic model
 - Decoder: language model
 - Combine acoustic and language models with single network
- Attention model is too flexible for alignment issues
- Not easy to combine the language model trained with a bunch of text data

Google's Experiments (12,500 hours)

- They use huge amount of training data (pair data)
 - 12,500 hours
- A lot of techniques in addition to a simple end-to-end ASR
- End-to-End ASR system finally achieved 5.8%
- Classical HMM system (Hybrid DNN/HMM system) 6.7%

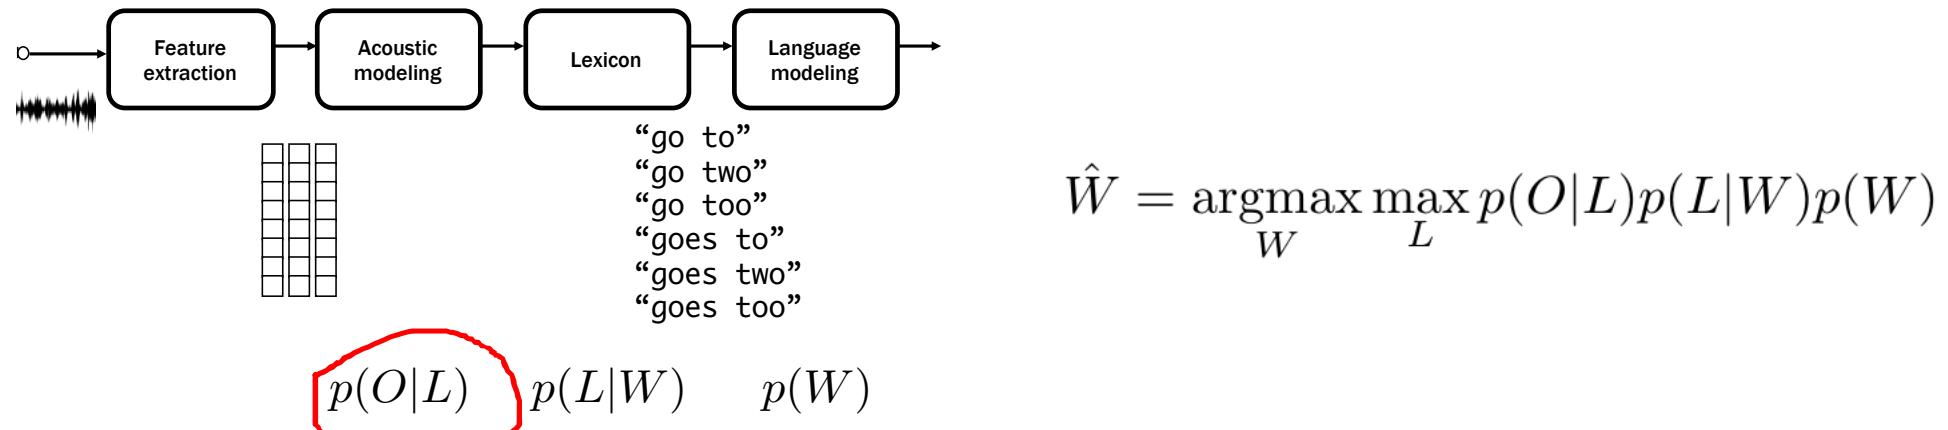
Experiments (< 80 hours, Nov 2018)

- Word Error Rate [%] in English Wall Street Journal (WSJ) task

Models	dev93	eval92	
ESPnet	7.0	4.7	Our best end-to-end
Attention model + word 3-gram LM [Bahdanau 2016]	-	9.3	
CTC + word 3-gram LM [Graves 2014]	-	8.2	
CTC + word 3-gram LM [Miao 2015]	-	7.3	
Attention model + word 3-gram LM [Chorowski 2016]	9.7	6.7	
Hybrid CTC/attention, multi-level LM	-	5.6	End-to-end best
Wav2Letter with gated convnet	-	5.6	
HMM/DNN + sMBR + word 3-gram LM	6.4	3.6	DNN/HMM (pipeline) best
HMM/DNN + sMBR + word RNN-LM	5.6	2.6	

Why HMM-based classical ASR is better than End-to-End ASR?

- Classical HMM



- We can **separately** train acoustic, lexicon, and language model
- We can **incorporate pronunciation dictionary information** through $p(L|W)$
- We can **train the language model $p(W)$ only with text** (newspaper, web, etc.)

On the other hand end-to-end ASR **always require the pair data** to train $p(W|O)$
(We can incorporate a language model $p(W)$ by $p(W)^\alpha p(W|O)$ heuristically)

HMM/DNN vs. CTC vs. Attention

- Conditional independence assumptions
- Language models
- Use of pronunciation lexicon information
- Implementation

Let's discuss the difference

CTC vs. HMM vs. Attention

HMM

$$p(W|O)$$

- With Bayes rule and CIA (separate acoustic, lexicon, and language models)

$$\sum_{Z,L} p(O, Z|L) p(L|W) p(W)$$

- 1st order Markov and frame-level decomposition
 $p(O|Z)p(Z|L) \rightarrow \prod_t p(o_t|z_t)p(z_t|z_{t-1}, L)$
- Replace the likelihood function $p(o_t|z_t)$ with a DNN based on the pseudo likelihood trick

CTC

$$p(C|O)$$

- No Bayes rule**, but CIA (separate acoustic and language model)

$$\sum_Z p(C|Z) p(Z|O)$$

- 1st order Markov and frame-level decomposition
 $p(Z|O)p(C|Z) \rightarrow \prod_t p(z_t|O)p(z_t|z_{t-1}, C)p(C)$
- Replace the frame-level posterior distribution $p(z_t|O)$ with a DNN

Attention

$$p(C|O)$$

- No CIA, no separate LM**

$$\prod_j p(c_j|c_{1:j-1}, \mathbf{v}_j), \mathbf{v}_j = \sum_j a_{jt} \mathbf{h}'_t$$