

Information Retrieval

Oct 1, 2024 @ Introduction to Human Language Technology

Eugene Yang eugene.yang@jhu.edu

Slides borrowed from SIGIR24 Tutorial
“Neural Methods for Cross-Language Information Retrieval”

What is Information Retrieval?

(relevant)

Retrieve information from a storage
based on user's information need

Don't we have Google?

Yes, **but Google is not all.**




What if I'm looking for the person?

Why in a list?

Google

Johns Hopkins

×



All

Images

News

Maps


Shopping

Videos

Web

More


Tools

 Johns Hopkins Federal Credit Union
<https://www.jhfcu.org>

⋮

[Johns Hopkins Federal Credit Union | Banking, Loans ...](#)


JHFCU is a full-service financial institution proudly serving retirees and employees of the Hopkins Community. Join us today!

 U.S. News & World Report
<https://www.usnews.com> › Education › Colleges

⋮

[Johns Hopkins University \(JHU\) - Profile, Rankings and Data](#)


Johns Hopkins University is a private institution that was founded in 1876. It has a total undergraduate enrollment of 6,090 (fall 2023), its setting is ...

 Facebook · Johns Hopkins University
349.7K+ followers

⋮

[Johns Hopkins University | Baltimore MD](#)


Welcome to the Facebook home of Johns Hopkins, a leader in teaching and research, with programs in arts and music, engineering, humanities, social and natural ...

 X · JohnsHopkins
242.3K+ followers

⋮

[Johns Hopkins University \(@JohnsHopkins\) ...](#)


America's first research university. Leading discovery and sharing knowledge to better the world since 1876.

 Johns Hopkins Center for Health Security
<https://www.centerforhealthsecurity.org>

⋮

[Johns Hopkins Center for Health Security](#)


The Johns Hopkins Center for Health Security explores how new policy approaches, scientific advances, and technological innovations can strengthen health ...

 Center for a Livable Future
<https://clf.jhsph.edu>

⋮

[Center for a Livable Future](#)

The Johns Hopkins Center for a Livable Future—working toward a healthy, equitable, resilient food system.

 GitHub
<https://github.com> › CSSEGISandData › COVID-19

⋮

[CSSEGISandData/COVID-19: Novel Coronavirus ...](#)

Attribute the data as the "COVID-19 Data Repository by the Center for Systems Science and Engineering (CSSE) at Johns Hopkins University" or "JHU CSSE COVID-19 ...

Why not read my mind?

Why asking me to read?

Why not group the results?

Google Search is just one implementation

Google trained us well!

- Even faster?
- Smarter?
- Cross language?

Hard Matching Problem

- Text to text
 - Search in notes
 - Cross language search
 - Cross domain search
- Text to other modalities
 - Image search
 - Video search

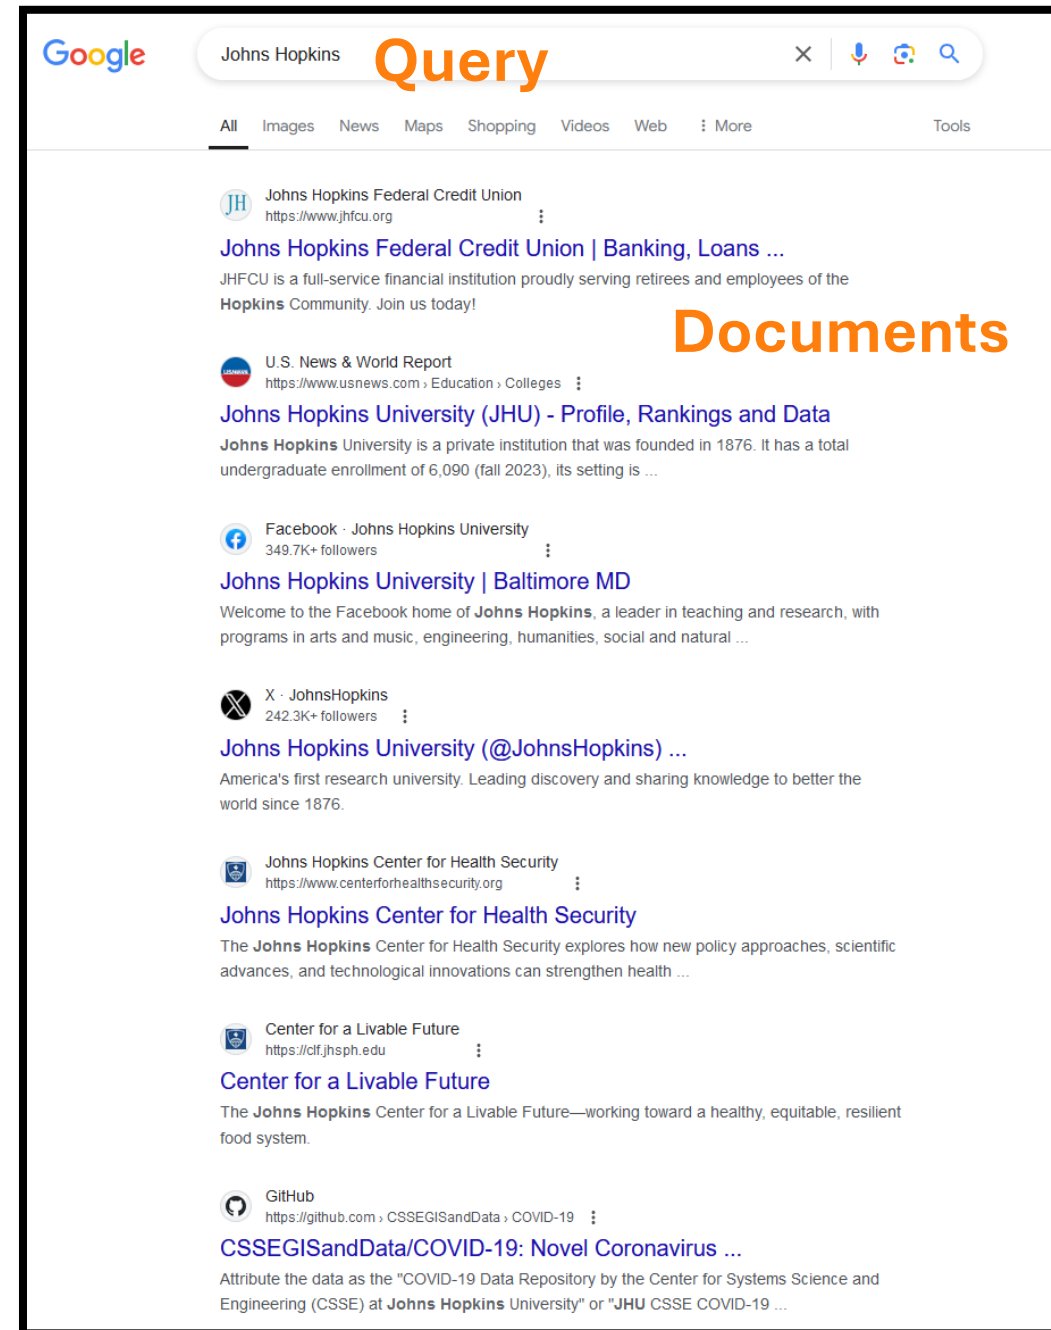
Different Search Process

- Iterative search
 - e.g., electronic discovery and systematic review
- Conversational search
 - Alexa search
- Recommendation systems
- (Set Retrieval)

Core Problem

- Matching problem
- Do it fast

Ranked List



The screenshot shows a Google search for "Johns Hopkins" with the word "Query" in orange. The search results are displayed as a ranked list of documents. The top result is "Johns Hopkins Federal Credit Union" with a URL and a brief description. The second result is "U.S. News & World Report" with a link to "Johns Hopkins University (JHU) - Profile, Rankings and Data". The third result is "Facebook · Johns Hopkins University" with a link to "Johns Hopkins University | Baltimore MD". The fourth result is "X · JohnsHopkins" with a link to "Johns Hopkins University (@JohnsHopkins) ...". The fifth result is "Johns Hopkins Center for Health Security" with a link to "Johns Hopkins Center for Health Security". The sixth result is "Center for a Livable Future" with a link to "Center for a Livable Future". The seventh result is "GitHub" with a link to "CSSEGISandData/COVID-19: Novel Coronavirus ...". The word "Documents" is written in large orange text on the right side of the search results.

Google Johns Hopkins Query

All Images News Maps Shopping Videos Web More Tools

Documents

Johns Hopkins Federal Credit Union
https://www.jhfcu.org
Johns Hopkins Federal Credit Union | Banking, Loans ...
JHFCU is a full-service financial institution proudly serving retirees and employees of the Hopkins Community. Join us today!

U.S. News & World Report
https://www.usnews.com › Education › Colleges
Johns Hopkins University (JHU) - Profile, Rankings and Data
Johns Hopkins University is a private institution that was founded in 1876. It has a total undergraduate enrollment of 6,090 (fall 2023), its setting is ...

Facebook · Johns Hopkins University
349.7K+ followers
Johns Hopkins University | Baltimore MD
Welcome to the Facebook home of Johns Hopkins, a leader in teaching and research, with programs in arts and music, engineering, humanities, social and natural ...

X · JohnsHopkins
242.3K+ followers
Johns Hopkins University (@JohnsHopkins) ...
America's first research university. Leading discovery and sharing knowledge to better the world since 1876.

Johns Hopkins Center for Health Security
https://www.centerforhealthsecurity.org
Johns Hopkins Center for Health Security
The Johns Hopkins Center for Health Security explores how new policy approaches, scientific advances, and technological innovations can strengthen health ...

Center for a Livable Future
https://clf.jhsph.edu
Center for a Livable Future
The Johns Hopkins Center for a Livable Future—working toward a healthy, equitable, resilient food system.

GitHub
https://github.com › CSSEGISandData › COVID-19
CSSEGISandData/COVID-19: Novel Coronavirus ...
Attribute the data as the "COVID-19 Data Repository by the Center for Systems Science and Engineering (CSSE) at Johns Hopkins University" or "JHU CSSE COVID-19 ...

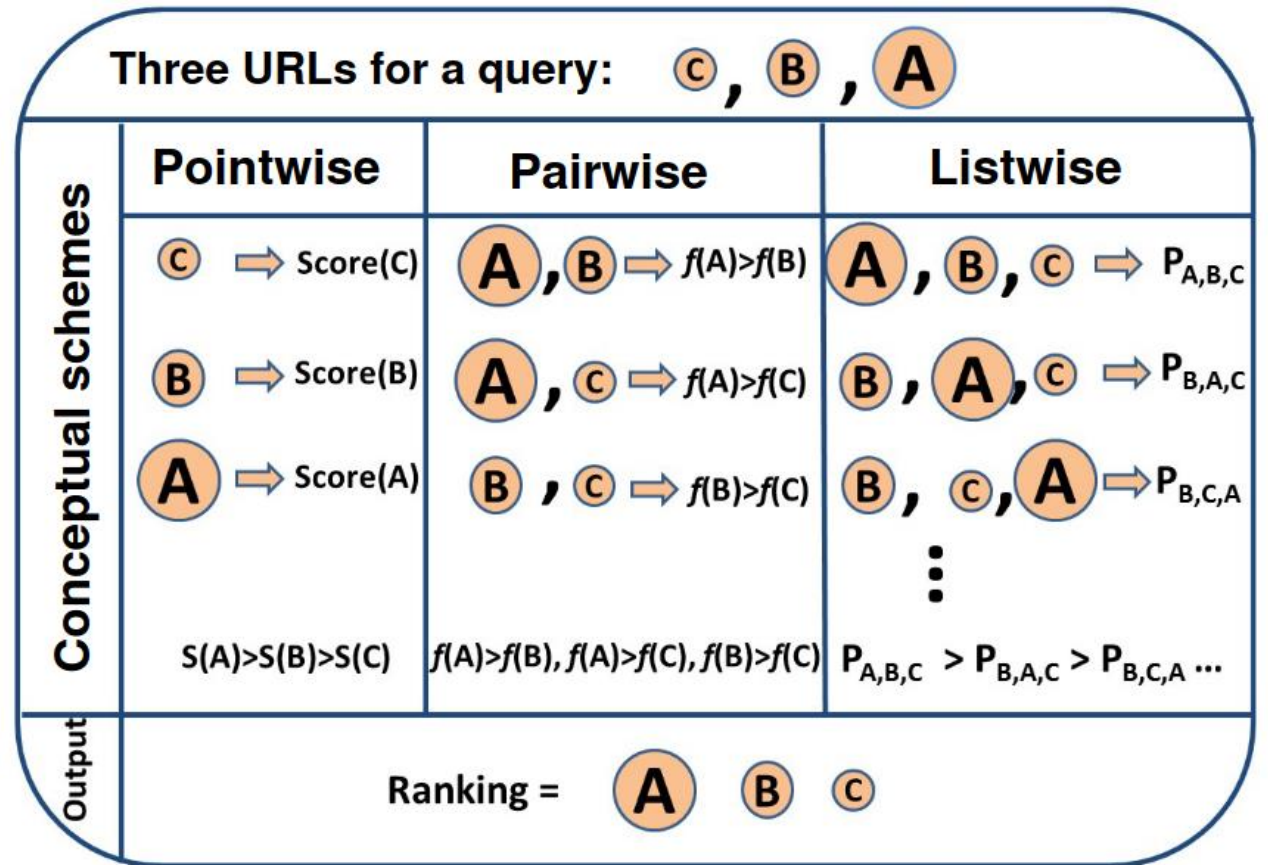
Agenda

- What is information retrieval?
- Retrieval Modeling and Pipeline
 - Statistical and Neural
- Evaluation
- State of IR Research and active research problems

Retrieval Modeling and Pipeline

Three main ways

- **Pointwise**
- Pairwise
- Listwise
- And combinations of them



<https://medium.com/vptech/learning-to-rank-at-veepee-ed420fd828e5>

Statistical Models

$$\text{score}(D, Q) = \sum_{\text{For each query term}} \boxed{\text{How important the term is}} \times \boxed{\text{How often the term appear in the D}}$$

$$\text{score}(D, Q) = \sum_{\text{For each query term}} \boxed{\text{Inverted document frequency}} \times \boxed{\text{Term frequency}}$$

TF-IDF

$$\text{score}(D, Q) = \sum_{i=1}^n \log \frac{N}{n_t} \times \log(f(q_i, D) + 1)$$


BM25

$$\text{score}(D, Q) = \sum_{i=1}^n \text{IDF}(q_i) \cdot \frac{f(q_i, D) \cdot (k_1 + 1)}{f(q_i, D) + k_1 \cdot \left(1 - b + b \cdot \frac{|D|}{\text{avgdl}}\right)}$$


How to make it fast?

- “Fast” in responding to queries
- Better data structure
- Preprocess the data

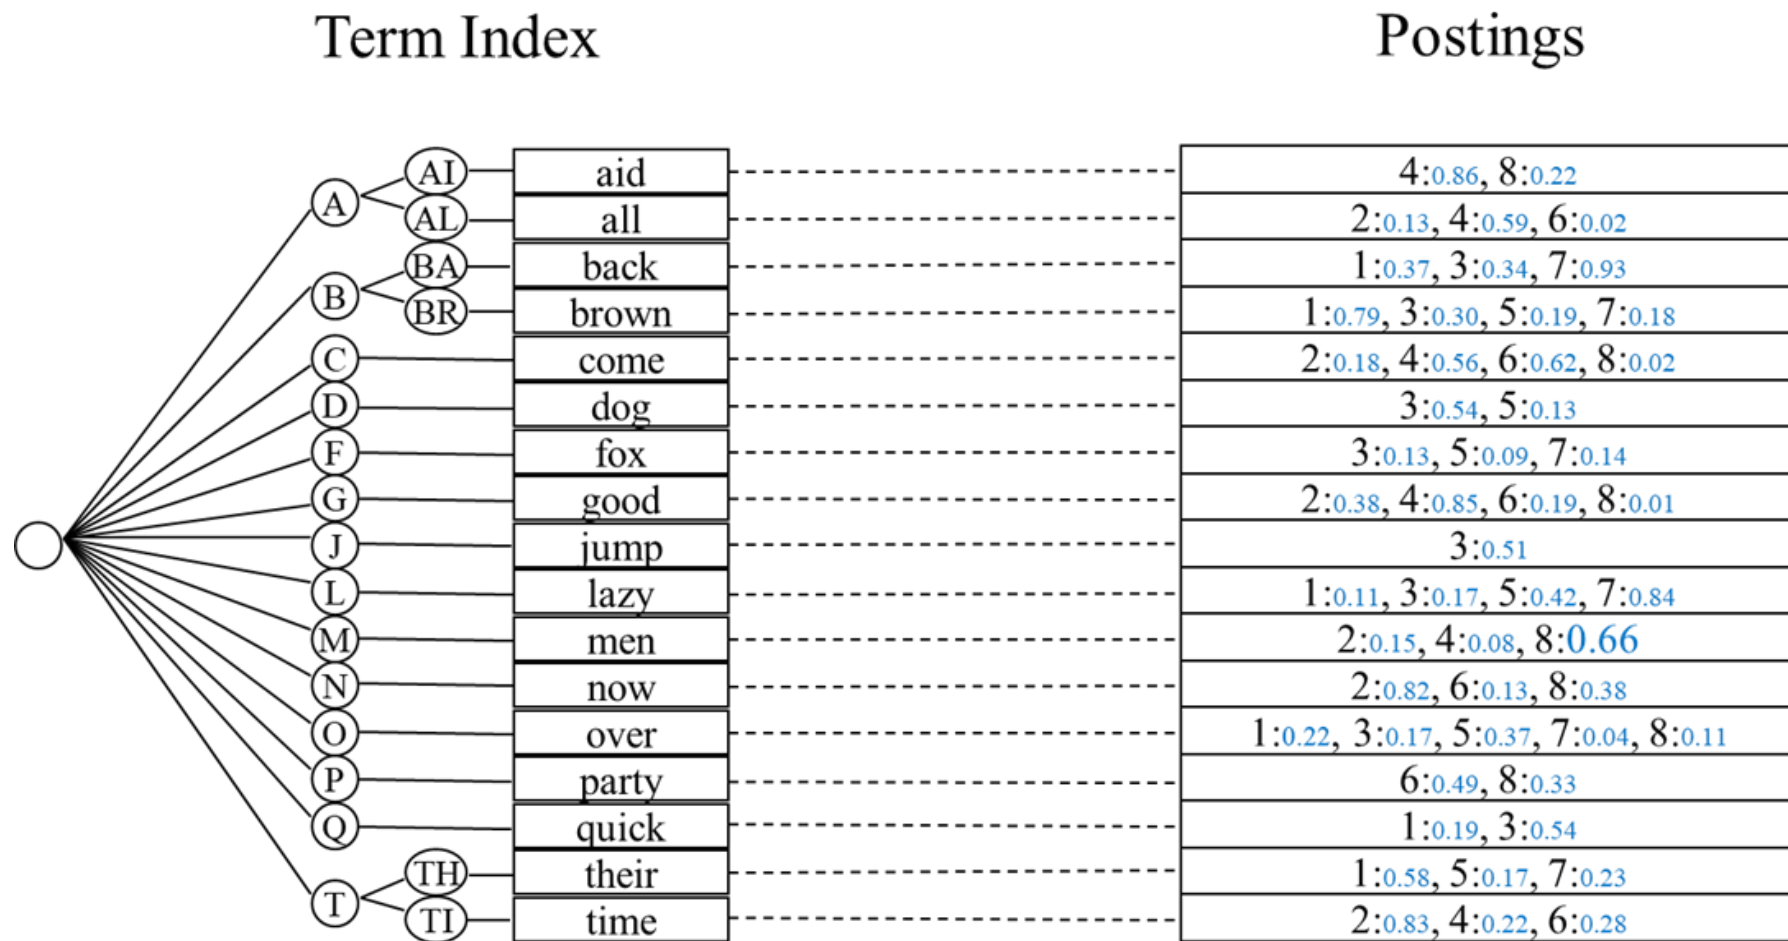
Inverted Index



Inverted Index



Inverted Index



Two-Stage System

- Offline preprocessing and indexing
 - Build the inverted index
- Online query serving
 - Traverse the inverted index and score it



Apache 2.0 licensed

Apache Lucene is distributed under a commercially friendly Apache Software license

Welcome to Apache Lucene

The Apache Lucene™ project develops open-source search software. The project releases a core search library, named Lucene™ core, as well as PyLucene, a python binding for Lucene.

Lucene Core is a Java library providing powerful indexing and search features, as well as spellchecking, hit highlighting and advanced analysis/tokenization capabilities. The **PyLucene** sub project provides Python bindings for Lucene Core.

Latest Lucene Core News

Apache Lucene™ 8.11.4 available (24.Sep)

Apache Lucene™ 9.11.1 available (27.Jun)

Apache Lucene™ 9.11.0 available (06.Jun)

Projects


[Lucene Core \(Java\)](#)
[PyLucene](#)
[Open Relevance](#)
(Discontinued)

About

[License](#)
[Who We are](#)
[TLP News](#)
[Code of Conduct](#)



ANNOUNCEMENT: The **Solr™** sub project has moved to a separate Top Level Project (TLP). All things Solr can now be found at <https://solr.apache.org/>. Mailing lists and git repositories have changed, please see details on the Solr website.



The heart of the free and open Elastic Stack

Elasticsearch is a distributed, RESTful search and analytics engine, scalable data store, and vector database capable of addressing a growing number of use cases. As the heart of the Elastic Stack, it centrally stores and manages data, providing fine-tuned relevancy, and powerful analytics.

[Start free trial](#)

[Download Elasticsearch](#)

elastic / elasticsearch (Public)

Notifications

Fork 24.7k

Star 69.7k

<> Code Issues 3.9k Pull requests 773 Actions Projects Security Insights

main 391 Branches 404 Tags

Go to file

Code

About

Free and Open Source, Distributed, RESTful Search Engine

www.elastic.co/products/elasticsearch

java search-engine elasticsearch

Readme

View license

Security policy

Activity

Custom properties

69.7k stars

2.7k watching

24.7k forks

Report repository

Releases 155

Elasticsearch 8.15.1 (Latest) 3 weeks ago

+ 154 releases

Packages


No packages published

smalyshev	Improve DateTime error handling and add some bad date tests...	5e06092 · 4 hours ago	79,984 Commits
.buildkite	Run snyk dependency checks on 8.x (#113117)		last week
.ci	Workaround packaging tests failures on debian10 (#113117)		9 hours ago
.github	Remove Analytical engine CODEOWNERS (#113178)		2 days ago
.idea	Don't apply IntelliJ illegal module dependency inspection ...		10 months ago
benchmarks	ESQL: Speed up CASE for some parameters (#112295)		yesterday
build-conventions	Add AGPLv3 as a supported license		2 weeks ago
build-tools-internal	Always use CLDR locale on ES v9 (#113184)		3 days ago
build-tools	Add AGPLv3 as a supported license		2 weeks ago
client	Add AGPLv3 as a supported license		2 weeks ago
dev-tools	Add AGPLv3 as a supported license		2 weeks ago
distribution	Always use CLDR locale on ES v9 (#113184)		3 days ago
docs-mdx/painless	[DOCS] Adds an MDX file for testing purposes. (#106165)		6 months ago
docs	Improve DateTime error handling and add some bad dat...		4 hours ago
gradle	Update Gradle wrapper to 8.10.1 (#112948)		last week
libs	Small performance improvement in h3 library (#113385)		2 days ago


Can we go beyond surface forms?

neural language models


Bi-Encoder



One Dense Vector
Per Sequence
e.g., DPR




One **Sparse** Vector
Per Sequence
e.g., SPLADE




Multiple Dense Vectors
Per Sequence
e.g., ColBERT

Cross Encoder




Joint Encoder
e.g., monoBERT



Online


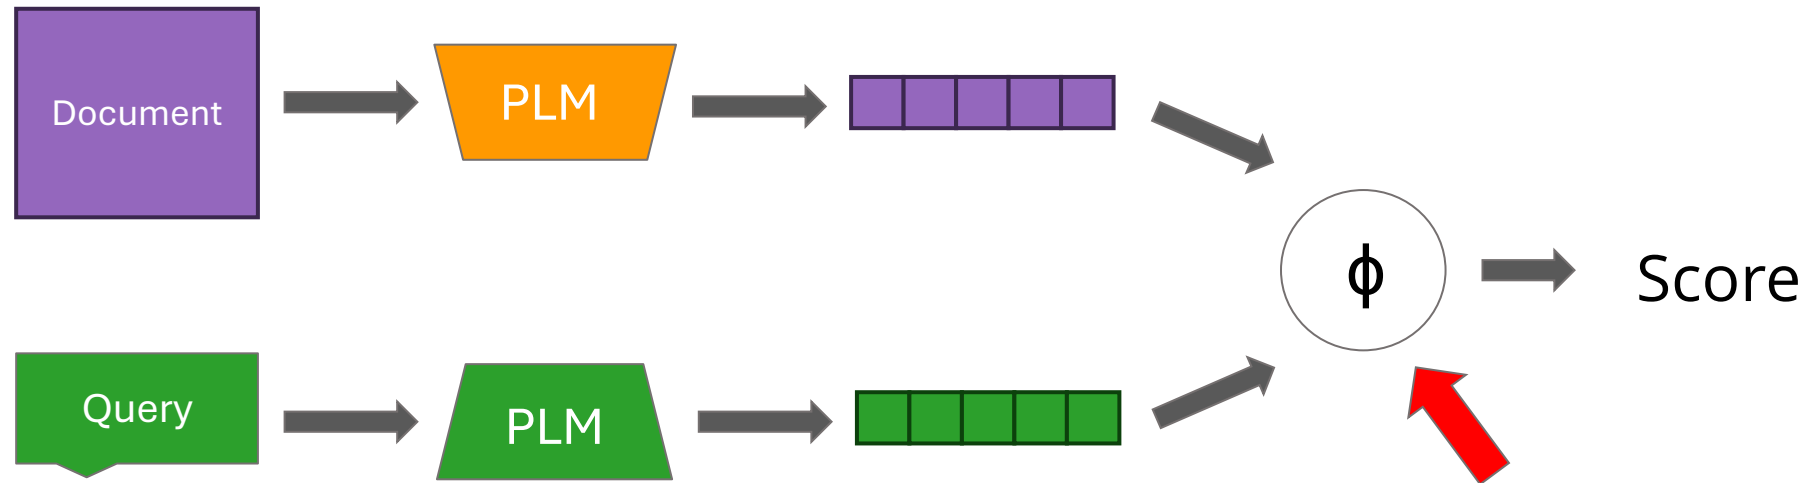
Offline




Both Online

Separate query and document processing


One Vector per Query, One Vector per Document



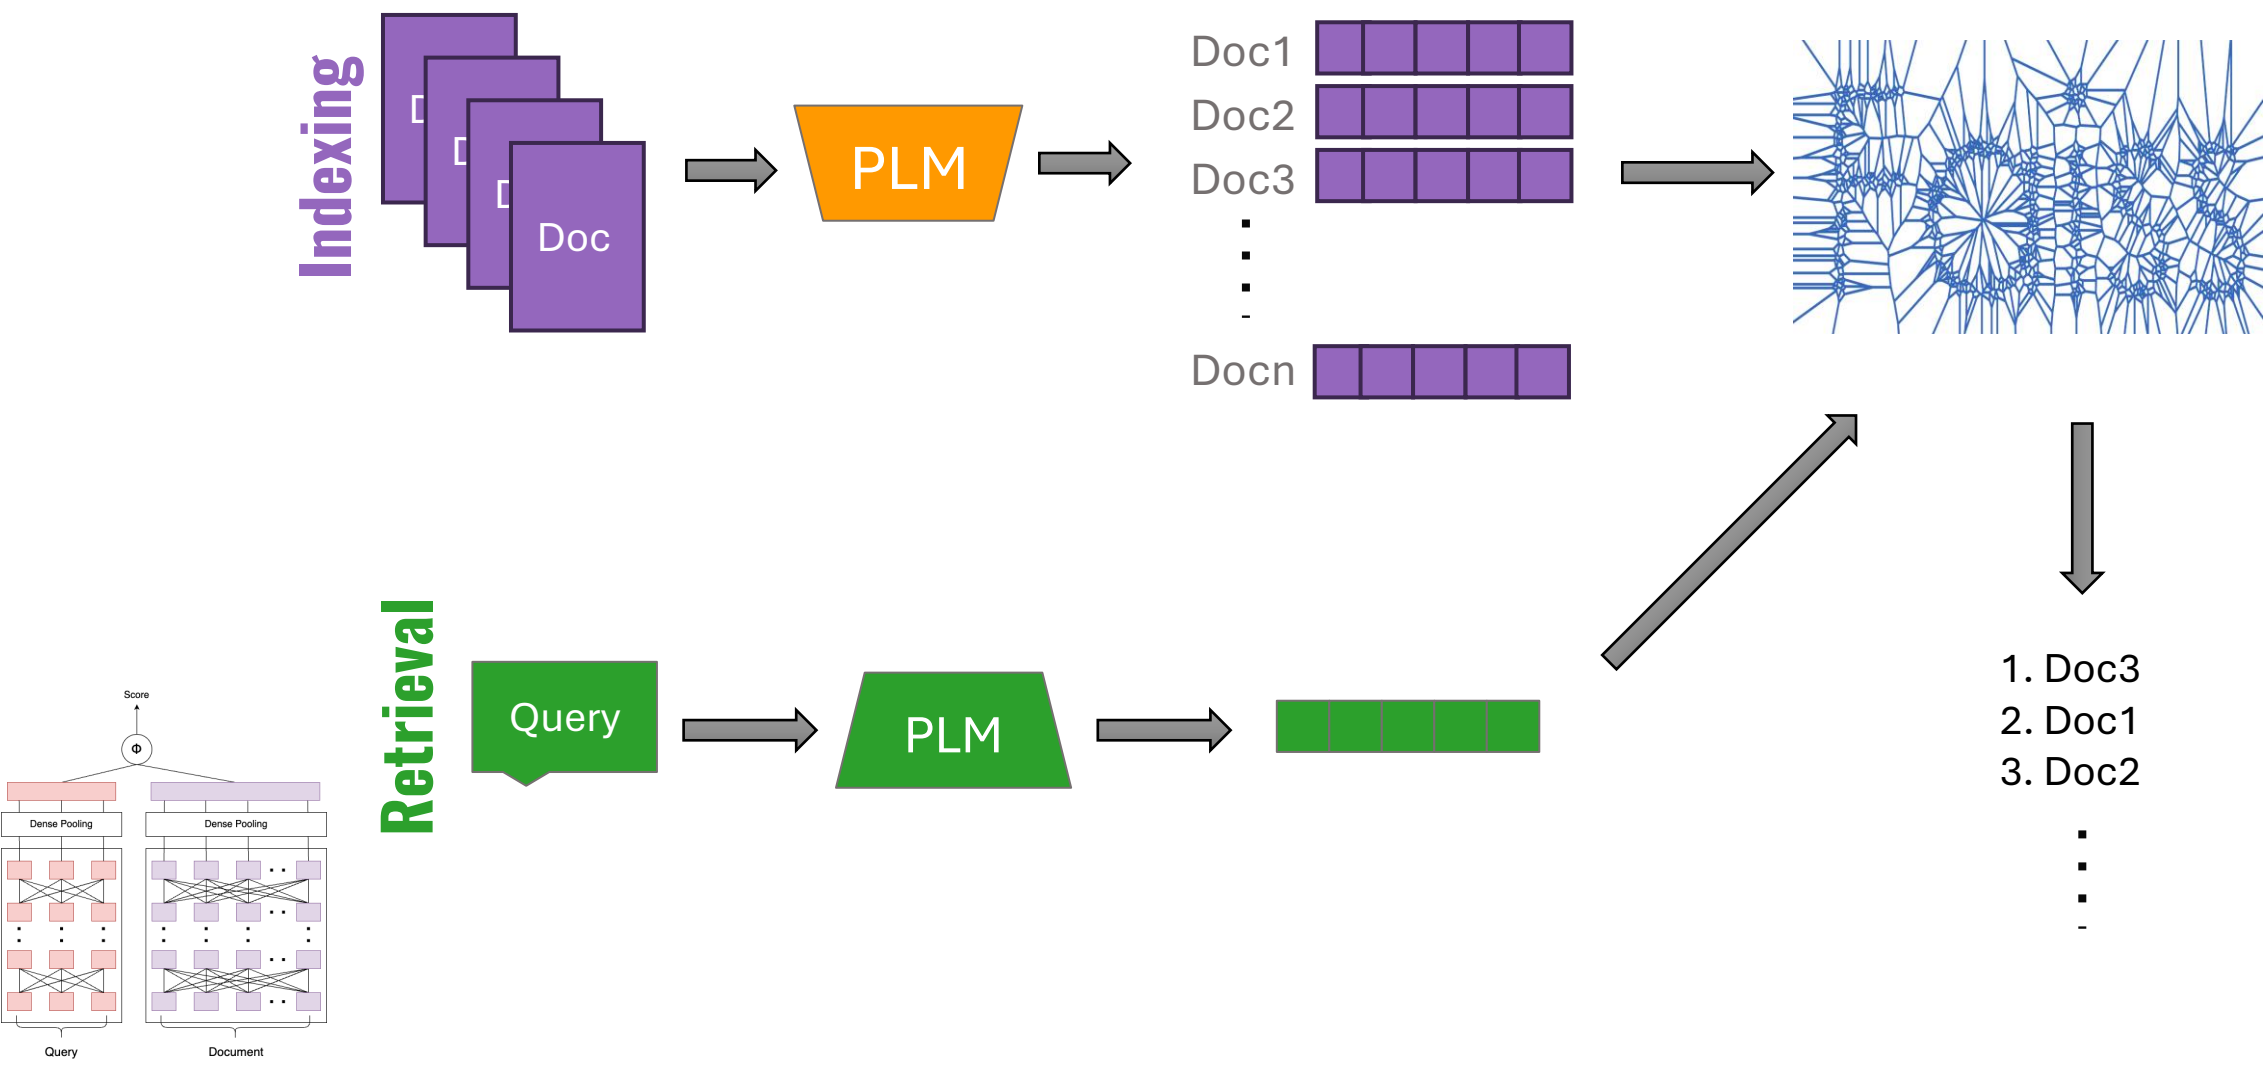
Nearest Vectors aka Neighbors




- Linear Search
 - Slow (scales linearly in size of document collection)
- Approximate Methods (e.g., Product Quantization) → **ANN**
 - Faster Search
- Runtime Efficiency vs Effectiveness




DPR Indexing and Retrieval




Bi-Encoder



One Dense Vector
Per Sequence
e.g., DPR




One **Sparse** Vector
Per Sequence
e.g., SPLADE



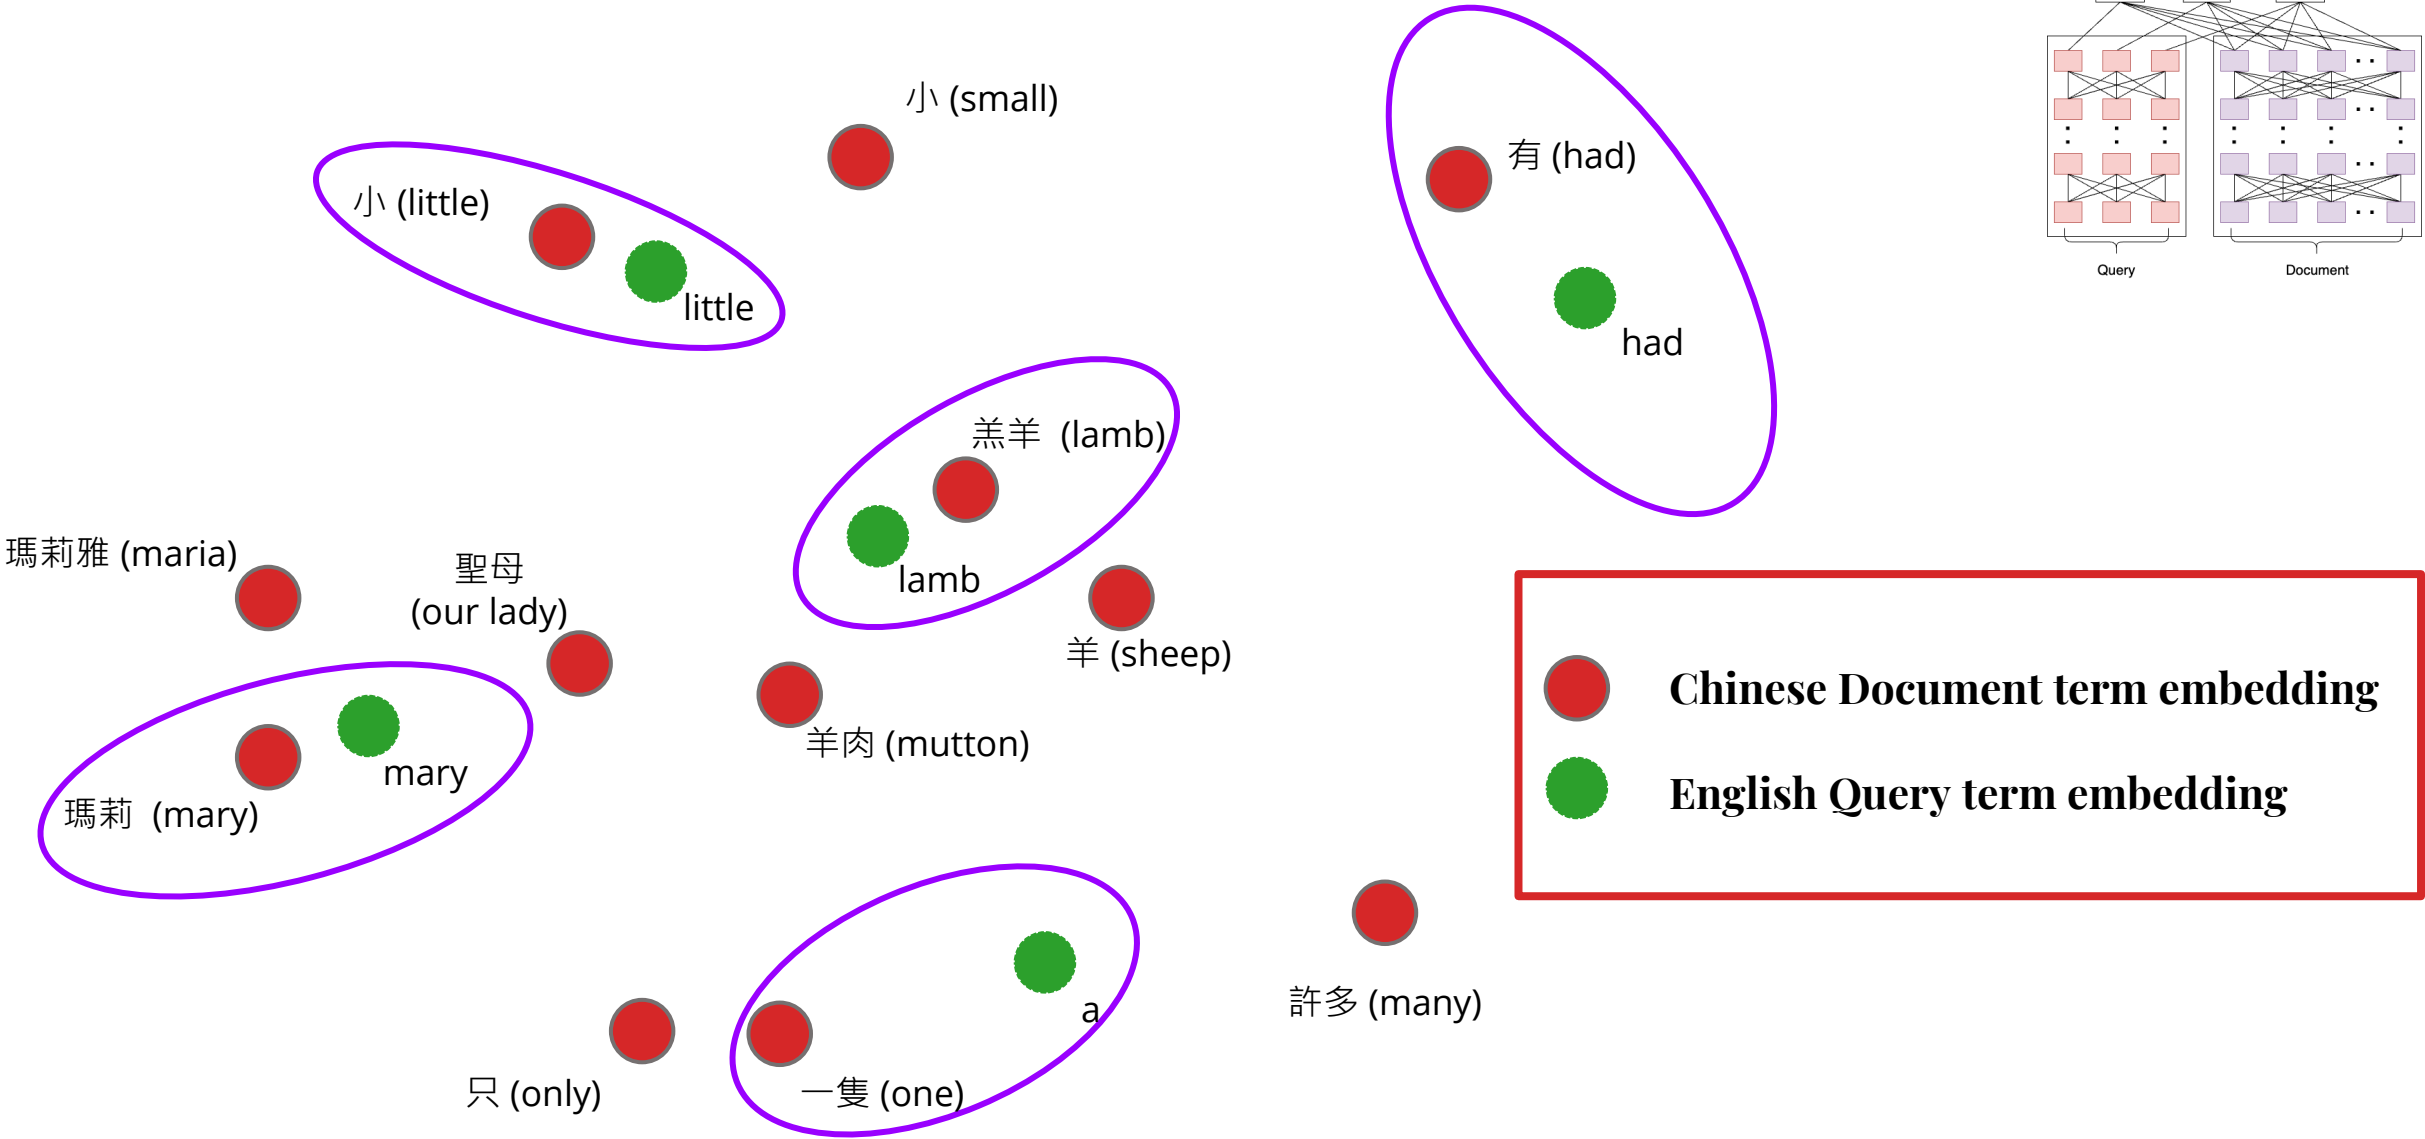
Multiple Dense Vectors
Per Sequence
e.g., ColBERT

Cross Encoder

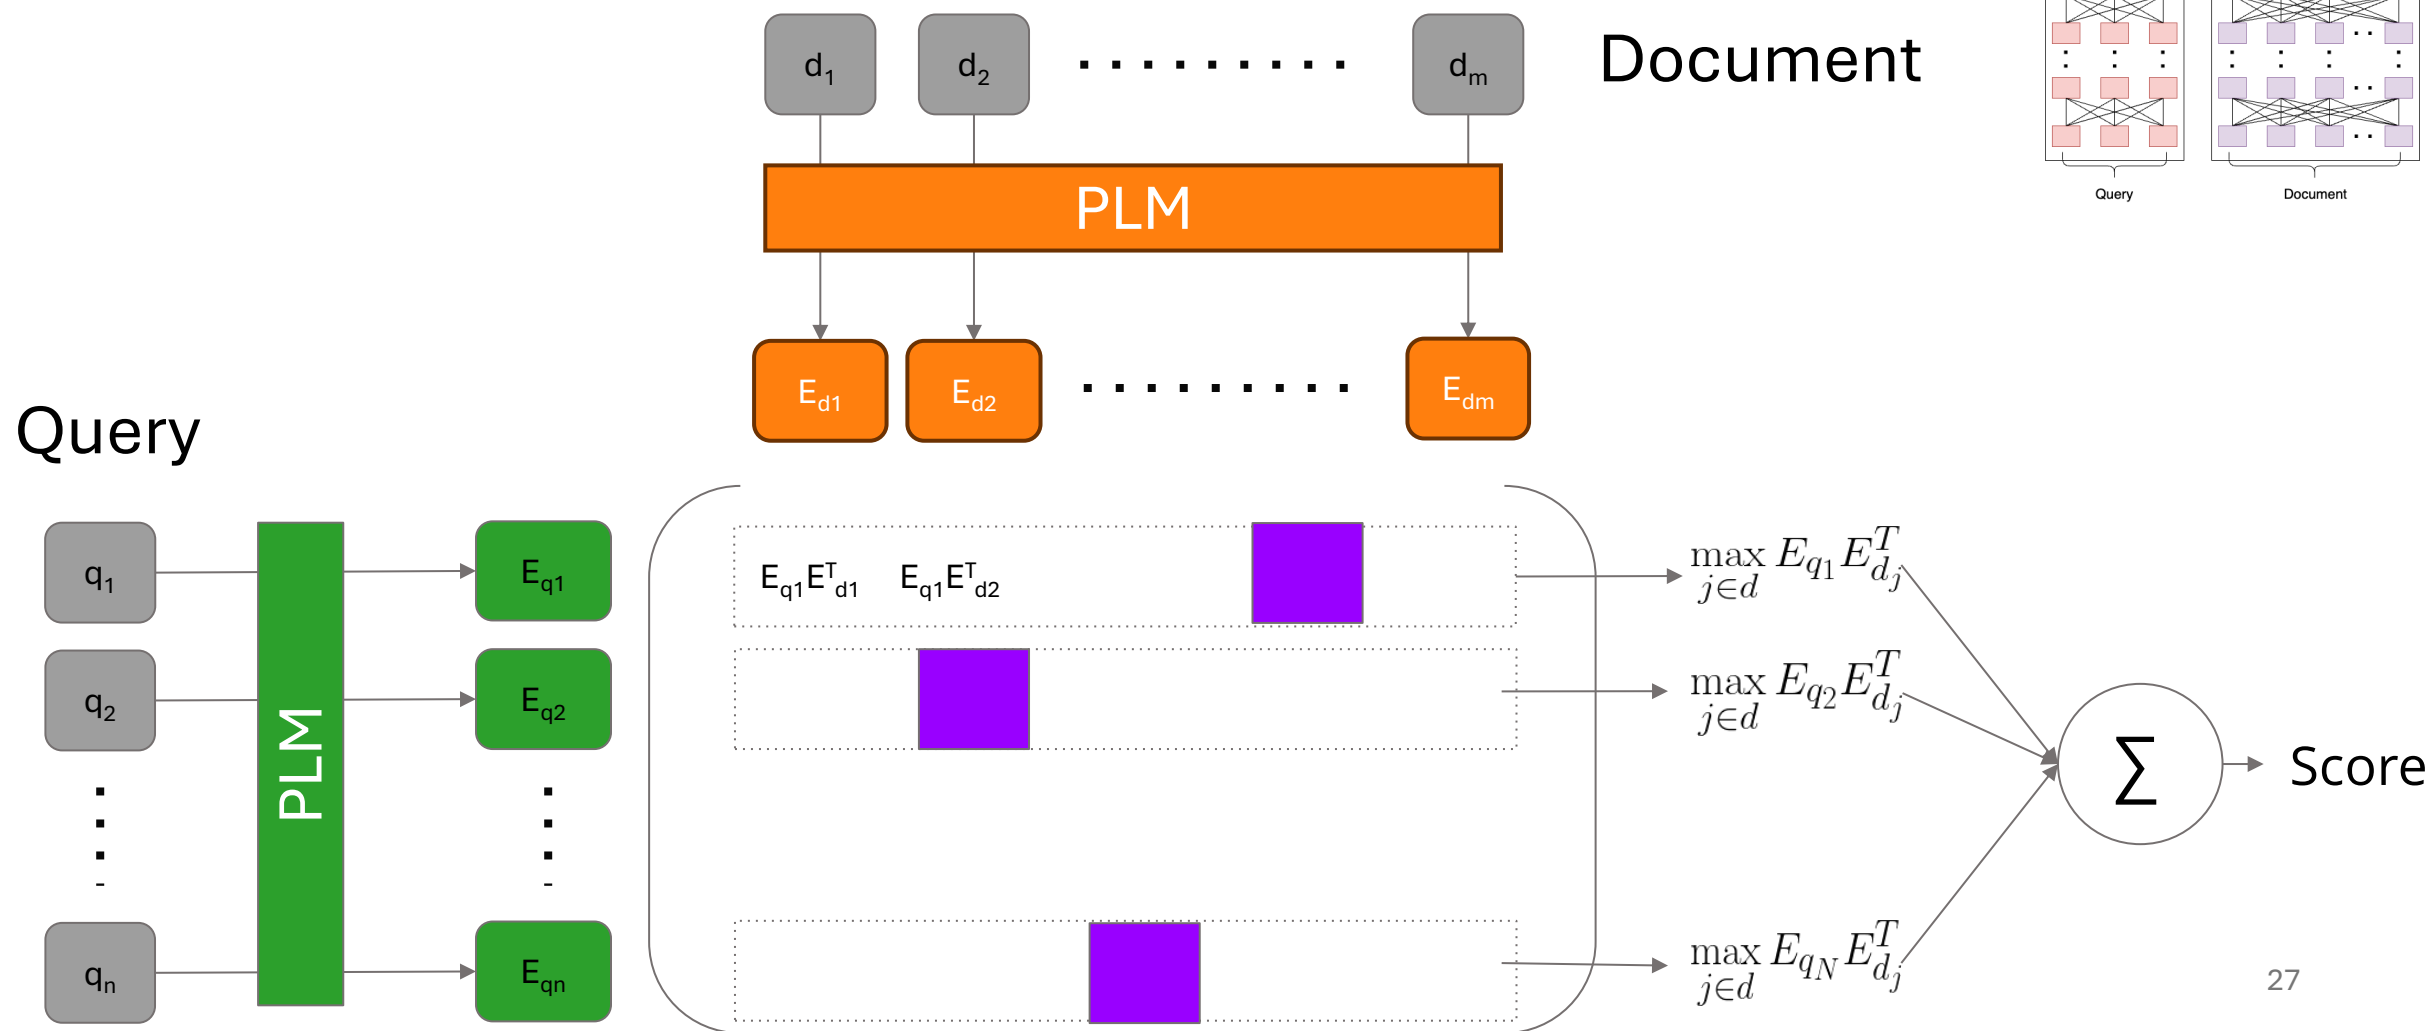


Joint Encoder
e.g., monoBERT

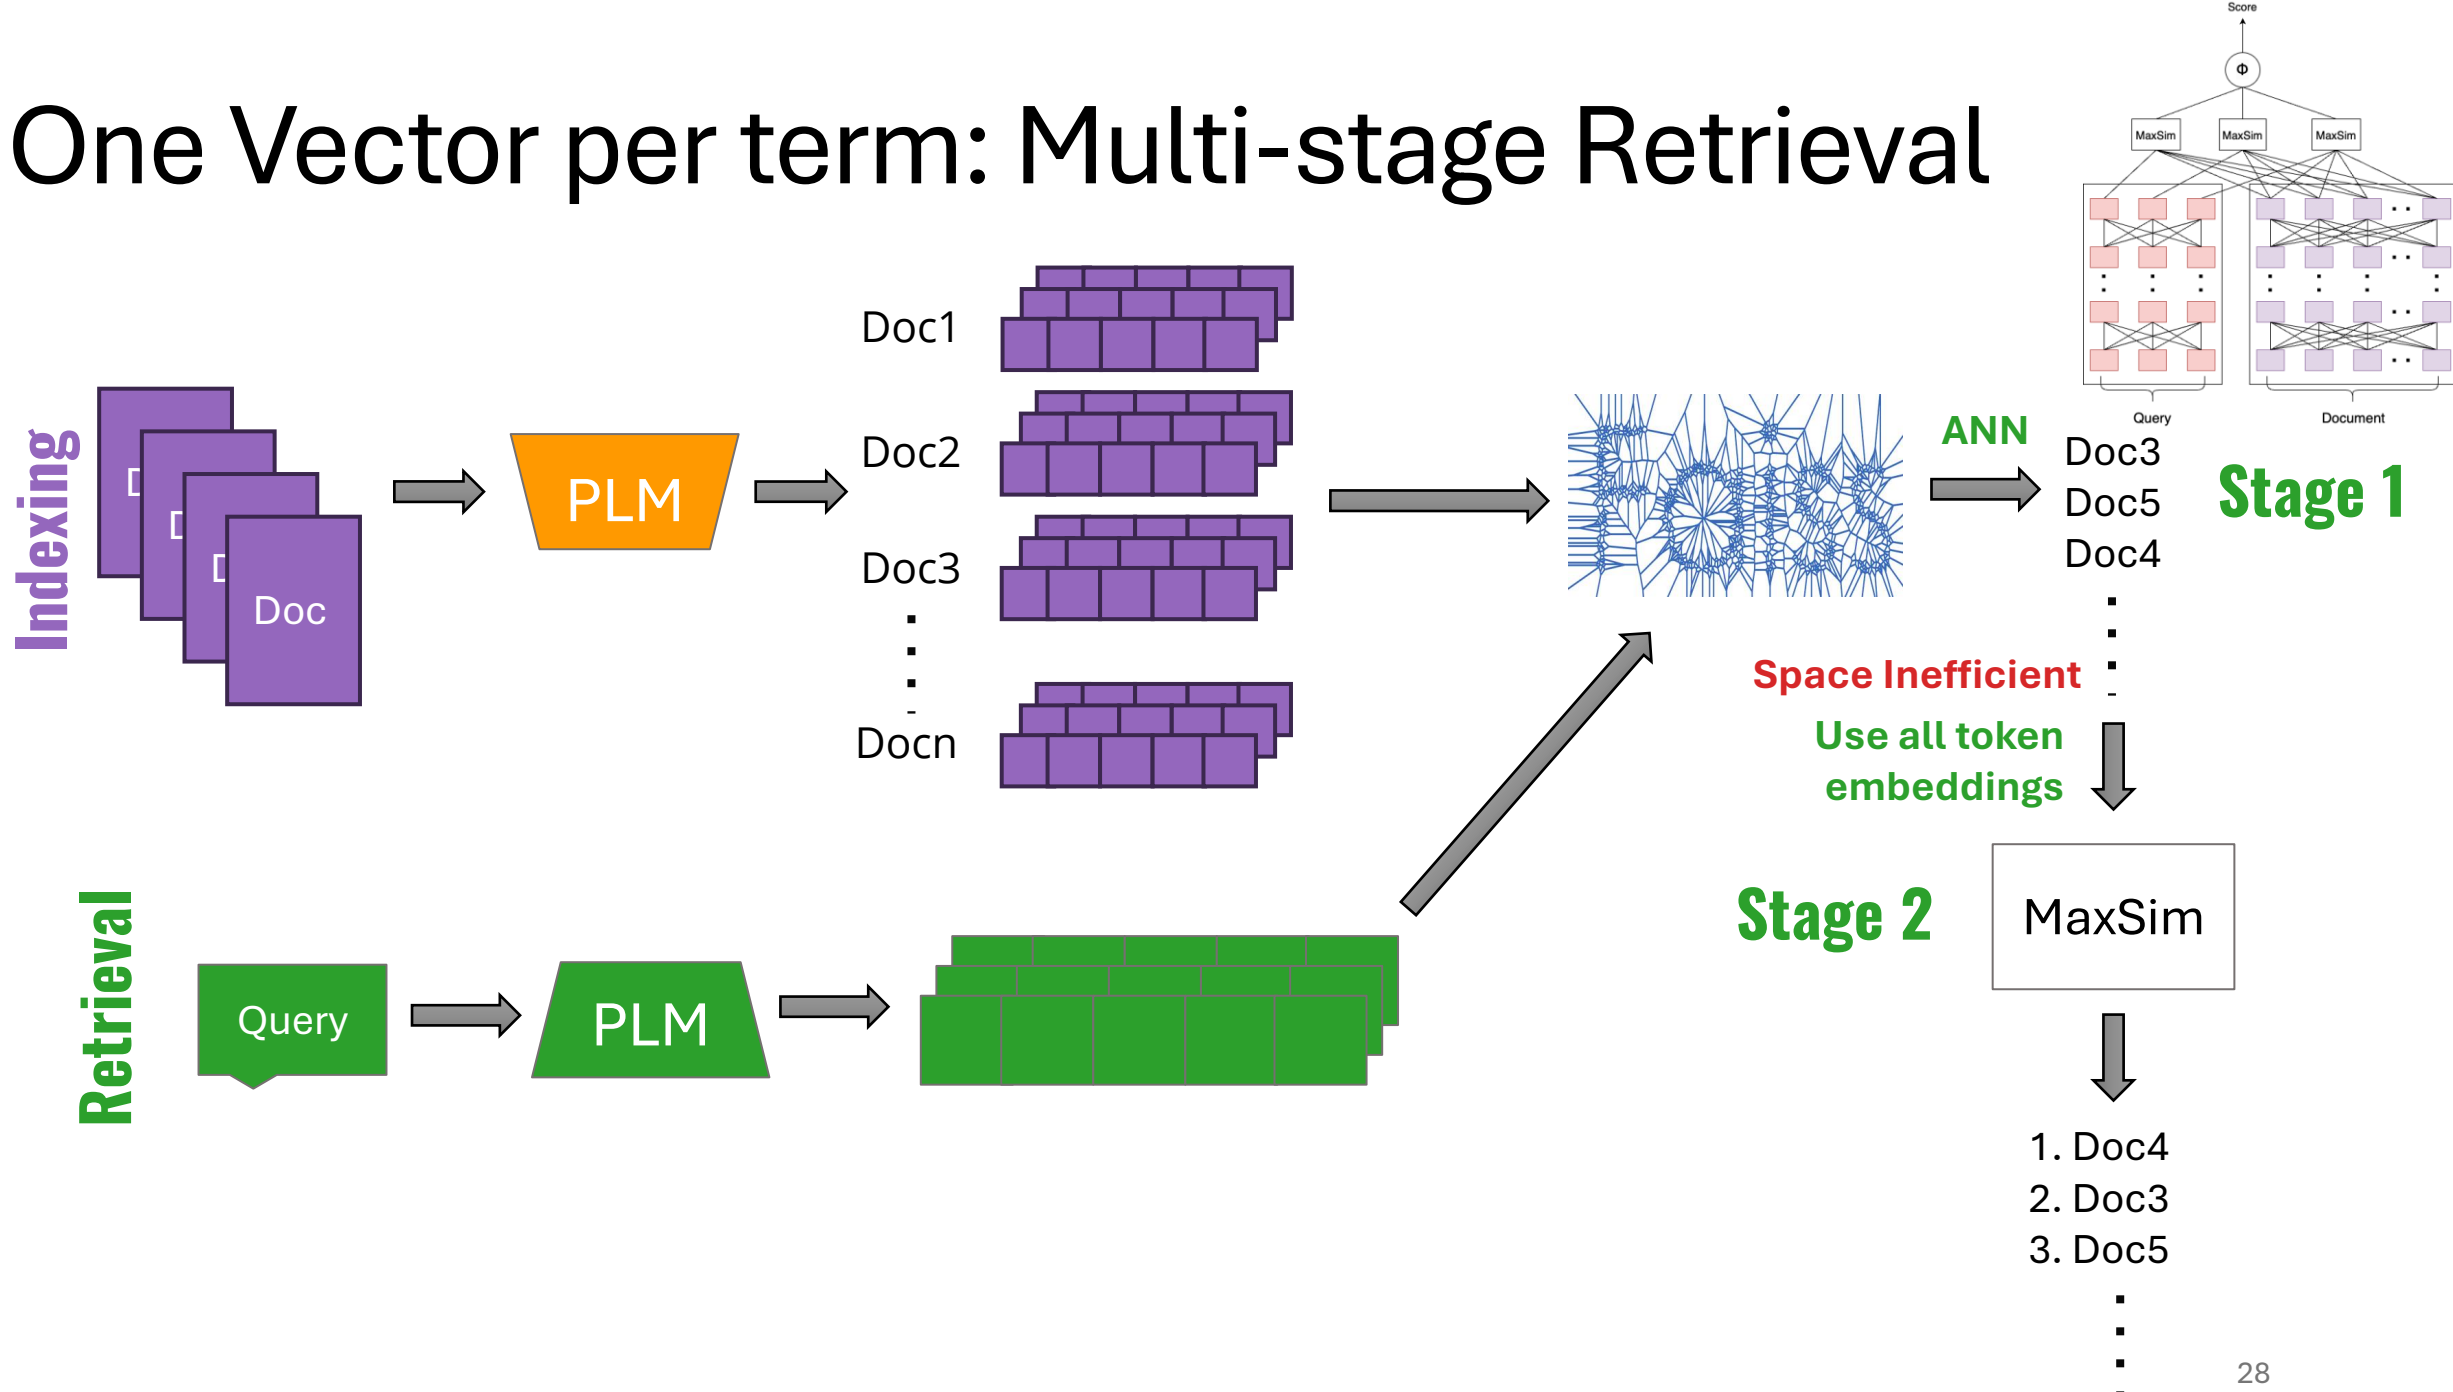
One Vector per Term: MaxSim



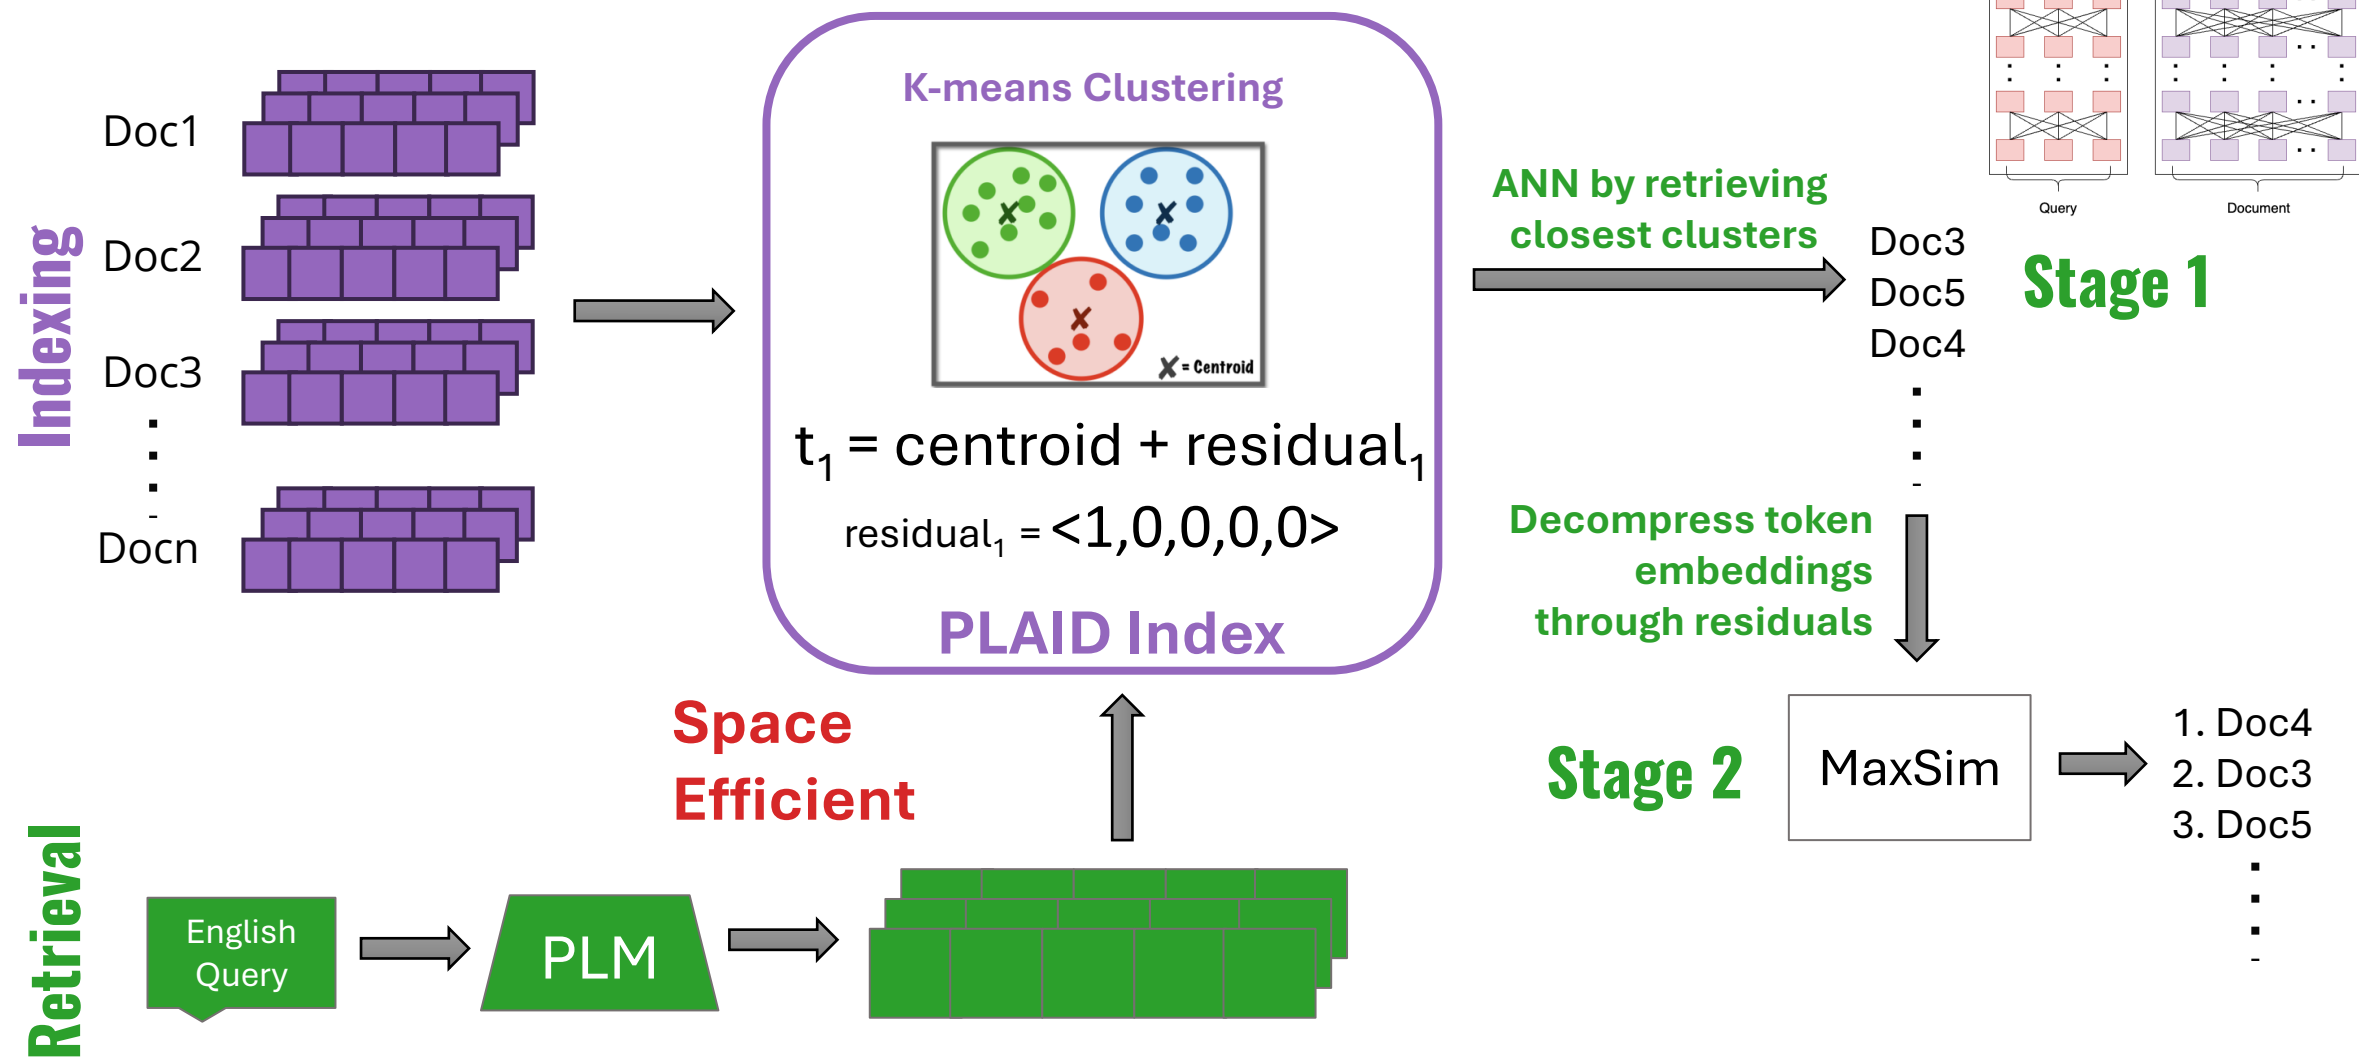
MaxSim in Action -- ColBERT



One Vector per term: Multi-stage Retrieval



Efficient PLAID Indexing Architecture



hltcoe / ColBERT-XPublic

Notifications

<> CodeIssuesPull requestsActionsProjectsSecurityIns

plaid-xGo to fileCode

eugene-yangUpdate collection_utils.py944b709 · 3 months ago

colbert	Update collection_utils.py	3 months ago
docs	update theme	2 years ago
.gitignore	update install	last year
LICENSE	Initial commit with the new API...	3 years ago
MANIFEST.in	Version 0.3.1 (#7)	7 months ago
README.md	Update README.md (#8)	4 months ago
conda_env.yml	commit the environment file	2 years ago
conda_env_cpu.yml	Make end to end test work wit...	2 years ago
requirements.txt	Version 0.3.1 (#7)	7 months ago
setup.py	Version 0.3.1 (#7)	7 months ago

READMEMIT license

PLAID-X

This is a generalized version of [PLAID](#) and the previous ColBERT-X for CLIR. The codebase supports models trained with the original ColBERT-X scripts, which are not compatible with the PLAID codebase released from the Stadford Futuredata Group.

stanford-futuredata / ColBERTPublic

NotificationsFork 376Star 2.9k

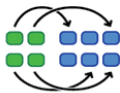
<> CodeIssues 71Pull requests 20DiscussionsActionsProjectsSecurity

mainGo to fileCodeAbout

okhatUpdate README.md85837b6 · 3 weeks ago

baleen	Minor updates to Baleen	3 years ago
colbert	Merge pull request #362 from ...	last month
data	Add data/ with 5k evals	last year
docs	Updated setup.py + intro nb	10 months ago
utility	remove spacy	8 months ago
.env	Add server.py, .env and depen...	last year
.gitignore	Add data/ with 5k evals	last year
LICENSE	Initial commit with the new AP...	3 years ago
LoTTE.md	Add LoTTE download link	2 years ago
MANIFEST.in	Updated setup.py + intro nb	10 months ago
README.md	Update README.md	3 weeks ago
conda_env.yml	remove spacy	8 months ago
conda_env_cpu.yml	remove spacy	8 months ago
server.py	Fix Index root is not used corre...	last year
setup.py	release: version bump (0.2.20)	3 months ago

READMEMIT license



ColBERT

ColBERT (v2)

hltcoe's Collections

Translate-DistillMultilingual Translate-Distill

Translate-Distillupdated Mar 27

Collection of trained model and teacher scores for distillation for paper "Translate-Distill" Code: <https://github.com/hltcoe/ColBERT-X>

Upvote 3

Translate-Distill: Learning Cross-Language Dense Retrieval by Translation and Distillation

Paper · 2401.04810 · Published Jan 9

hltcoe/plaidx-large-zho-tdist-mt5xxl-engzho

Updated Mar 15 · 2

hltcoe/plaidx-large-zho-tdist-mt5xxl-zhozho

Updated Mar 15 · 1


hltcoe/plaidx-large-zho-tdist-t53b-engeng

Updated Mar 15 · 1


hltcoe/plaidx-large-zho-tdist-mt5xxl-engeng

Updated Mar 15 · 25


Bi-Encoder



One Dense Vector
Per Sequence
e.g., DPR




One **Sparse** Vector
Per Sequence
e.g., SPLADE




Multiple Dense Vectors
Per Sequence
e.g., ColBERT

Cross Encoder



Joint Encoder
e.g., monoBERT

High-dimensional Vector: Masked LM



SPLADE




Baltimore Orioles clinch playoff berth for 2nd straight season

The Baltimore Orioles are headed to the playoffs in consecutive years for the first time since the 1990s, clinching no worse than a wild-card berth with a 5-3 win over the New York Yankees paired with Minnesota's loss to Miami




Credit: AP

Baltimore Orioles' Anthony Santander runs the bases after hitting a home run during the sixth inning of a baseball game against the New York Yankees, Tuesday, Sept. 24, 2024, in New York. (AP Photo/Bryan Woolston)



Sparse Pooling

Predicted Vocabulary



baltimore (1.2)

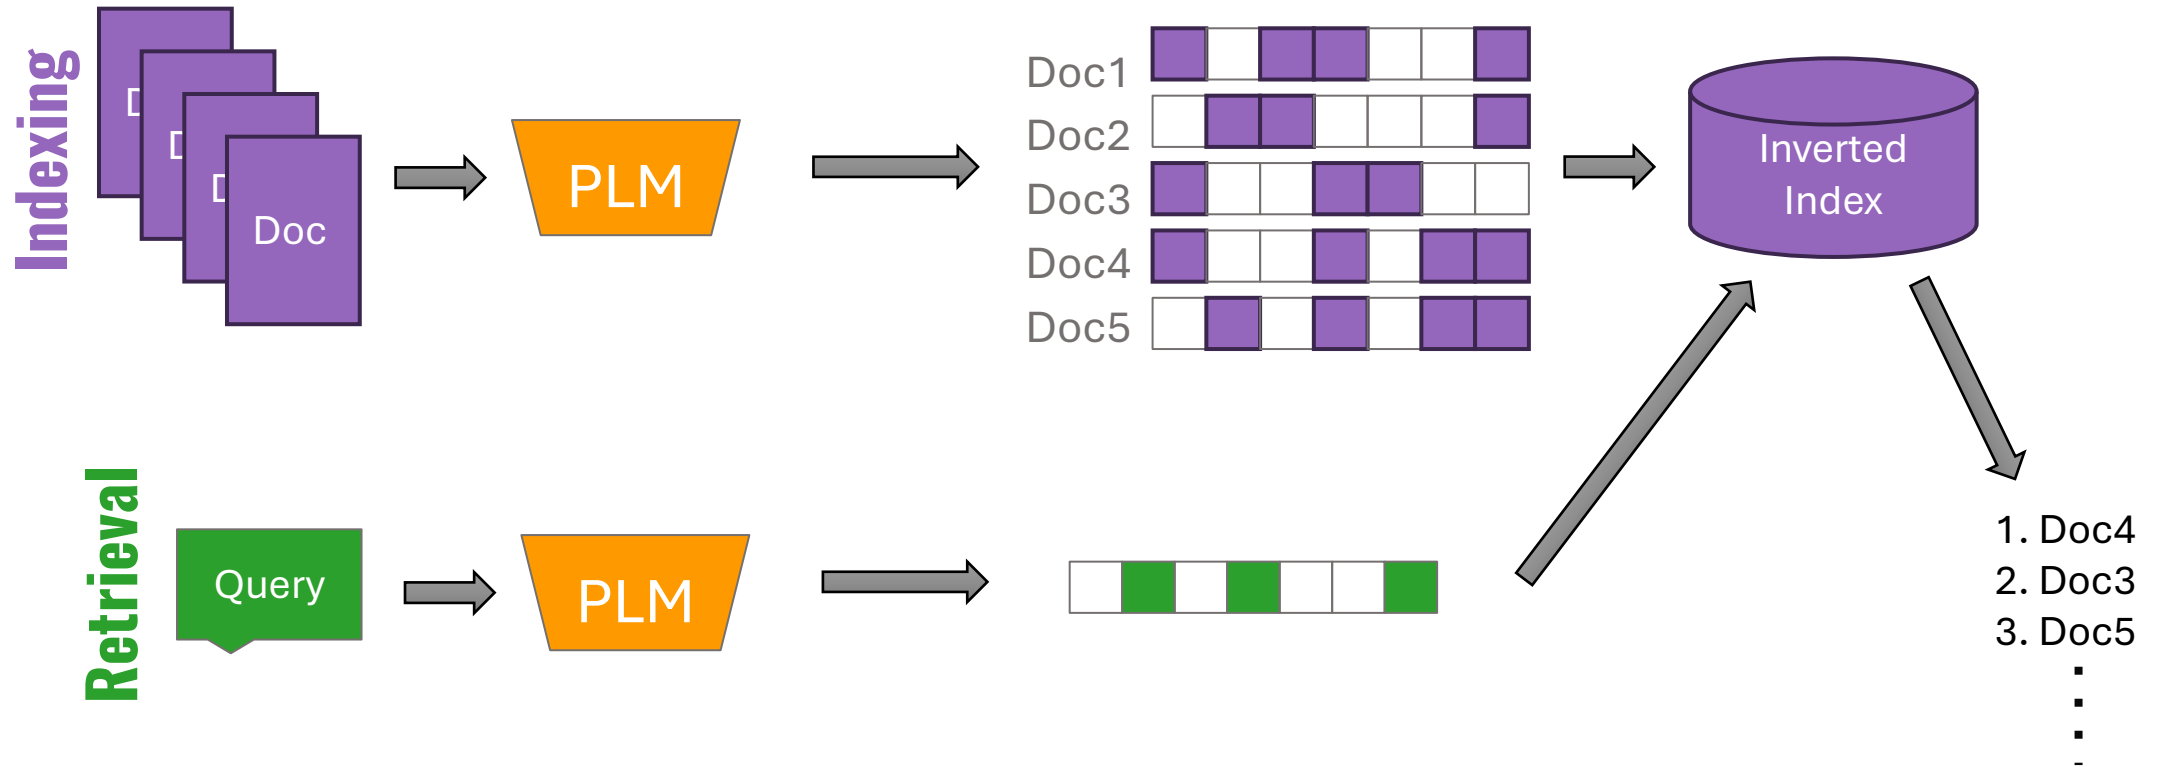
orioles (2.5)

season (0.2)

berth (0.9)

playoff (1.9)

SPLADE Search Pipeline



arXiv > cs > arXiv:2303.13416

Search... All fields Search

Help | Advanced Search

Computer Science > Information Retrieval

[Submitted on 23 Mar 2023]

A Unified Framework for Learned Sparse Retrieval

Thong Nguyen, Sean MacAvaney, Andrew Yates

Learned sparse retrieval (LSR) is a family of first-stage retrieval methods to generate sparse lexical representations of queries and documents from an inverted index. Many LSR methods have been recently introduced, achieving state-of-the-art performance on MSMarco. Despite their success, many LSR methods show substantial differences in their architectures, many LSR methods show substantial differences in their efficiency. Differences in the experimental setups and configurations make it difficult to compare the methods and derive insights. In this work, we unify LSR methods and identify key components to establish an LSR framework. We then reproduce LSR methods under the same perspective. We then reproduce LSR methods using a common codebase and re-train them in the same environment to quantify how components of the framework affect effectiveness. We find that (1) including document term weighting is most important for effectiveness, (2) including query weighting has a small positive effect, (3) document expansion and query expansion have a cancellation effect.

Access Paper:

View PDF

TeX Source

Other Formats

thongnt99 / learned-sparse-retrieval Public

Notifications Fork 5

<> Code Issues 4 Pull requests Actions Projects Security Insights

main Go to file <> Code

Thong Nguyen raw float weights ✓ d702026 · 7 months ago

docs	update skeleton	last year
images	add logo	last year
lsr	raw float weights	7 months ago
.gitignore	Initial commit	2 years ago
LICENSE	Create LICENSE	last year
README.md	Add DOI	7 months ago
beir.sh	Add beir to lsr	last year
clean.py	add file to clean beir trec file	last year
requirements.txt	Merge pull request #7 from ca...	last year
run_all_beir.sh	Add beir to lsr	last year

Readme

Apache-2.0 license

Activity

57 stars

4 watching

5 forks

Report repository

Releases 1

v1.0.0 Latest on Feb 14

Contributors 3

thongnt99 T...

seanmacavane...

cadurosar Ca...

Languages

README

Apache-2.0 license

lsr instructions python 3.9.12 DOI 10.5281/zenodo.10659500

LSR: A unified framework for efficient and effective learned sparse retrieval

TusKANNy / seismic Public

Notifications Fork 1 Star 38

<> Code Issues Pull requests Actions Projects Security Insights

main Go to file <> Code

rossanoventurini Update README.md 5efa741 · 2 months ago

imgs	code	3 months ago
scripts	update conversion script	3 months ago
src	code	3 months ago
.gitignore	code	3 months ago
.pre-commit-config.yaml	code	3 months ago
Cargo.toml	Update Cargo.toml	2 months ago
LICENSE.md	code	3 months ago
README.md	Update README.md	2 months ago
pyproject.toml	code	3 months ago
rust-toolchain.toml	code	3 months ago

Readme

MIT license

Activity

38 stars

7 watching

1 fork

Report repository

Releases 1

SIGIR2024 Latest on Jul 4

Packages

No packages published

Contributors 3


rossanoventurini Rossano Ven...

francomarianardini Franco Ma...

CosimoRulli Cosimo Rulli

README

MIT license




Seismic


paper SIGIR 2024 arXiv 2404.18812

crates.io v0.1.0 downloads 377 license MIT


Bi-Encoder



One Dense Vector
Per Sequence
e.g., DPR




One **Sparse** Vector
Per Sequence
e.g., SPLADE



Multiple Dense Vectors
Per Sequence
e.g., ColBERT

Cross Encoder

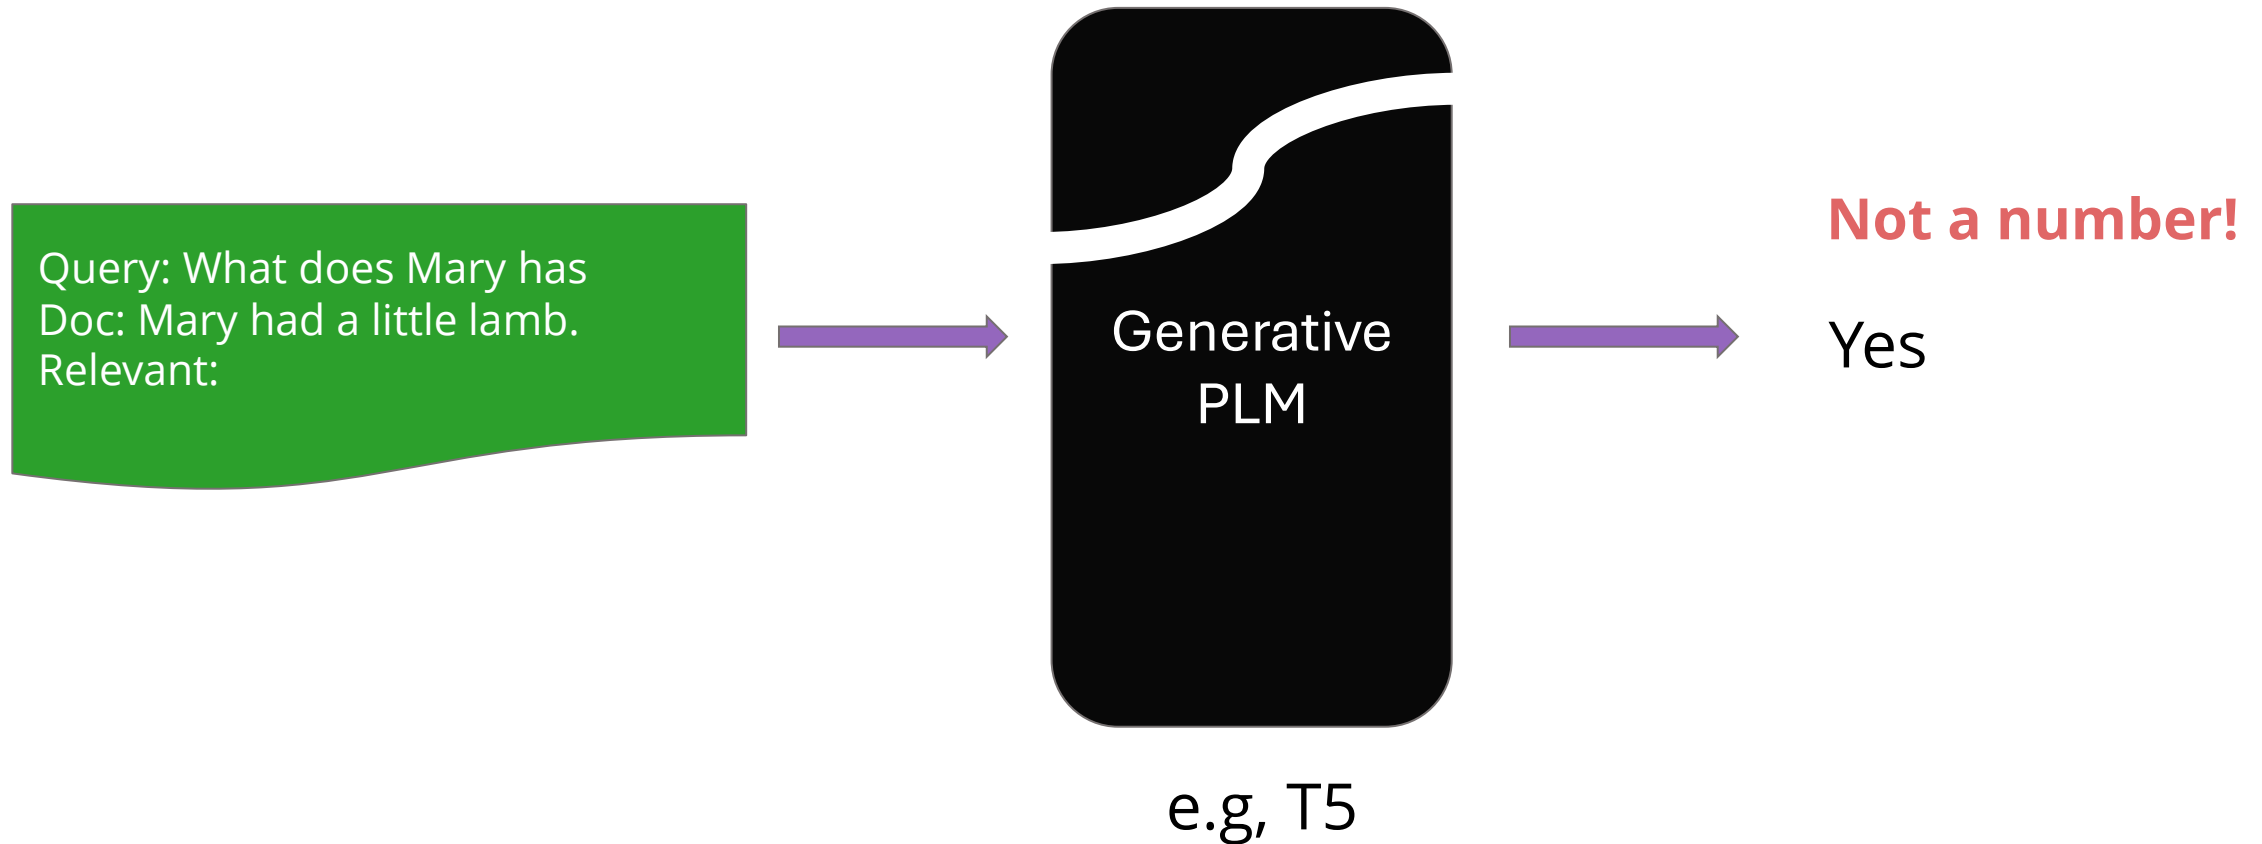


Joint Encoder
e.g., monoBERT

Cross-Encoder

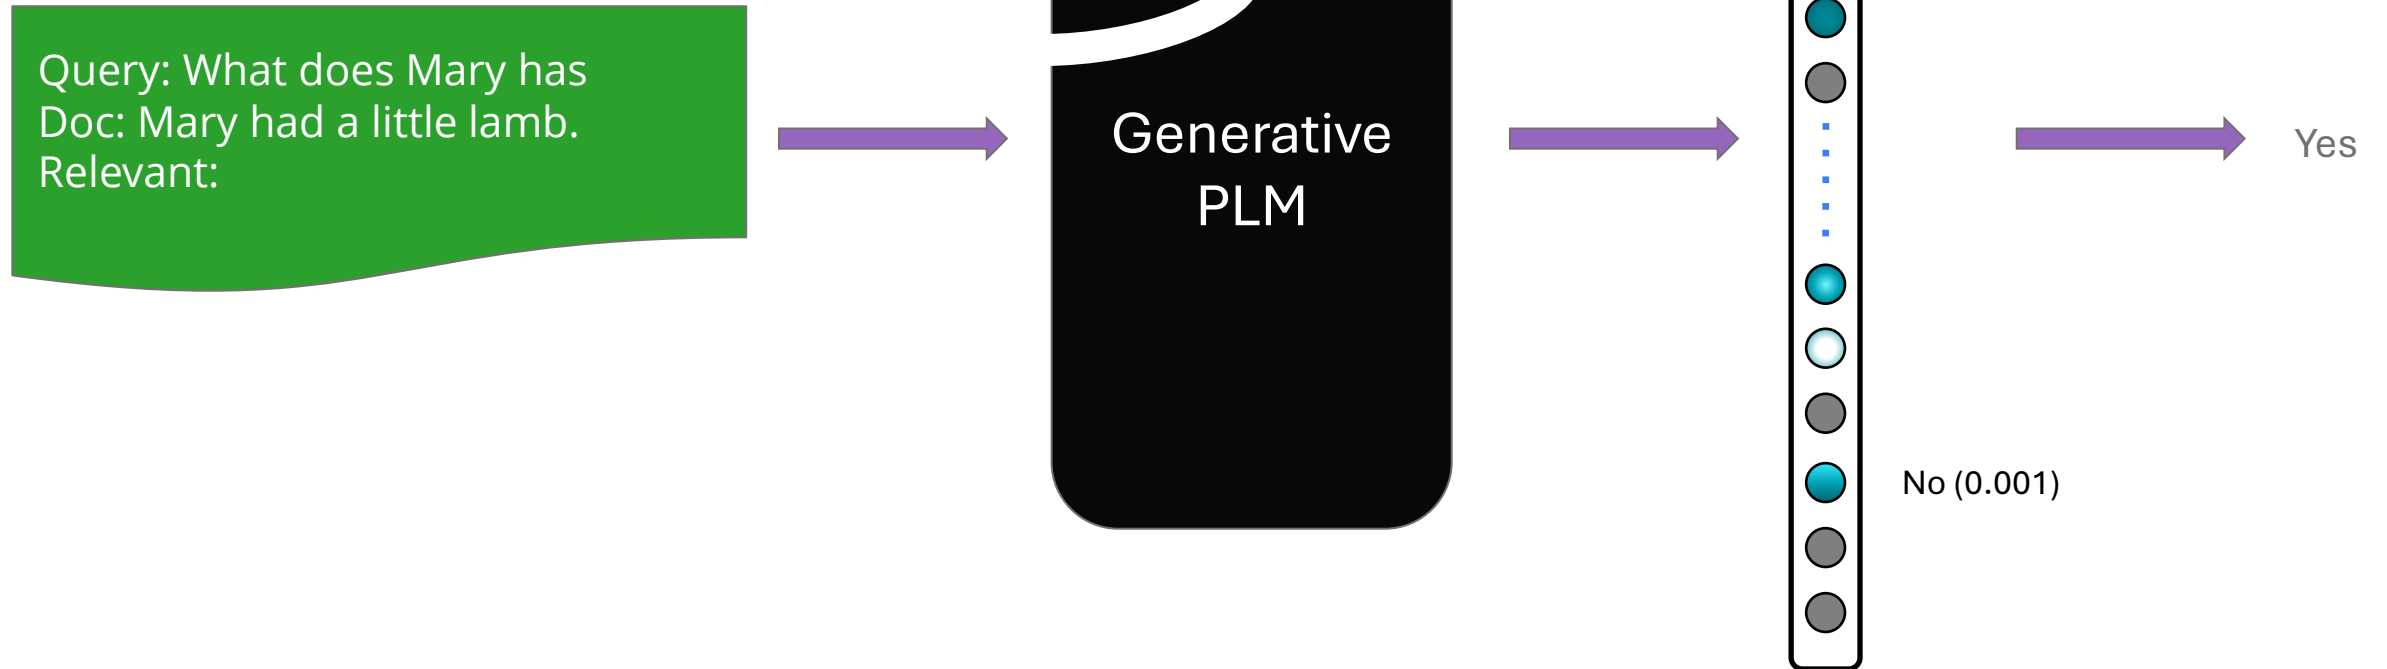


Using Generative Models



Pradeep, Ronak, Rodrigo Nogueira, and Jimmy Lin. "The expando-mono-duo design pattern for text ranking with pretrained sequence-to-sequence models." arXiv preprint arXiv:2101.05667 (2021).

Using Generative Models




Pradeep, Ronak, Rodrigo Nogueira, and Jimmy Lin. "The expando-mono-duo design pattern for text ranking with pretrained sequence-to-sequence models." arXiv preprint arXiv:2101.05667 (2021).

Using Generative Models

Pointwise score

Query: What does Mary has
Doc: Mary had a little lamb.
Relevant:

Generative
PLM



Yes (0.08)

No (0.001)




0.52

Pradeep, Ronak, Rodrigo Nogueira, and Jimmy Lin. "The expando-mono-duo design pattern for text ranking with pretrained sequence-to-sequence models." arXiv preprint arXiv:2101.05667 (2021).

Using Generative Models

Pairwise score

Query: What does Mary has
Doc0: JHU is in Baltimore
Doc1: Mary had a little lamb.
Relevant:



Yes (0.001)


No (0.02)




0.49

Pradeep, Ronak, Rodrigo Nogueira, and Jimmy Lin. "The expando-mono-duo design pattern for text ranking with pretrained sequence-to-sequence models." arXiv preprint arXiv:2101.05667 (2021).


Bi-Encoder



One Dense Vector
Per Sequence
e.g., DPR




One **Sparse** Vector
Per Sequence
e.g., SPLADE




Multiple Dense Vectors
Per Sequence
e.g., ColBERT

Cross Encoder




Joint Encoder
e.g., monoBERT


Bi-Encoder



One Dense Vector
Per Sequence
e.g., DPR




One **Sparse** Vector
Per Sequence
e.g., SPLADE



Multiple Dense Vectors
Per Sequence
e.g., ColBERT

Cross Encoder



Joint Encoder
e.g., monoBERT

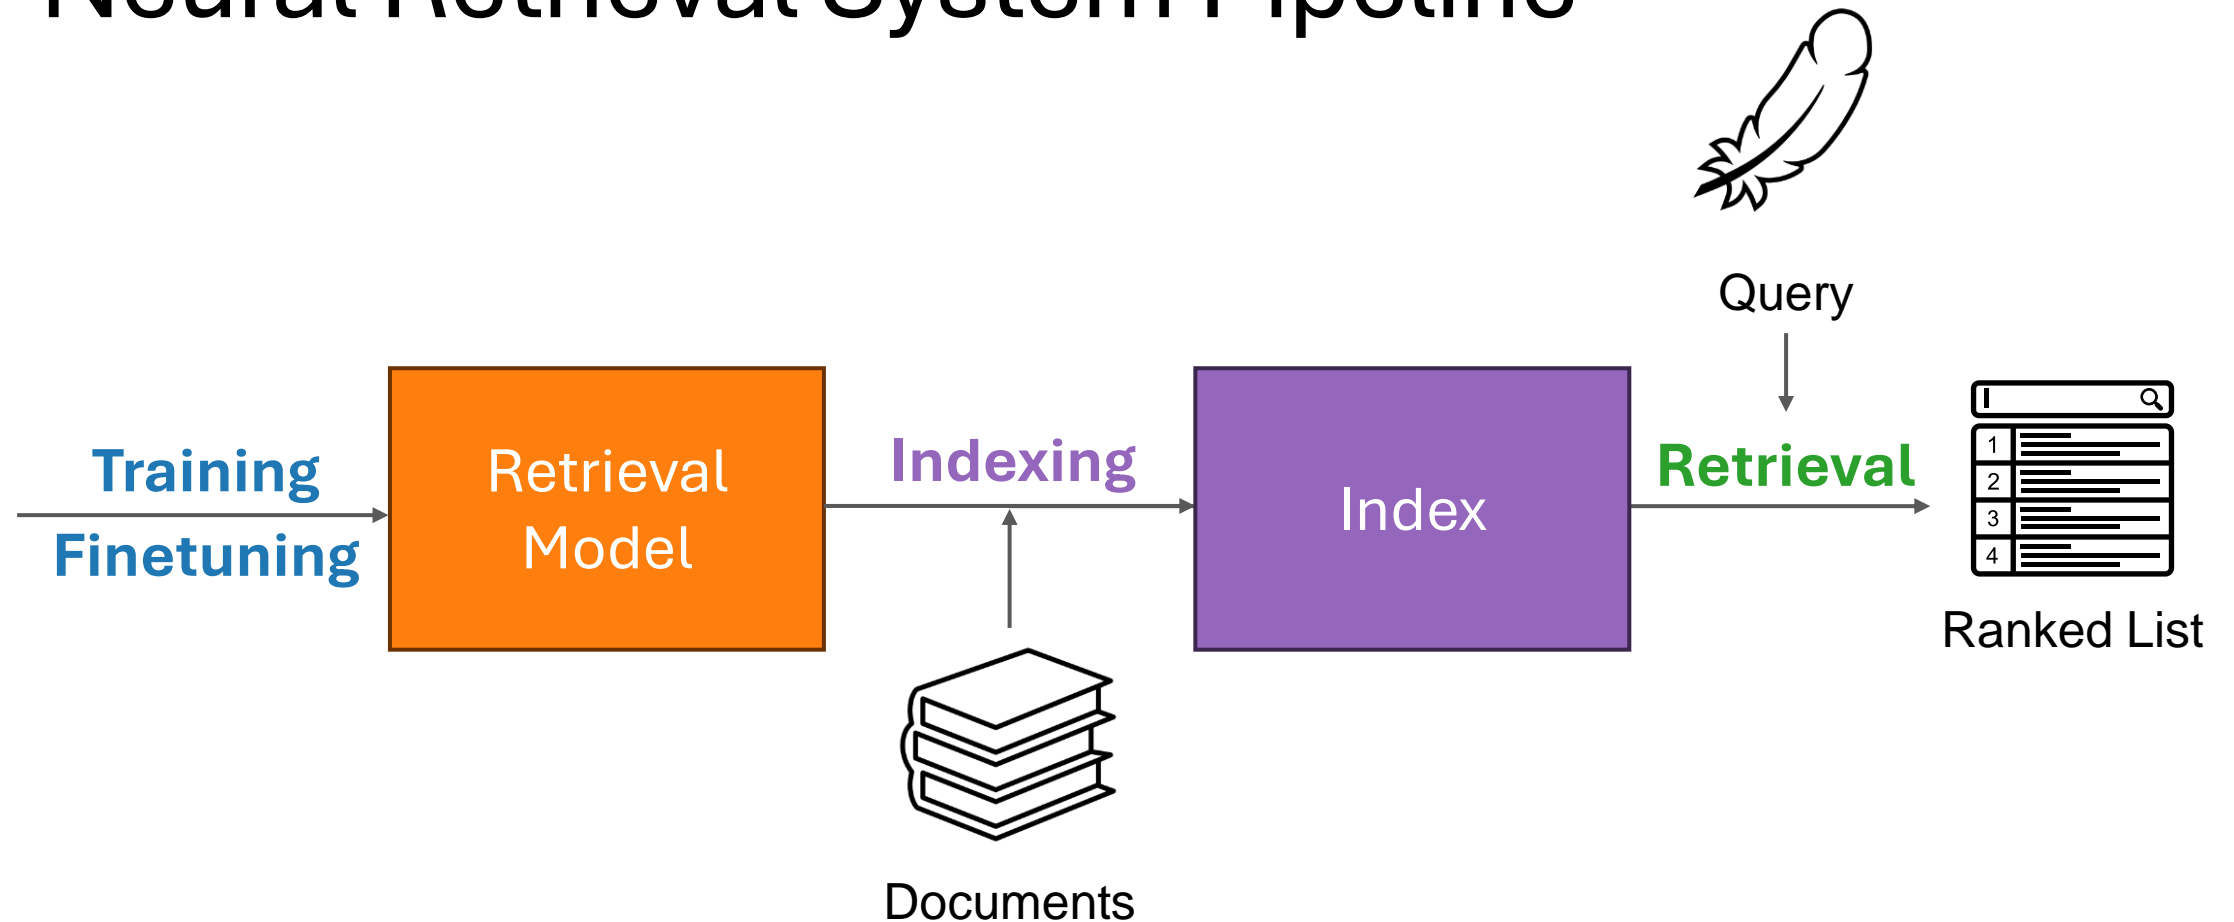
More Effective

More Efficient


Retrieve-and-Rerank System Combinations



Neural Retrieval System Pipeline



PLM to IR Model



- Align the representation
- Model “relevancy”

Evaluation

Which system is better?

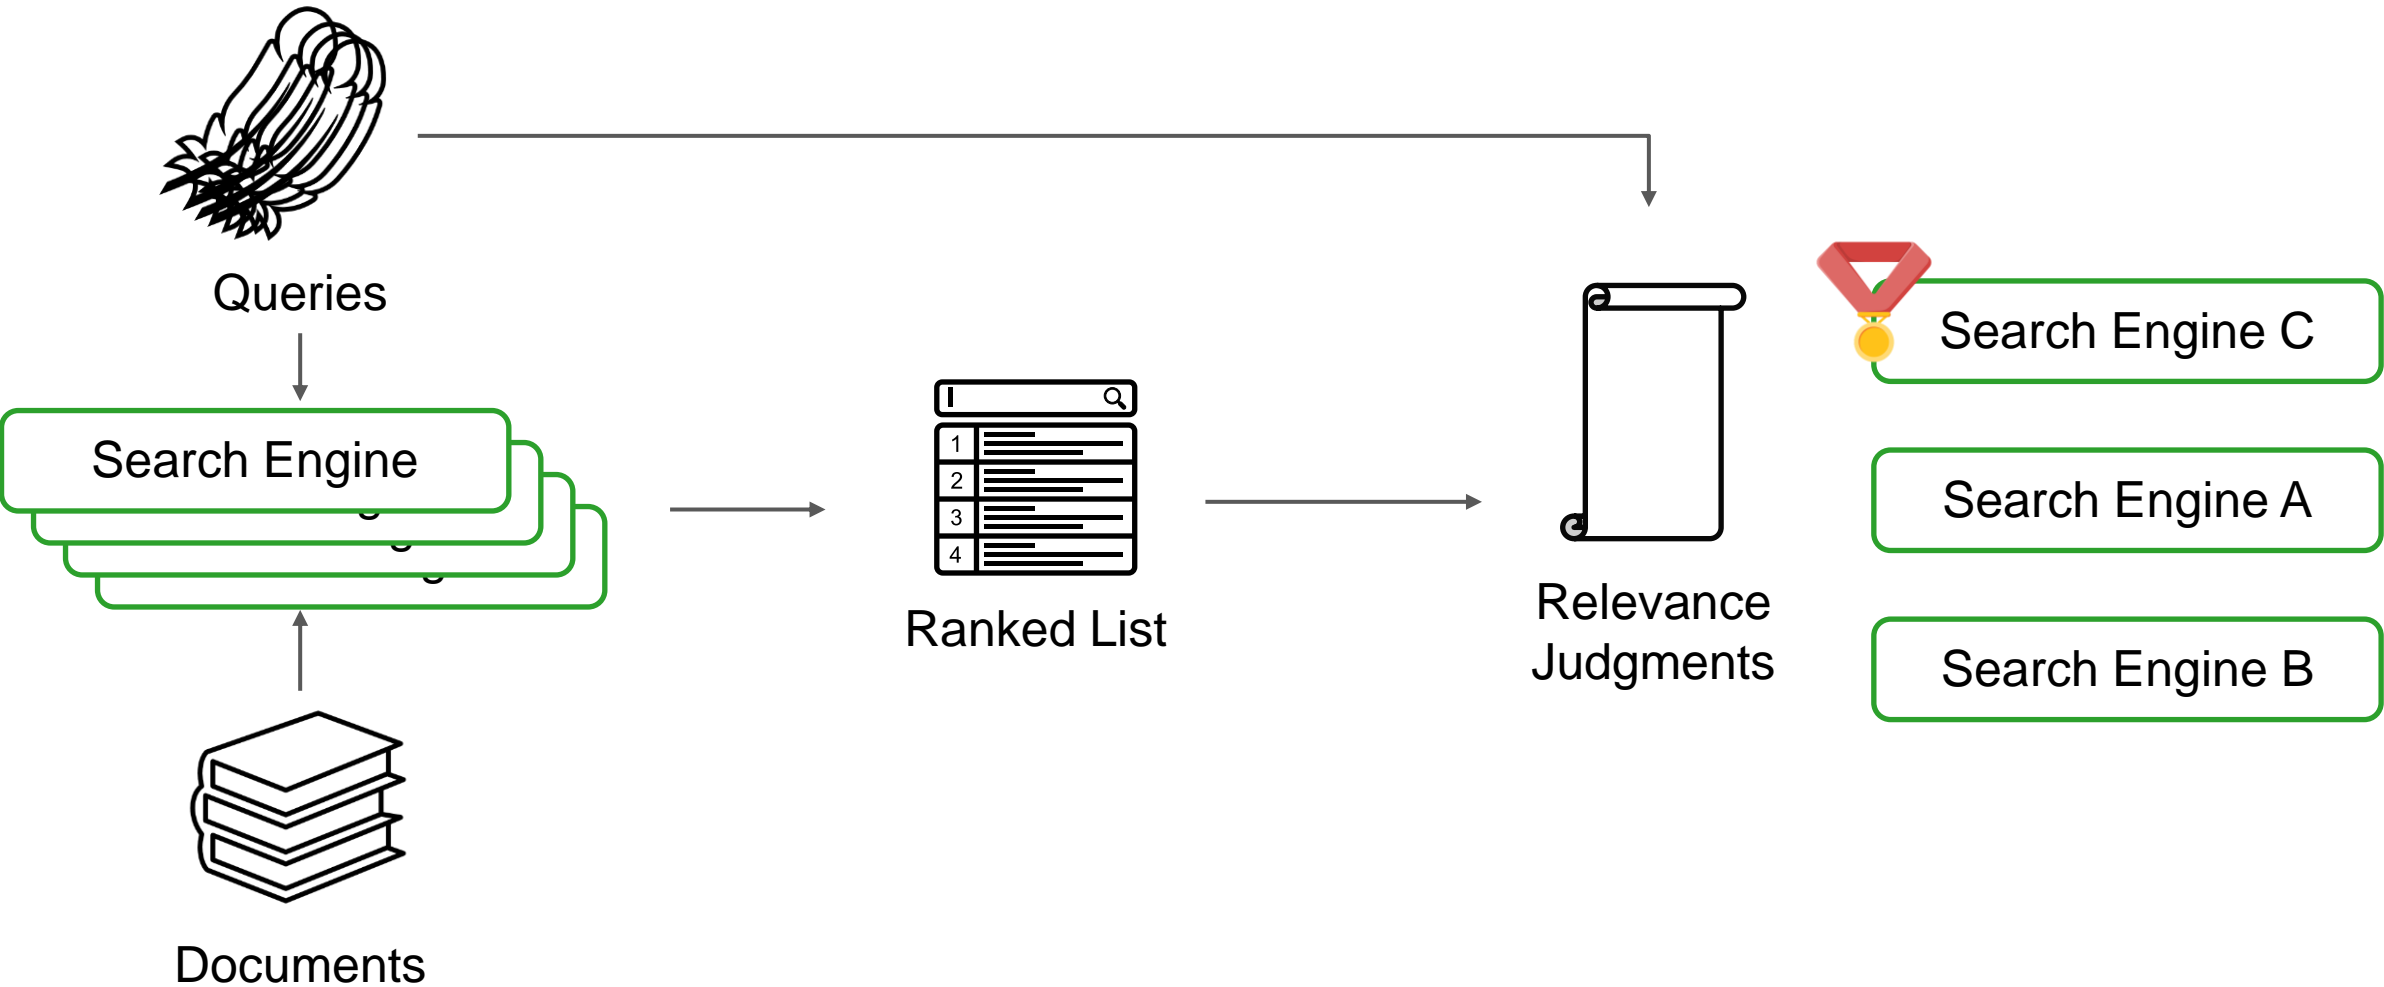
What is Information Retrieval?

(relevant)

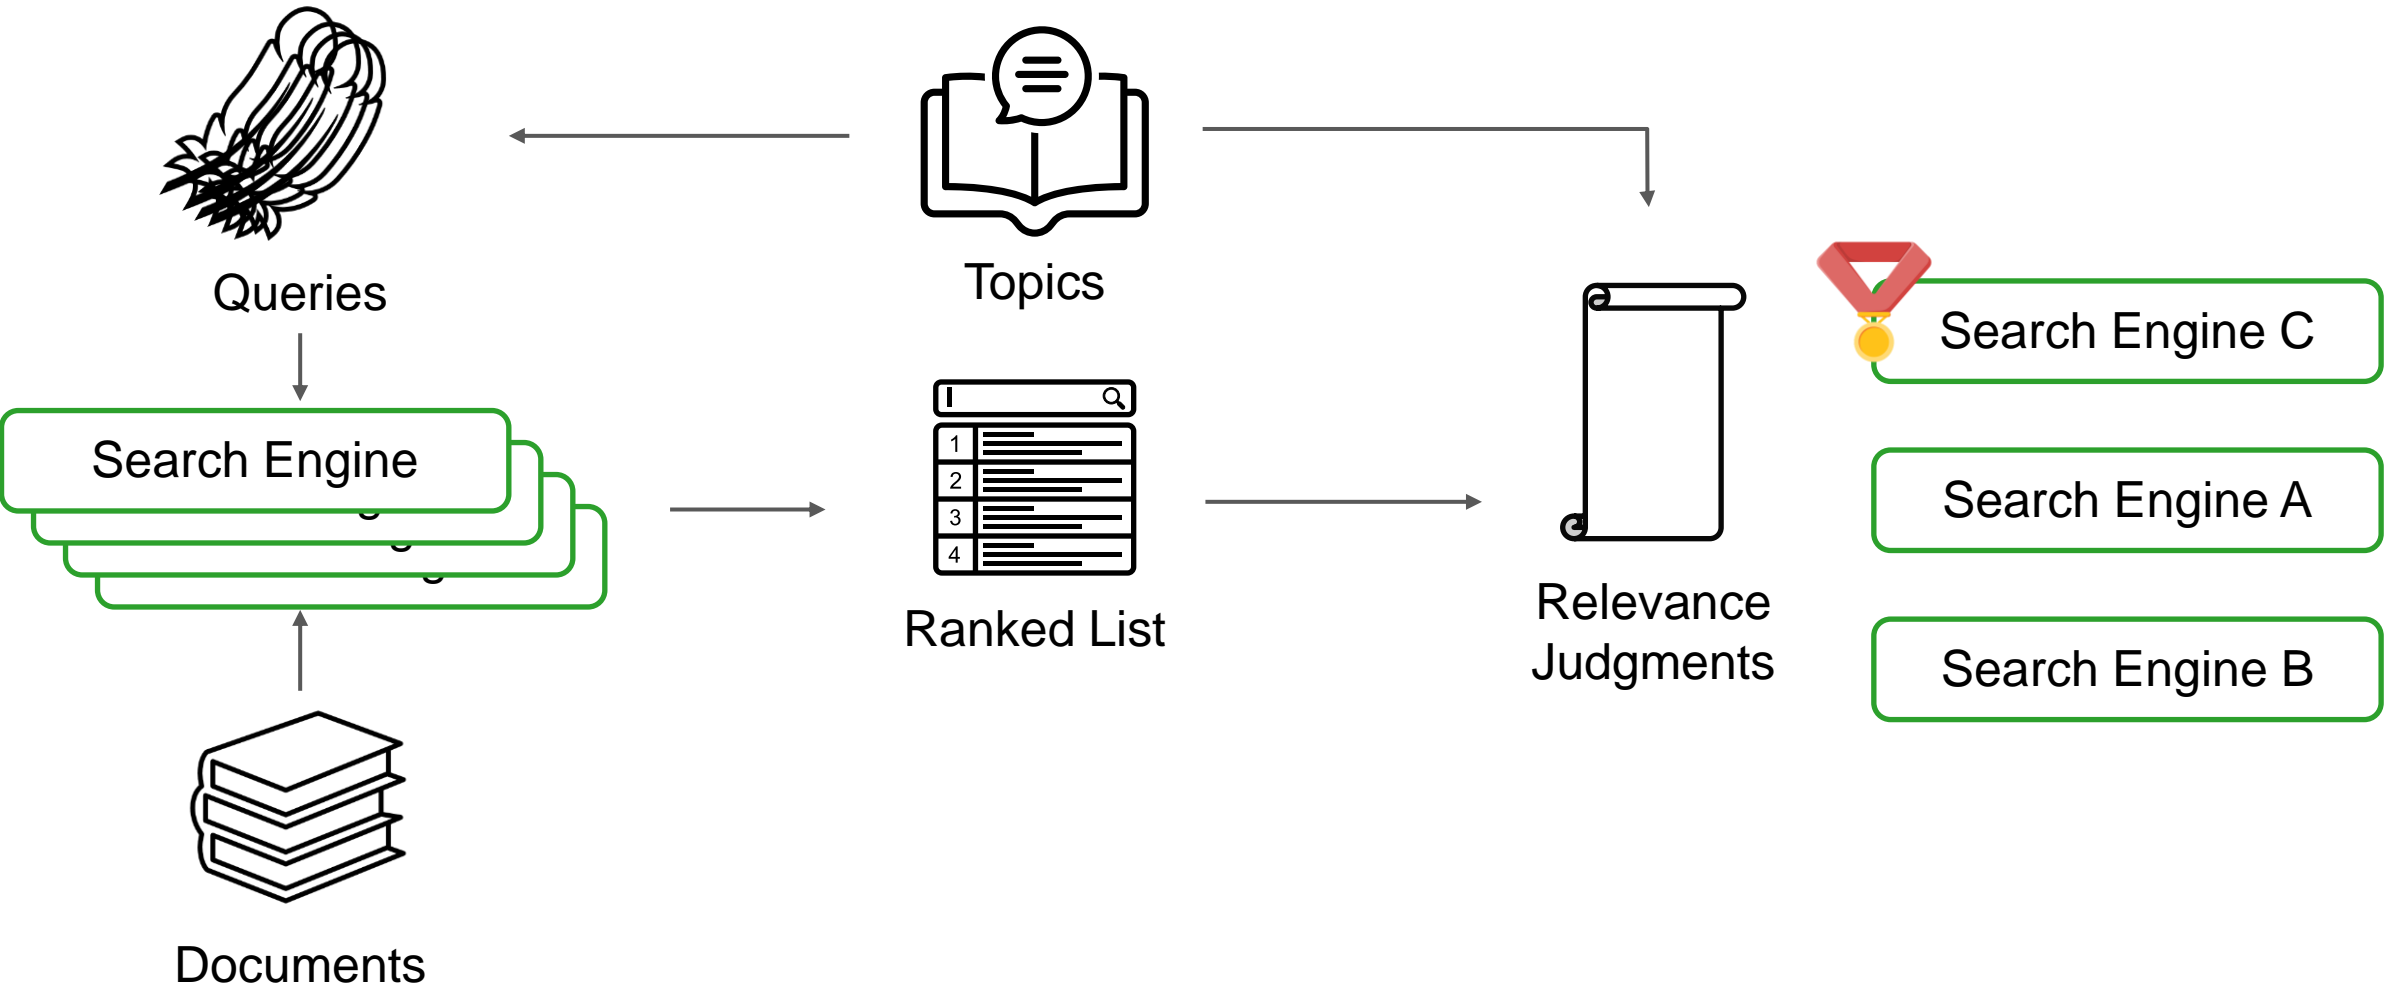
Retrieve information from a storage
based on user's information need

Which system retrieve more relevant information?

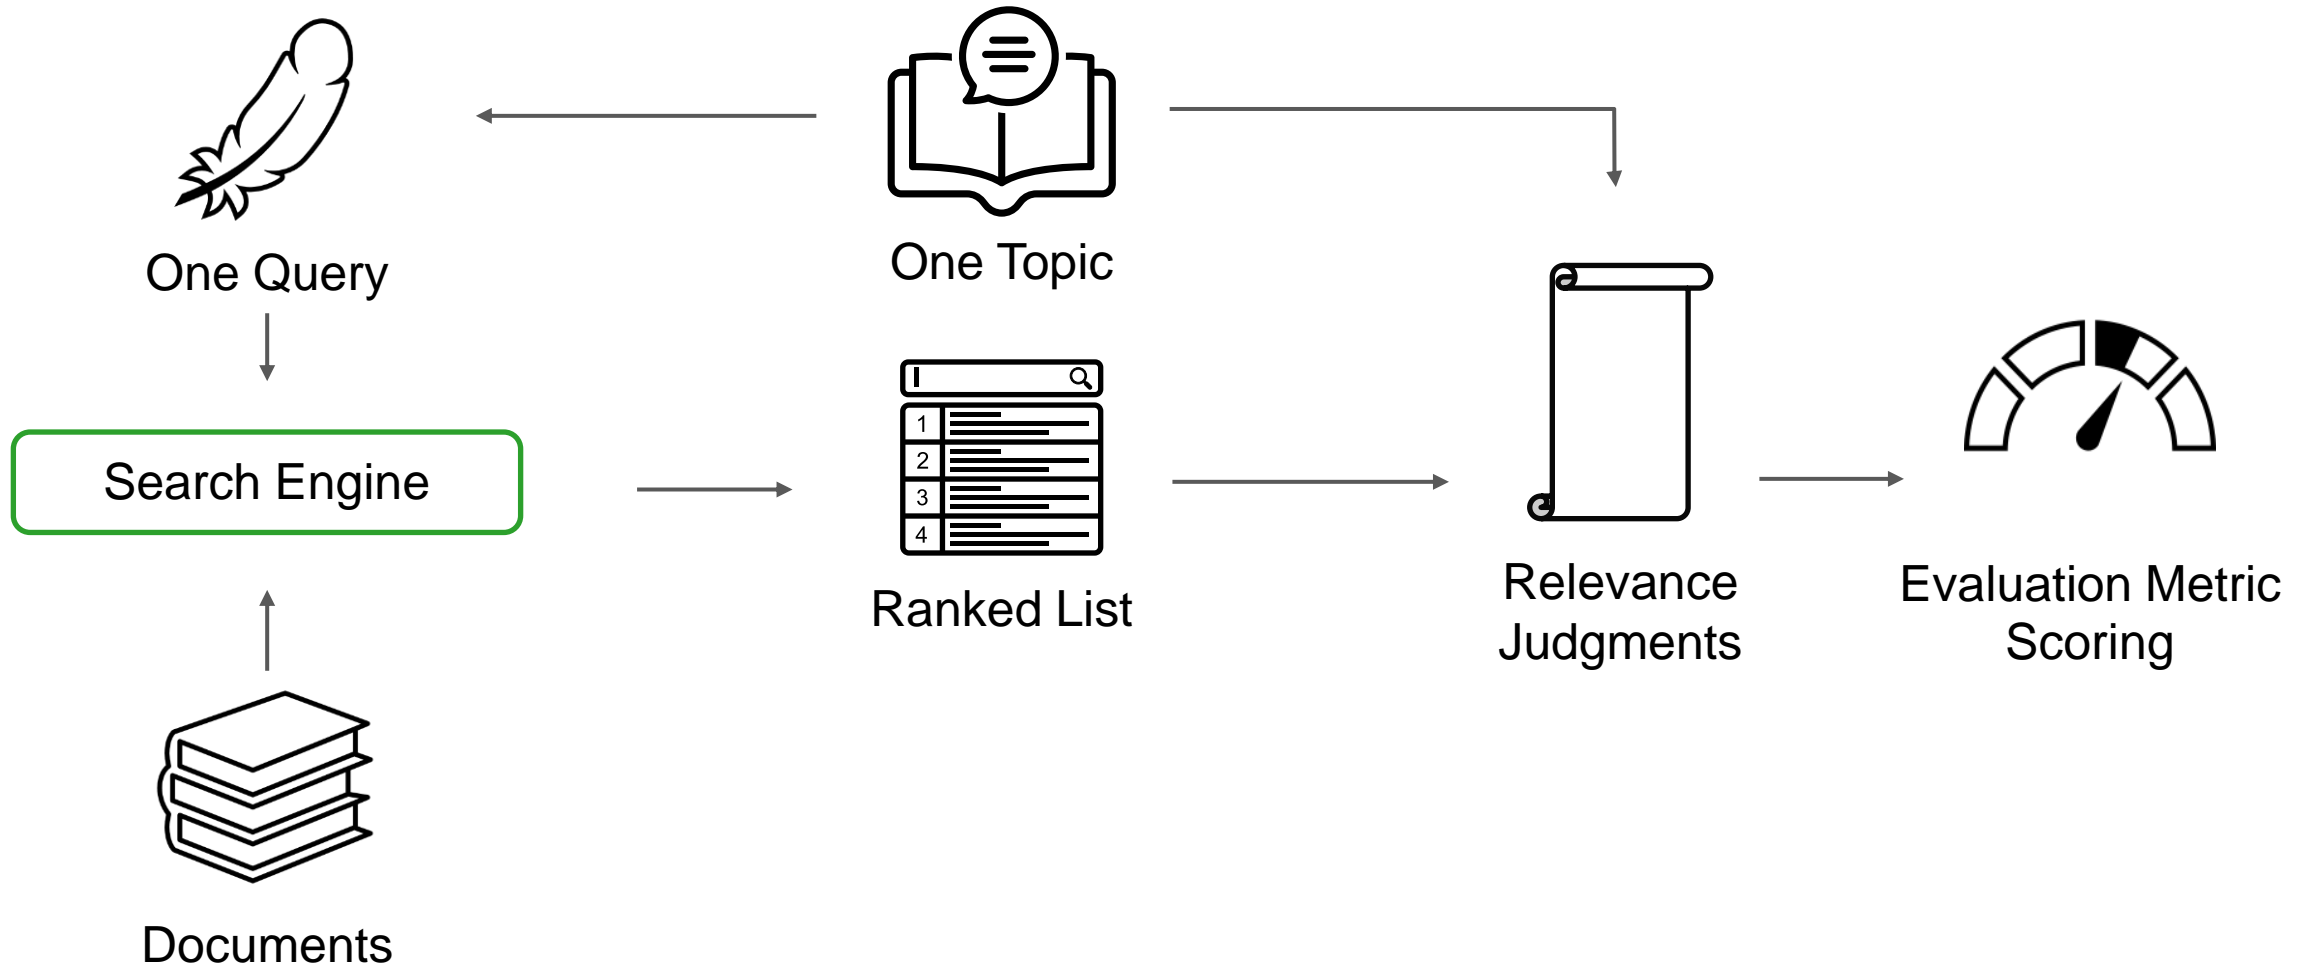
Cranfield Paradigm Evaluation



Cranfield Paradigm Evaluation



Cranfield Paradigm Evaluation



Differences

- Topics vs Queries
 - Clear intent vs an expression of such intent
- Relevance Judgements vs Labels
 - Opinion vs “fact”
- Ranked retrieval metrics
 - Measuring the quality/effectiveness of a ranked list

IR Metrics

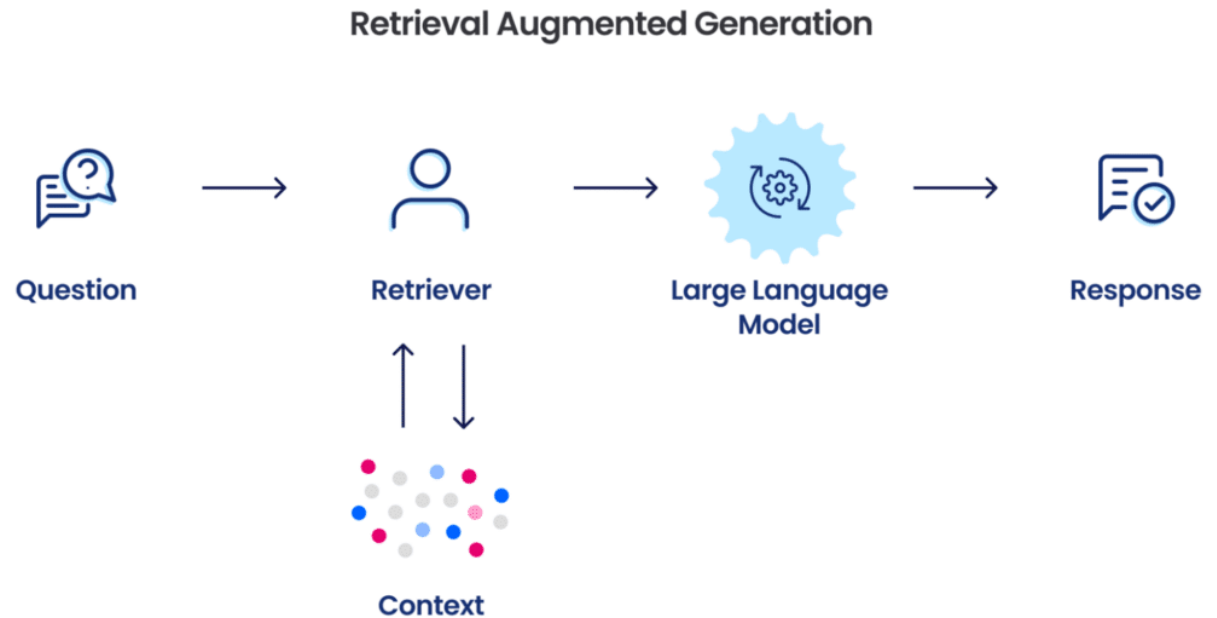


- Effective Metrics
 - Mean Average Precision
 - Normalized Discounted Cumulative Gain
 - Recall@k
- Efficiency Metrics
 - Indexing time
 - Index disk space
 - Query latency (average search time per query)


State of IR Research

Retrieval-Augmented Generation

- Is everything a RAG problem?
- What is the right retrieval model/system for RAG?
- IR going away?




<https://snorkel.ai/which-is-better-retrieval-augmentation-rag-or-fine-tuning-both/>



Better Retrieval Models

- More effective
 - Better/larger neural models
 - Better architecture?
 - Under harder setup, e.g., scholar search, multilingual, cross-modal, etc
- More efficient
 - Faster at query time
 - Less resource footprint, e.g., memory, storage, compute, etc
- Other qualities
 - Fairness, diversity, etc



Other Retrieval Problems

- Conversational
 - Guessing intent, finding the “right” information to serve
- Iterative/interactive/human-in-the-loop
 - Rounds of interactions
- Generative
 - Returning a piece of text



Evaluation

- What to measure
 - and when would it fail
- How to measure
 - Generative text? Citations?
- “Better” evaluation collection
 - Not necessarily larger